



KVM PV DEVICES

dor.laor@qumranet.com


Agenda

- Introduction & brief history
- VirtIO
- Enhanced VirtIO with KVM support
- Further implementation

General & history

- Fully virtualized devices performs bad
 - 55 Mbps for RTL
 - Lots of io-exits per packet
- Decided to implement a modern e1000
 - Advantage:
 - Only Qemu coding
 - no guest tools involved
 - IRQ coalescing
 - Only 2-3 io-exits per packet
 - Can be the base of user-space PV
- But then came Ingo...

PV driver architecture



General & history

- V0 - leveraging Ingo Molnar's PV code
 - Make loadable module
 - Add HVM support
 - Add NAPI
 - Add memory barriers and improved ring
 - Keep running after performance & stability
 - Merge to the kernel?
- Alternatives
 - Xen
 - Polished drivers
 - Xen specific
 - VirtIO was just published.


PV driver requirements

- Close to native performance
- Merge with the kernel
- Leverage existing code
- Have usermode implementation
 - Block device for qcow, vmdk formats
 - Ability to function without KVM (-no-kvm)
- Ability to run HVM guests
- Robust code
 - Easily add drivers

VirtIO

- An API for virtual I/O
 - Implements network & block driver logic
 - Written by Rusty Russell
- Motivation
 - Many hypervisors of all types
 - Hard to tune and maintain each one
 - Code reuse – The KVM way ;)
- Implementations
 - Lguest
 - KVM
 - Possible (Xen, UML, Qemu, VMware?..)

VirtIO




- scatterlist sg[] for skb/blk_req data

VirtIO – Hypervisor specifics

- The front end logic is implemented by VirtIO

- The backend needs
 - Probing & Bus services
 - Enumeration
 - IRQ
 - Parameters (mac,...)
 - Shared memory with remote side
 - Hypercalls
 - Host driver/userspace backend




Enhanced VirtIO

■ Motivation


- Increase re-use
- Allow operation with various bus types
- Make new devices code smallest

■ Components

- Shared memory code
 - With per hypervisor I/O hypercalls
- Bus (pci, virtual bus)
- Host backend



Enhanced VirtIO




Enhanced VirtIO

■ Status:

- Interface needs polishing
- KVM support
- PCI bus support


■ Result:

- Makes backend driver tiny
- 620 Mbps throughput for network
 - HVM Linux guest
 - Before optimization
 - Userspace backend driver




Enhanced VirtIO – shared memory backend

- VirtIO backend
 - Implements VirtIO interface
 - Callbacks to hypervisor and
- Code consists of
 - `add_buf`, `get_buf`, `restart`, `detach_buf`
 - Only shared memory logic needed
 - `sync`
 - Ring logic
 - IO pending hypercall
 - `be_virtqueue_interrupt handler`




Enhanced VirtIO – Day in a life of packet



Enhanced VirtIO – shared memory details

- Based on Iguest
- 1-1 shared memory
- Data structure
 - Page of descriptors for rx, tx.
 - Available pointers page controlled by guest
 - Used pointers page controlled by host
- SG list is currently internal to descriptors
 - Descriptors are chained by next pointer



Enhanced VirtIO – network be driver

- Implements kvm_virtnet_probe for pci bus

- Creates tx,rx be_new_virtqueue


- Probes virtnet

- Request_irq

- IRQ# taken from bus


- Register hypercall - shared memory pfn

- Device key for enumeration taken from bus



Further work

- Basic
 - Add readv/writev handlers to Qemu
 - Complete the user-space block device
 - Update with VirtIO gso.
 - Complete migration support
- Advanced (also simple)
 - Publish the enhanced interface
 - Optimize and stabilize
 - Add host back end drivers
 - Add virtual bus
 - HVM improvements
 - Test with PV kernel



Q&A

Thank you ;)