



Using and Understanding the Real-Time Cyclictest Benchmark

Cyclictest results are the most frequently cited real-time Linux metric. The core concept of Cyclictest is very simple. However the test options are very extensive. The meaning of Cyclictest results appear simple but are actually quite complex. This talk will explore and explain the complexities of Cyclictest. At the end of the talk, the audience will understand how Cyclictest results describe the potential real-time performance of a system.

What Cyclicttest Measures

Latency of response to a stimulus.

external interrupt triggers (clock expires)

- possible delay until IRQs enabled
- IRQ handling
- cyclicttest is woken
- possible delay until preemption enabled
- possible delay until cyclicttest is highest priority
- possible delay until other process is preempted
- scheduler overhead

transfer control to cyclicttest

What Cyclictest Measures

Latency of response to a stimulus.

Causes of delay list on previous slide is simplified:

- order will vary
- may occur multiple times
- there are additional causes of delay

Many factors can increase latency

- additional external interrupts
- SMI
- processor emerging from sleep states
- cache migration of data used by woken process
- block on sleeping lock
 - lock owner gets priority boost
 - lock owner schedules
 - lock owner completes scheduled work
 - lock owner releases lock, loses priority boost

How Cyclictest Measures Latency

(Cyclictest Pseudocode)

The source code is nearly 3000 lines,
but the algorithm is trivial

Test Loop

```
clock_gettime(&now)
next = now + par->interval

while ( !shutdown) {

    clock_nanosleep(&next)

    clock_gettime(&now)
    diff = calcdiff(now, next)

    # update stat-> min, max, total latency, cycles
    # update the histogram data

    next += interval
}
```

The Magic of Simple

This trivial algorithm captures all of the factors that contribute to latency.

Mostly. Caveats will follow soon.

Cyclictest Program

```
main() {  
  
    for (i = 0; i < num_threads; i++) {  
        pthread_create((timerthread))  
  
    while (!shutdown) {  
        for (i = 0; i < num_threads; i++)  
            print_stat((stats[i]), i))  
        usleep(10000)  
    }  
  
    if (histogram)  
        print_hist(parameters, num_threads)  
}
```

timerthread()

```
*timerthread(void *par) {  
    # thread set up  
    # test loop  
}
```

Thread Set Up

```
stat = par->stats;
pthread_setaffinity_np((pthread_self()))
setscheduler(({par->policy, par->priority})
sigprocmask((SIG_BLOCK))
```

Test Loop (as shown earlier)

```
clock_gettime(&now)
next = now + par->interval

while ( !shutdown) {

    clock_nanosleep(&next)

    clock_gettime(&now)
    diff = calcdiff(now, next)

    # update stat-> min, max, avg, cycles
    # Update the histogram

    next += interval
}
```

Why show set up pseudocode?

The timer threads are not in lockstep from time zero.

Multiple threads will probably not directly impact each other.

The Magic of Simple

This trivial algorithm captures all of the factors that contribute to latency.

Mostly. Caveats, as promised.

Caveats

Measured maximum latency is a floor of the possible maximum latency

- Causes of delay may be partially completed when timer IRQ occurs
- Cyclictest wakeup is on a regular cadence, may miss delay sources that occur outside the cadence slots

Caveats

Does not measure the IRQ handling path of the real RT application

- timer IRQ handling typically fully in IRQ context
- normal interrupt source IRQ handling:
 - irq context, small handler, wakes IRQ thread
 - IRQ thread eventually executes, wakes RT process

Caveats

Cyclictest may not exercise latency paths that are triggered by the RT application, or even non-RT applications

- SMI to fixup instruction errata
- stop_machine()
 - module load / unload
 - hotplug

Solution 1

Do not use cyclictest. :-)

Instrument the RT application to measure latency

Solution 2

Run the normal RT application and non-RT applications as the system load

Run cyclictest with a higher priority than the RT application to measure latency

Solution 2

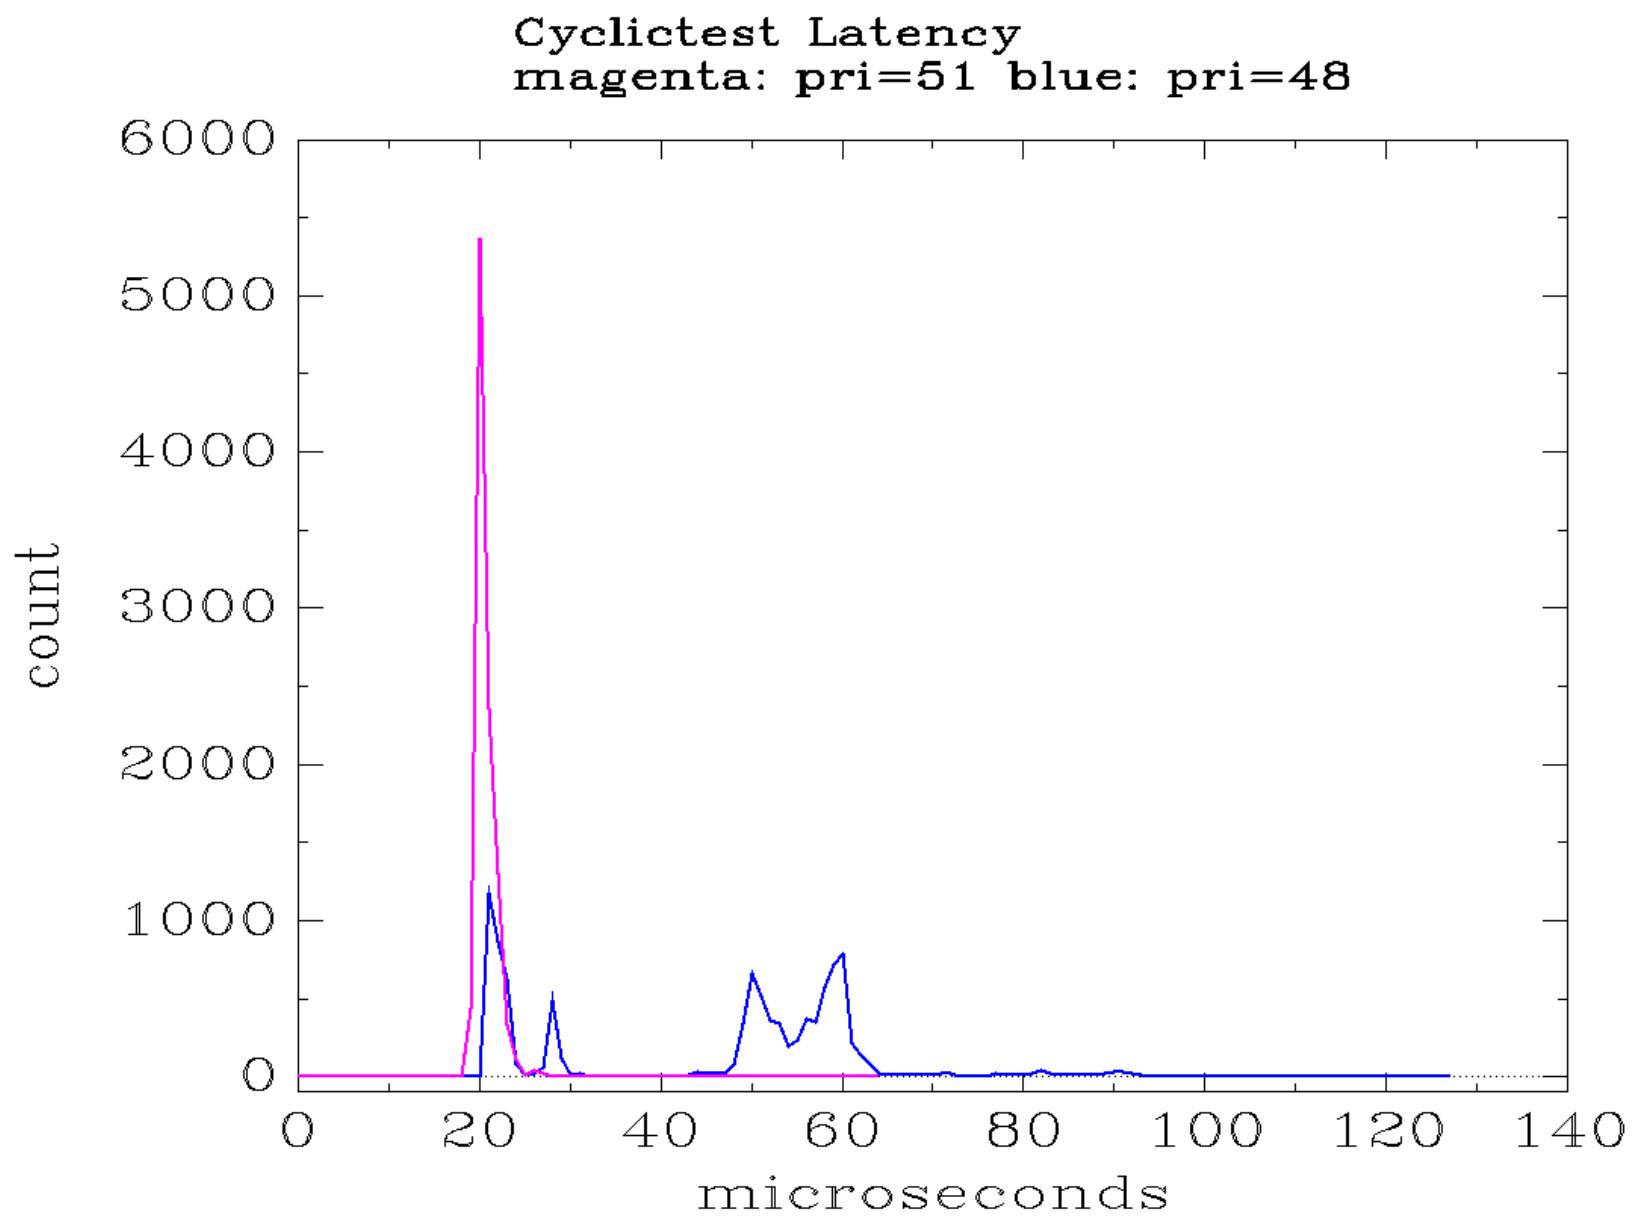
Typical real time application will consist of multiple threads, with differing priorities and latency requirements

To capture latencies of each of the threads, run separate tests, varying the cyclictest priority

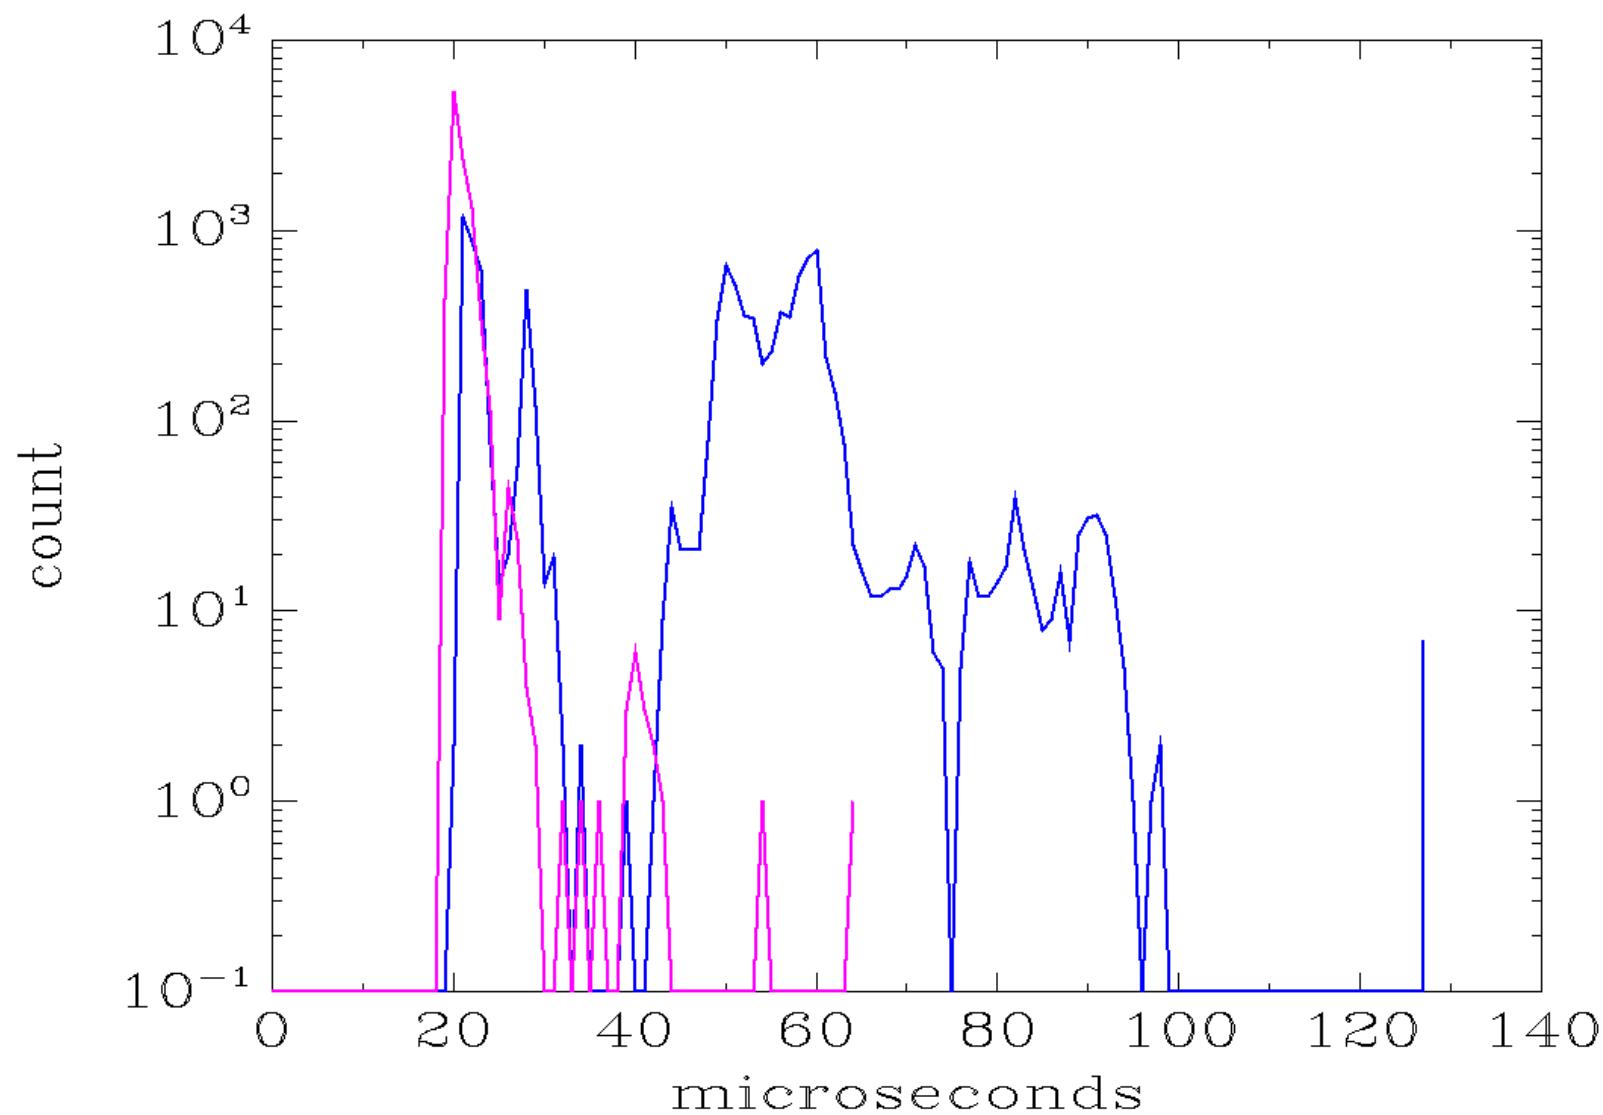
Solution 2

Example

<u>RT app thread</u>	<u>deadline constraint</u>	<u>latency constraint</u>	<u>RT app scheduler priority</u>	<u>cyclictest priority</u>
A	critical	80 usec	50	51
B	0.1% miss	100 usec	47	48



Cyclictest Latency
magenta: pri=51 blue: pri=48



Aside:

Cyclictest output in these slides is edited to fit on the slides

Original:

```
$ cyclictest_0.85 -l1000000 -q -p80 -S  
T: 0 ( 460) P:80 I:1000 C: 100000 Min: 37 Act: 43 Avg: 45 Max: 68  
T: 1 ( 461) P:80 I:1500 C: 66675 Min: 37 Act: 49 Avg: 42 Max: 72
```

Example of edit:

```
$ cyclictest_0.85 -l1000000 -q -p80 -S  
T:0 I:1000 Min: 37 Avg: 45 Max: 68  
T:1 I:1500 Min: 37 Avg: 42 Max: 72
```

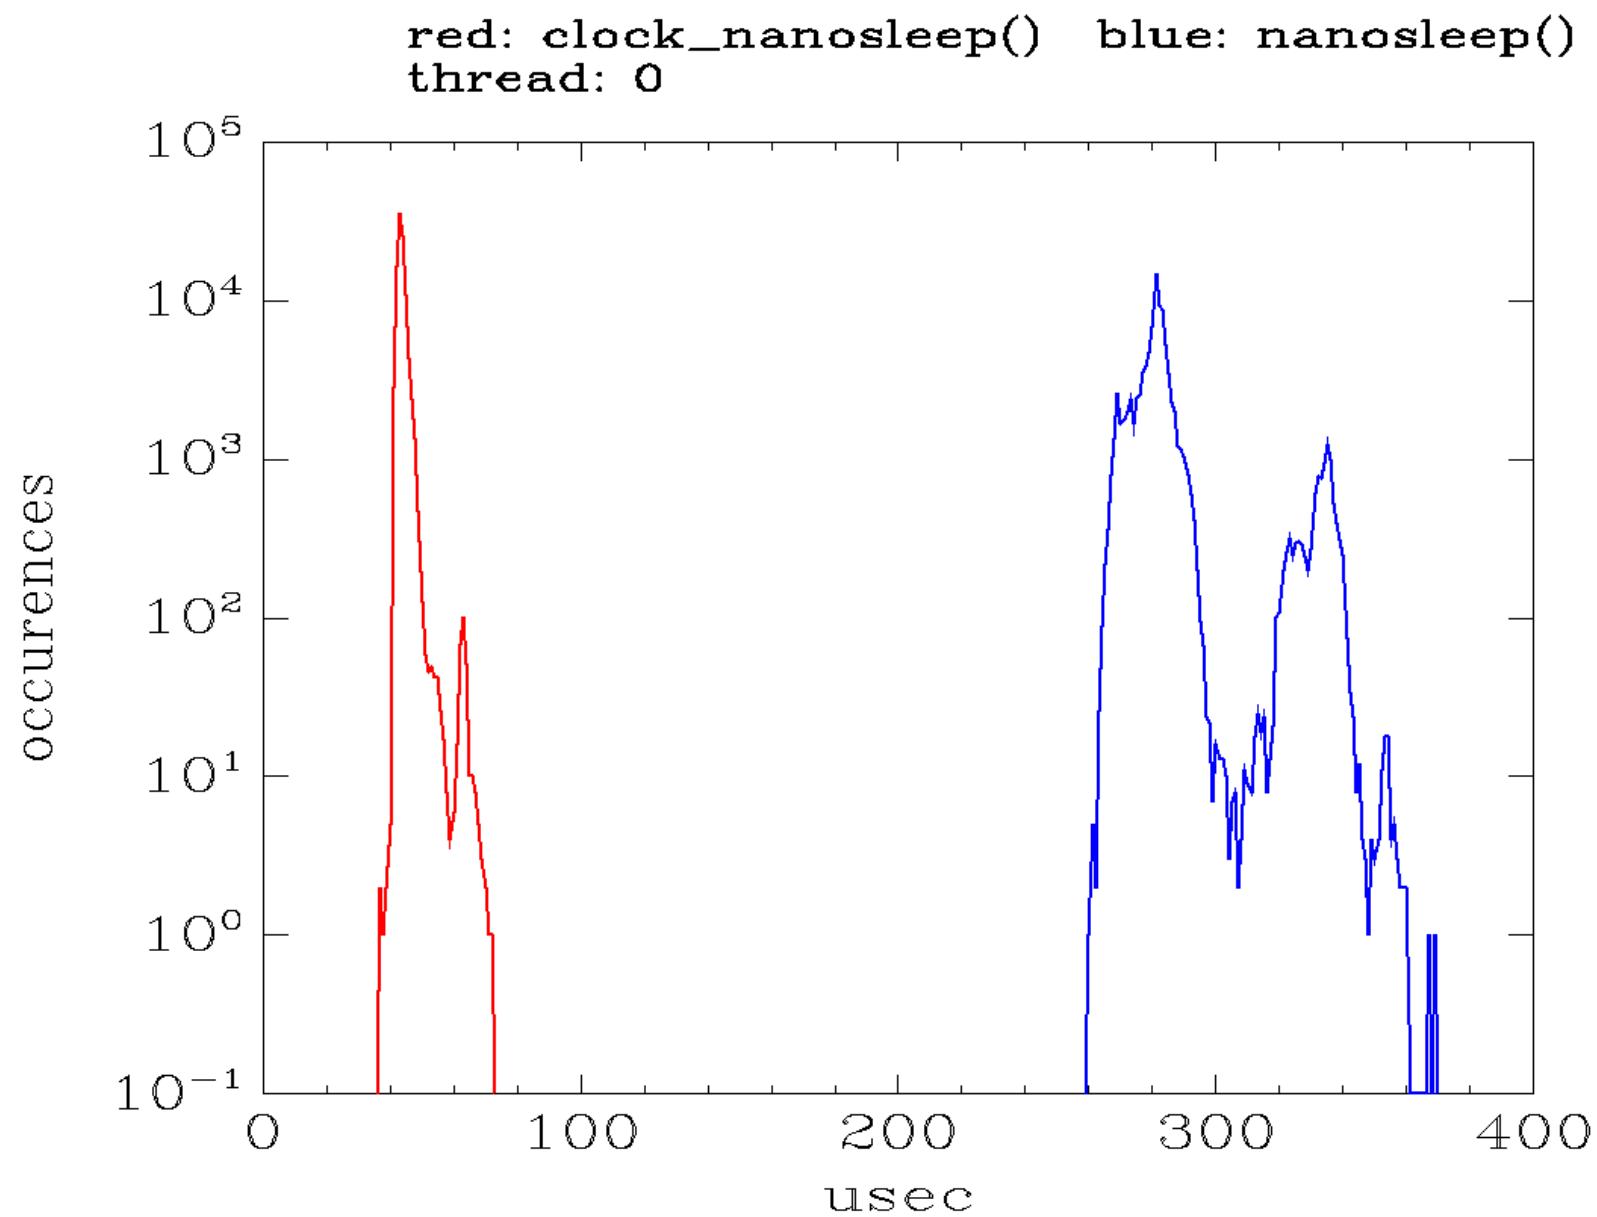
Cyclictest Command Line Options

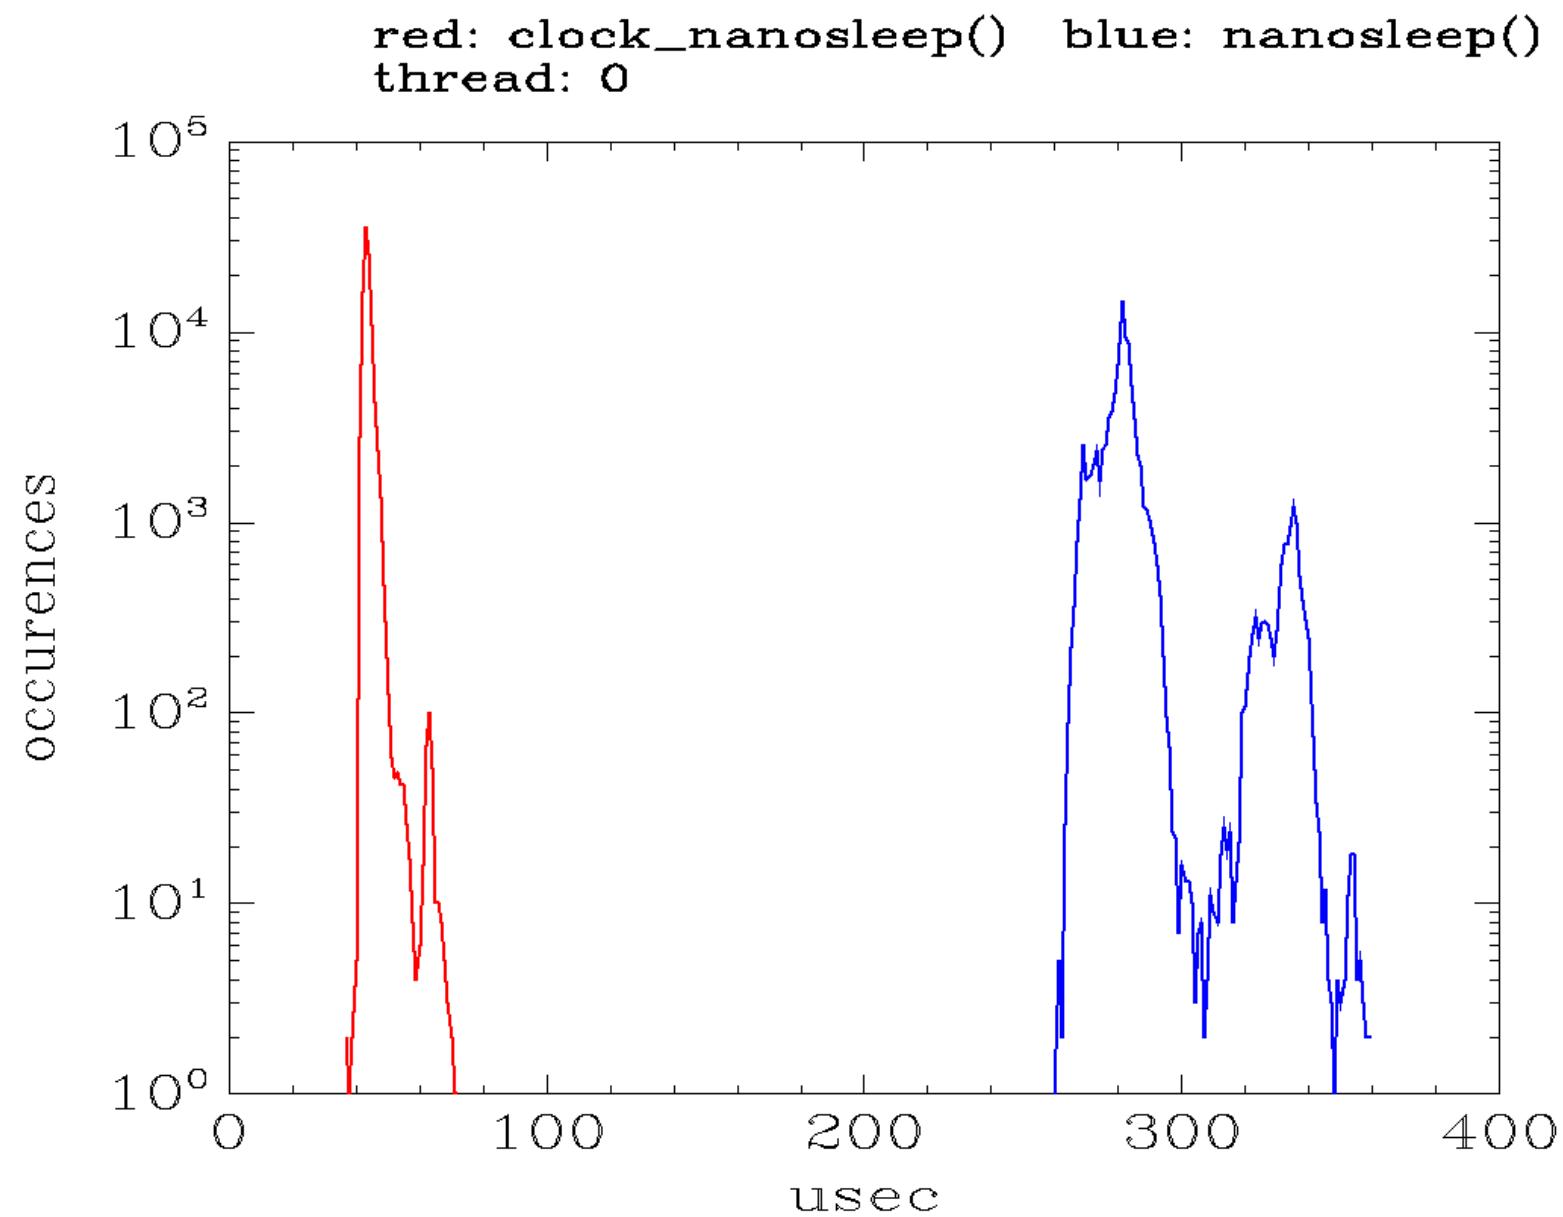
Do I really care???

Can I just run it with the default options???

Do I really care???

```
$ cyclictest_0.85 -l100000 -q -p80  
T:0 Min: 262 Avg: 281 Max: 337  
  
$ cyclictest_0.85 -l100000 -q -p80 -n  
T:0 Min: 35 Avg: 43 Max: 68  
  
-l100000 stop after 100000 loops  
-q quiet  
-p80 priority 80, SCHED_FIFO  
-n use clock_nanosleep()  
instead of nanosleep()
```





Impact of Options

More examples

Be somewhat skeptical of maximum latencies due to the short test duration.

Examples are:

100,000 loops

1,000,000 loops

Arbitrary choice of loop count. Need large values to properly measure maximum latency!!!

Priority of Real Time kernel threads for next two slides

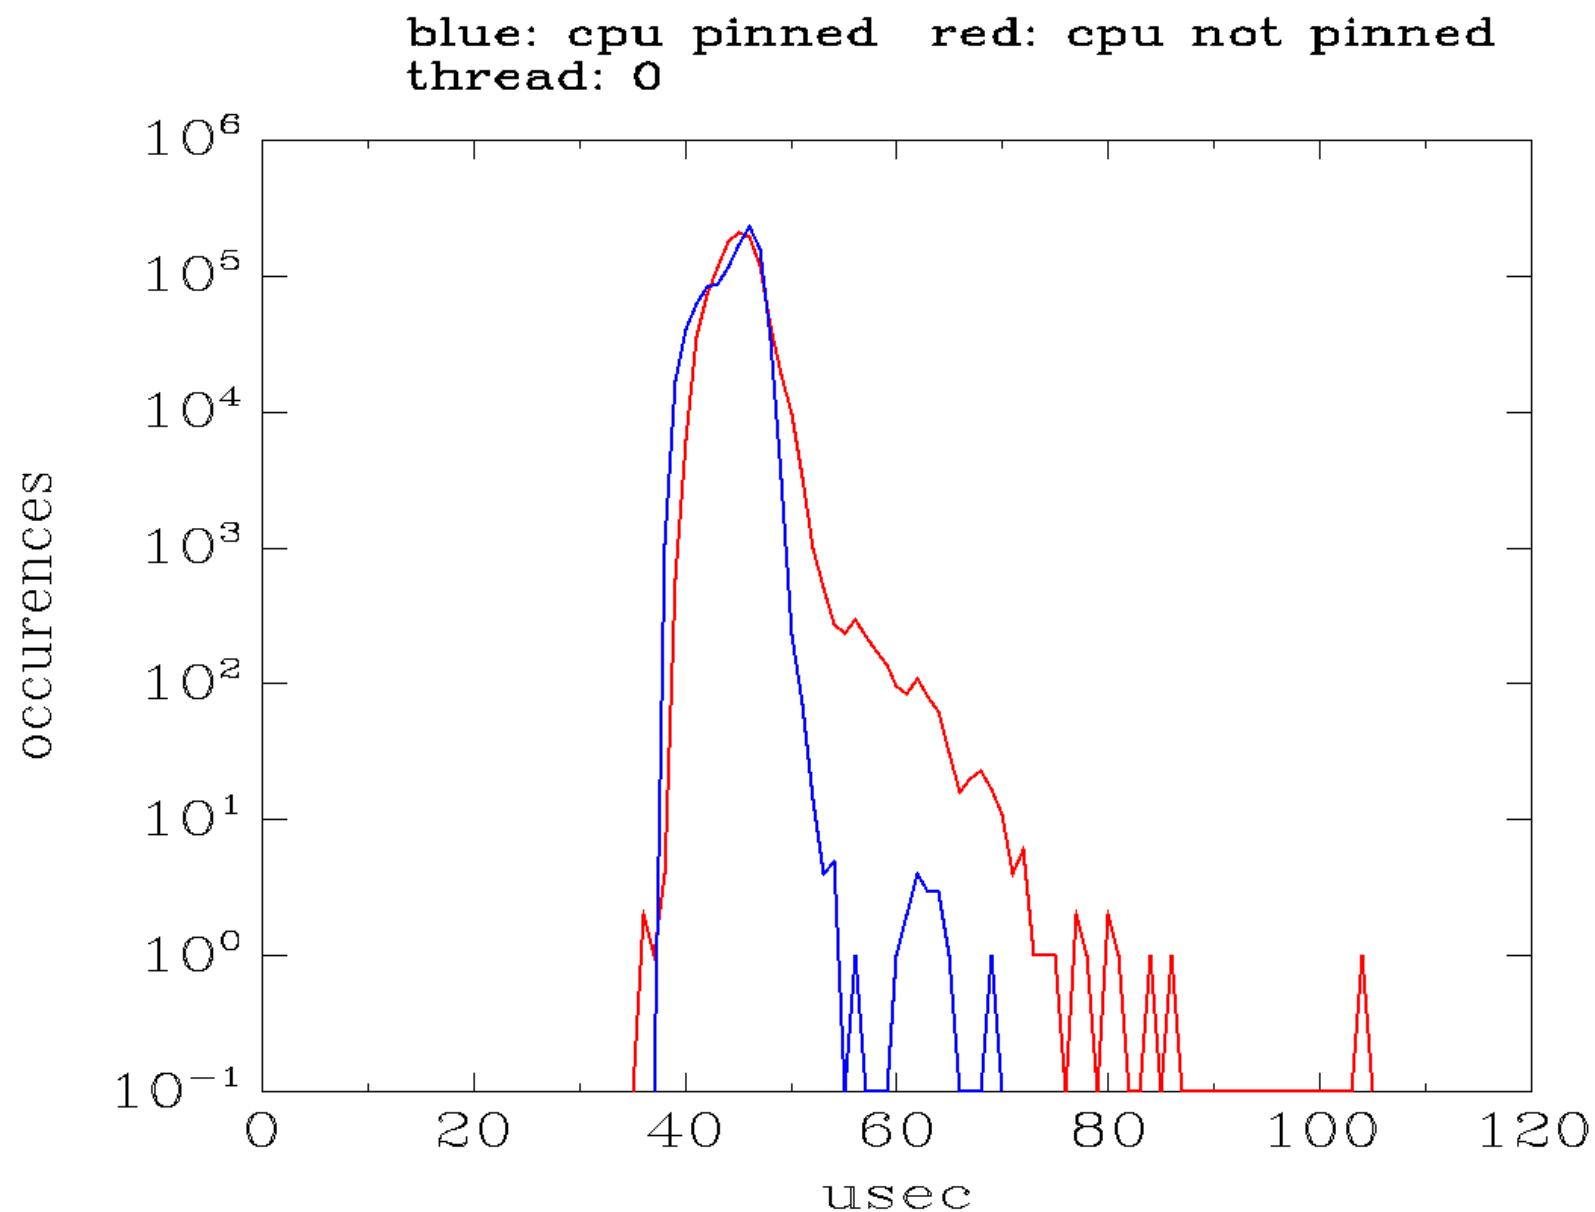
PID	PPID	S	RTPRIO	CLS	CMD
3	2	S	1	FF	[ksoftirqd/0]
6	2	S	70	FF	[posixcpu/0]
7	2	S	99	FF	[migration/0]
8	2	S	70	FF	[posixcpu/1]
9	2	S	99	FF	[migration/1]
11	2	S	1	FF	[ksoftirqd/1]
353	2	S	50	FF	[irq/41-eth%d]
374	2	S	50	FF	[irq/46-mmci-p1]
375	2	S	50	FF	[irq/47-mmci-p1]
394	2	S	50	FF	[irq/36-uart-p10]

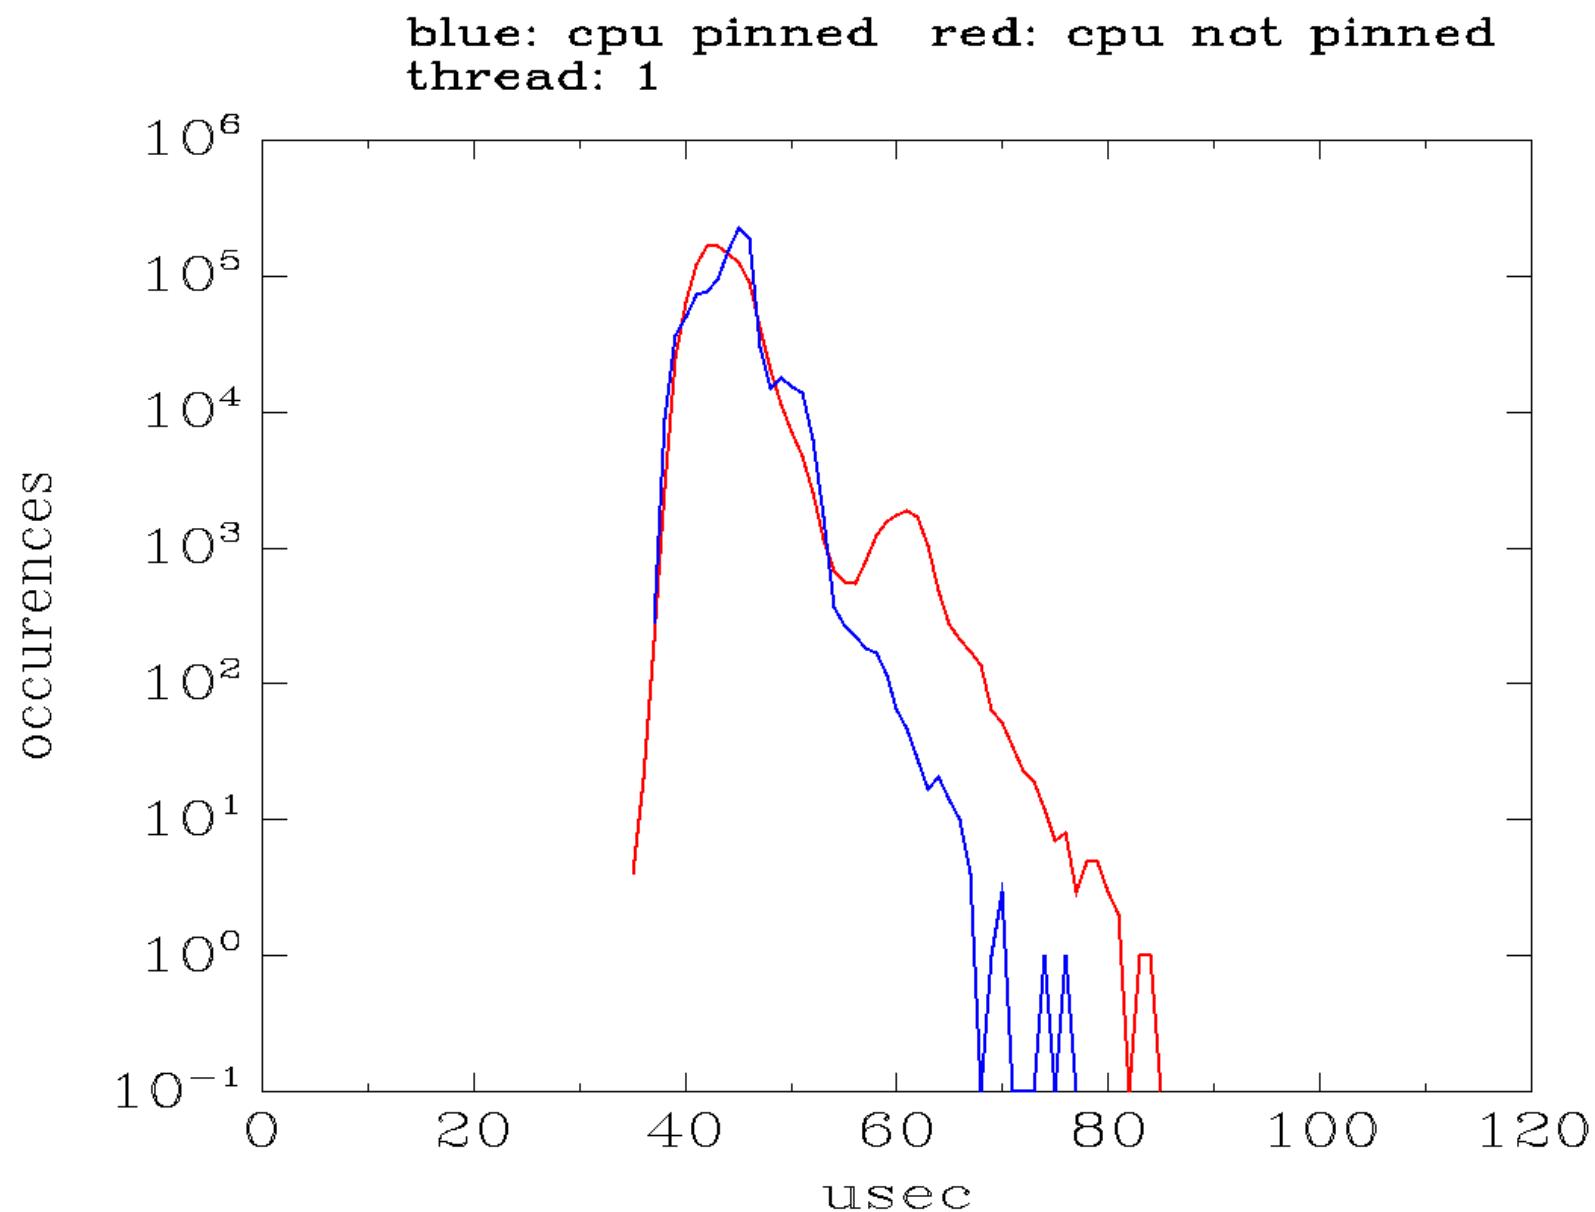
-l100000

T:0	Min:	128	Avg:	189	Max:	2699		<u>live update</u>
T:0	Min:	125	Avg:	140	Max:	472	-q	<u>no live update</u>
T:0	Min:	262	Avg:	281	Max:	337	-p80	<u>SCHED_FIFO 80</u>
T:0	Min:	88	Avg:	96	Max:	200	-n	<u>clock_nanosleep</u>
T:0	Min:	246	Avg:	320	Max:	496	-q -p80 -a -t	<u>pinned</u>
T:1	Min:	253	Avg:	315	Max:	509		
T:0	Min:	35	Avg:	43	Max:	68	-q -p80 -n	<u>SCHED_FIFO, c_n</u>
T:0	Min:	34	Avg:	44	Max:	71	-q -p80 -a -n	<u>pinned</u>
T:0	Min:	38	Avg:	43	Max:	119	-q -p80 -a -n -m	<u>mem locked</u>
T:0	Min:	36	Avg:	43	Max:	65	-q -p80 -t -n	<u>not pinned</u>
T:1	Min:	37	Avg:	45	Max:	78		
T:0	Min:	36	Avg:	44	Max:	91	-q -p80 -a -t -n	<u>pinned</u>
T:1	Min:	37	Avg:	45	Max:	111		
T:0	Min:	34	Avg:	44	Max:	94	-q -p80 -S	<u>=> -a -t -n</u>
T:1	Min:	34	Avg:	43	Max:	104		

-l1000000

T:0	Min:	123	Avg:	184	Max:	3814		<u>live update</u>
T:0	Min:	125	Avg:	150	Max:	860	-q	<u>no live update</u>
T:0	Min:	257	Avg:	281	Max:	371	-q -p80	<u>SCHED_FIFO 80</u>
T:0	Min:	84	Avg:	94	Max:	319	-q -n	<u>clock_nanosleep</u>
T:0	Min:	247	Avg:	314	Max:	682	-q -p80 -a -t	<u>pinned</u>
T:1	Min:	228	Avg:	321	Max:	506		
T:0	Min:	38	Avg:	44	Max:	72	-q -p80 -n	<u>SCHED_FIFO, c_n</u>
T:0	Min:	33	Avg:	42	Max:	95	-q -p80 -a -n	<u>pinned</u>
T:0	Min:	36	Avg:	42	Max:	144	-q -p80 -a -n -m	<u>mem locked</u>
T:0	Min:	36	Avg:	44	Max:	84	-q -p80 -t -n	<u>not pinned</u>
T:1	Min:	37	Avg:	45	Max:	94		
T:0	Min:	36	Avg:	43	Max:	87	-q -p80 -a -t -n	<u>pinned</u>
T:1	Min:	36	Avg:	43	Max:	91		
T:0	Min:	36	Avg:	43	Max:	141	-q -p80 -S	<u>=> -a -t -n</u>
T:1	Min:	34	Avg:	42	Max:	88		





Simple Demo -- SCED_NORMAL

- single thread
- `clock_nanosleep()`, one thread per cpu, pinned
- `clock_nanosleep()`, one thread per cpu
- `clock_nanosleep()`, one thread per cpu,
memory locked
- `clock_nanosleep()`, one thread per cpu,
memory locked, non-interactive

What Are Normal Results?

What should I expect the data to look like for my system?

Examples of Maximum Latency

[https://rt.wiki.kernel.org/index.php/CONFIG_PREEMPT_RT_Patch
#Platforms_Tested_and_in_Use_with_CONFIG_PREEMPT_RT](https://rt.wiki.kernel.org/index.php/CONFIG_PREEMPT_RT_Patch#Platforms_Tested_and_in_Use_with_CONFIG_PREEMPT_RT)

Platforms Tested and in Use with CONFIG_PREEMPT_RT

Comments sometimes include avg and max latency
table is usually stale

linux-rt-users email list archives

<http://vger.kernel.org/vger-lists.html#linux-rt-users>

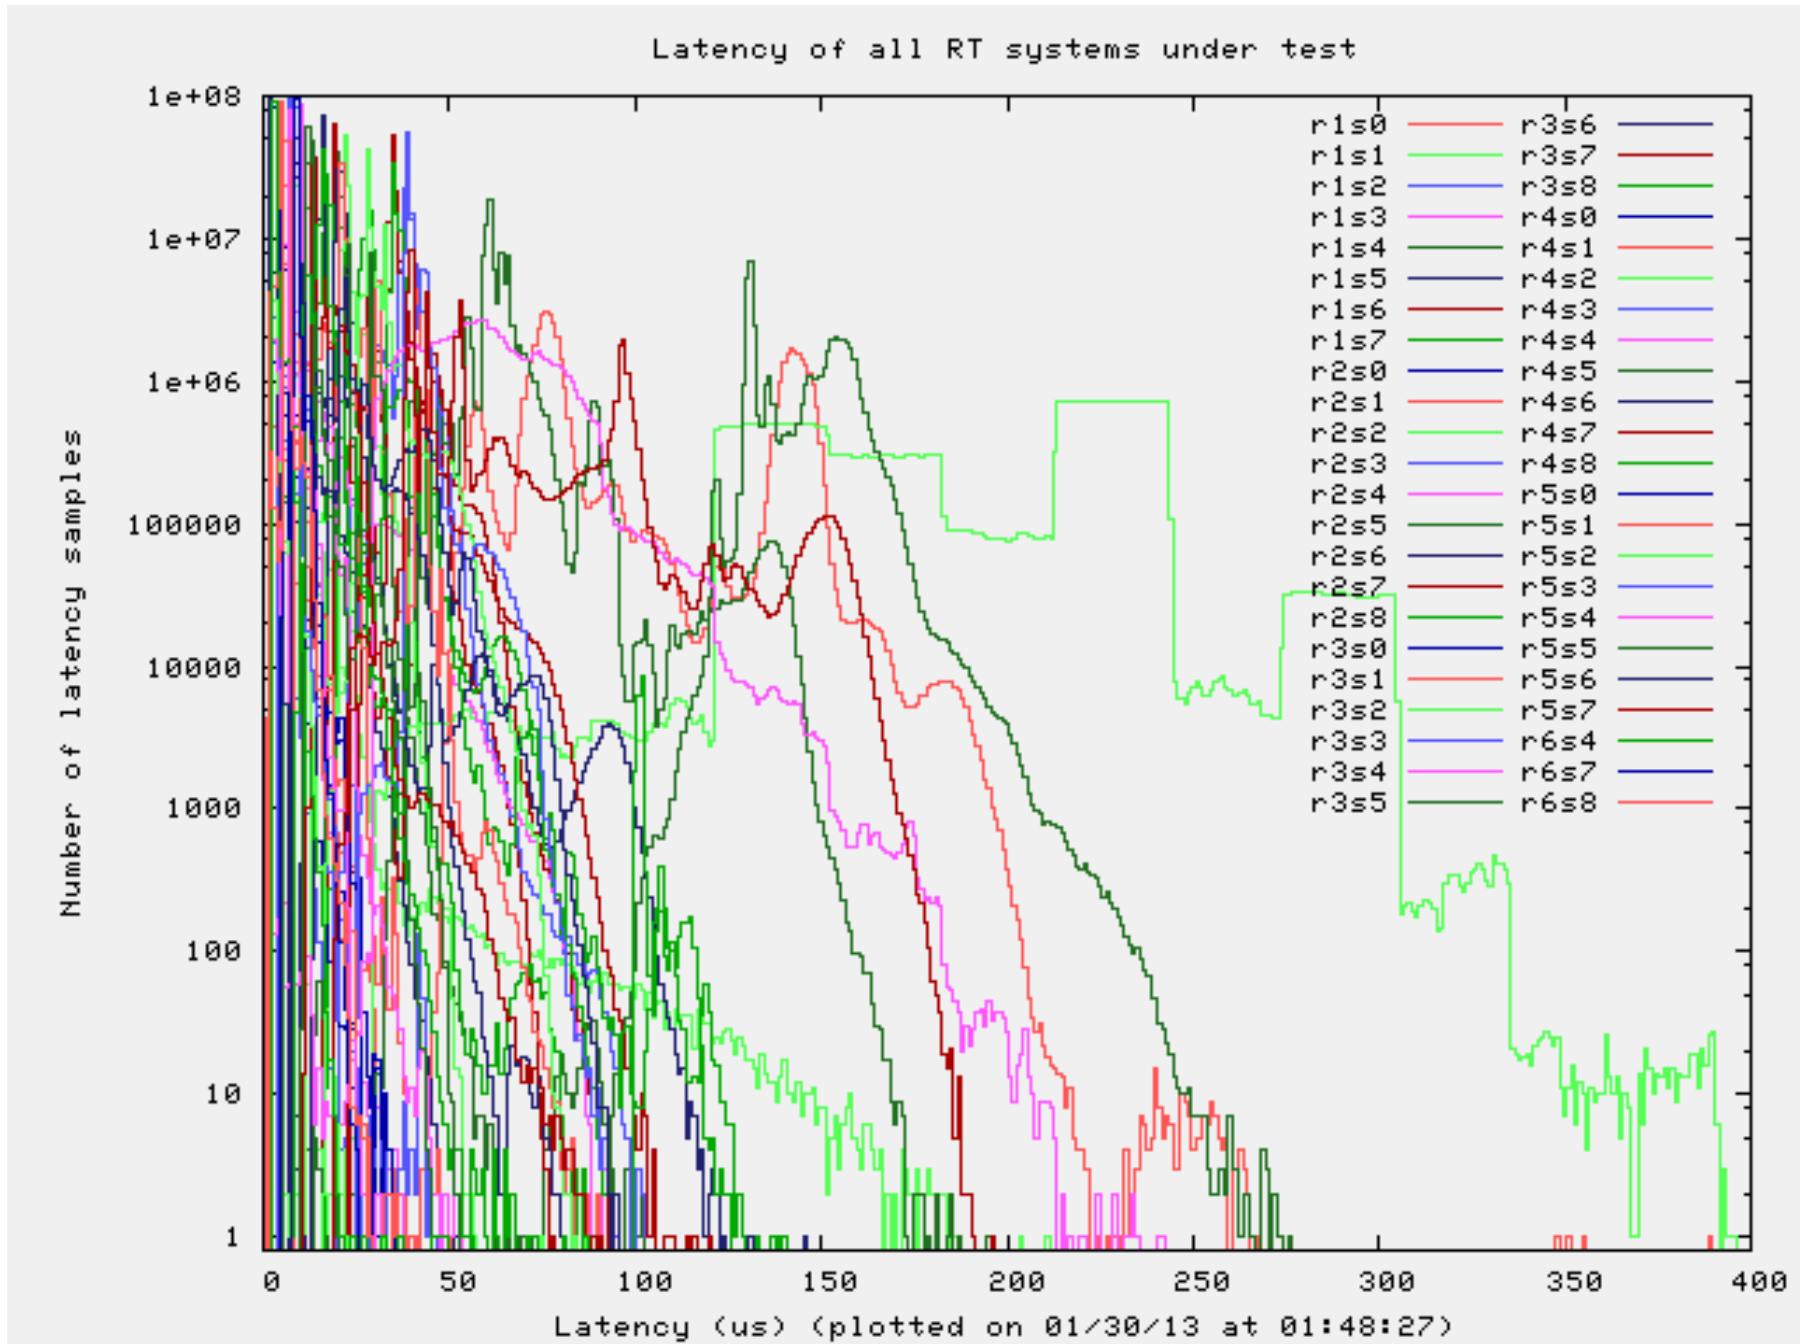
Graphs of Maximum Latency

OSADL.org

Graphs for a wide variety of machines

List of test systems:

<https://www.osadl.org/Individual-system-data.qa-farm-data.0.html>



Full URL of previous graph

<https://www.osadl.org/Combined-latency-plot-of-all-RT-systems.qa-latencyplot-allrt.0.html?latencies=&showno=>

Typical command:

```
cyclictest -l100000000 -m -Sp99 -i200 -h400 -q
```

OSADL Realtime QA Farm:

<https://www.osadl.org/QA-Farm-Realtime.qa-farm-about.0.html>

OSADL Latency plots:

<https://www.osadl.org/Latency-plots.latency-plots.0.html>

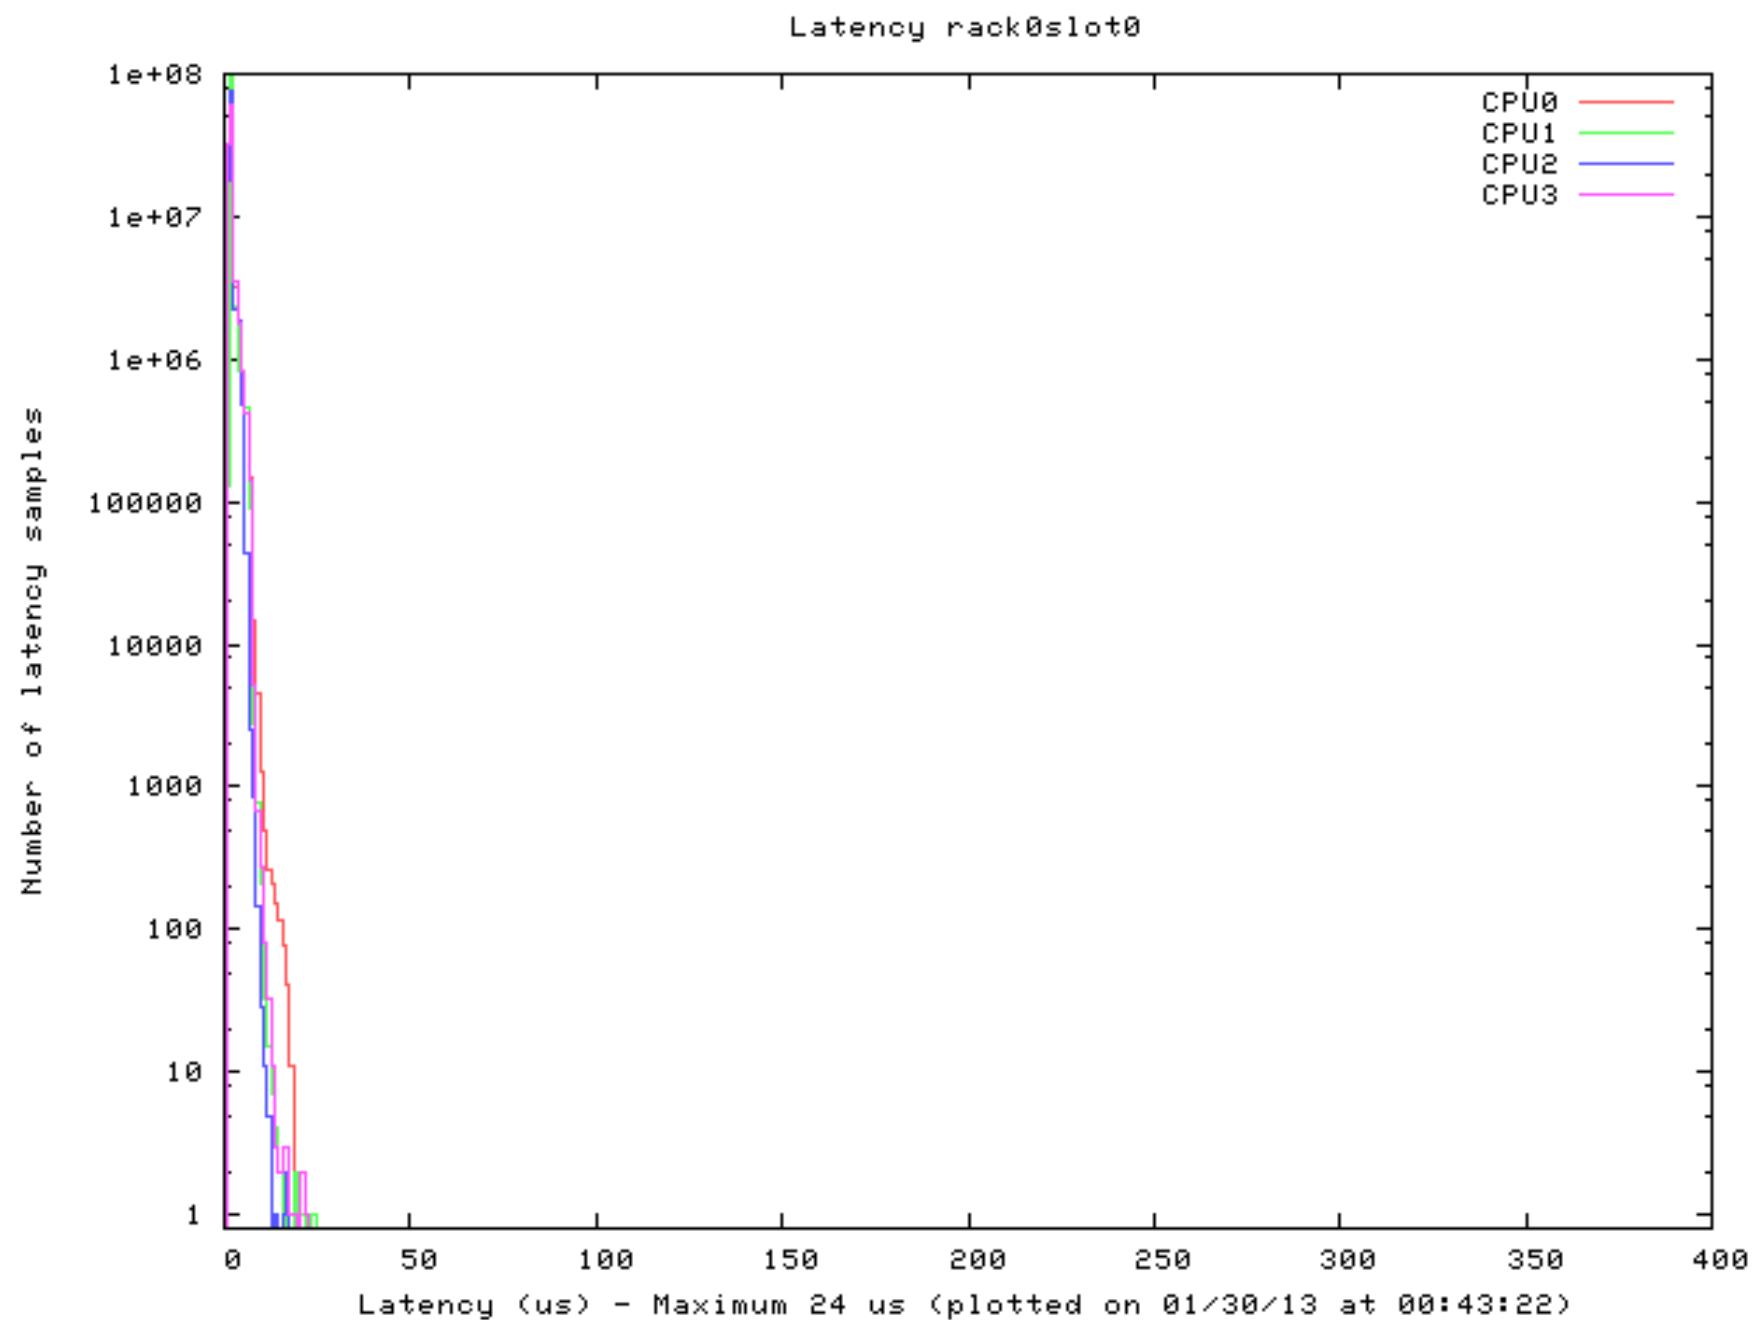
Additional OSADL Data

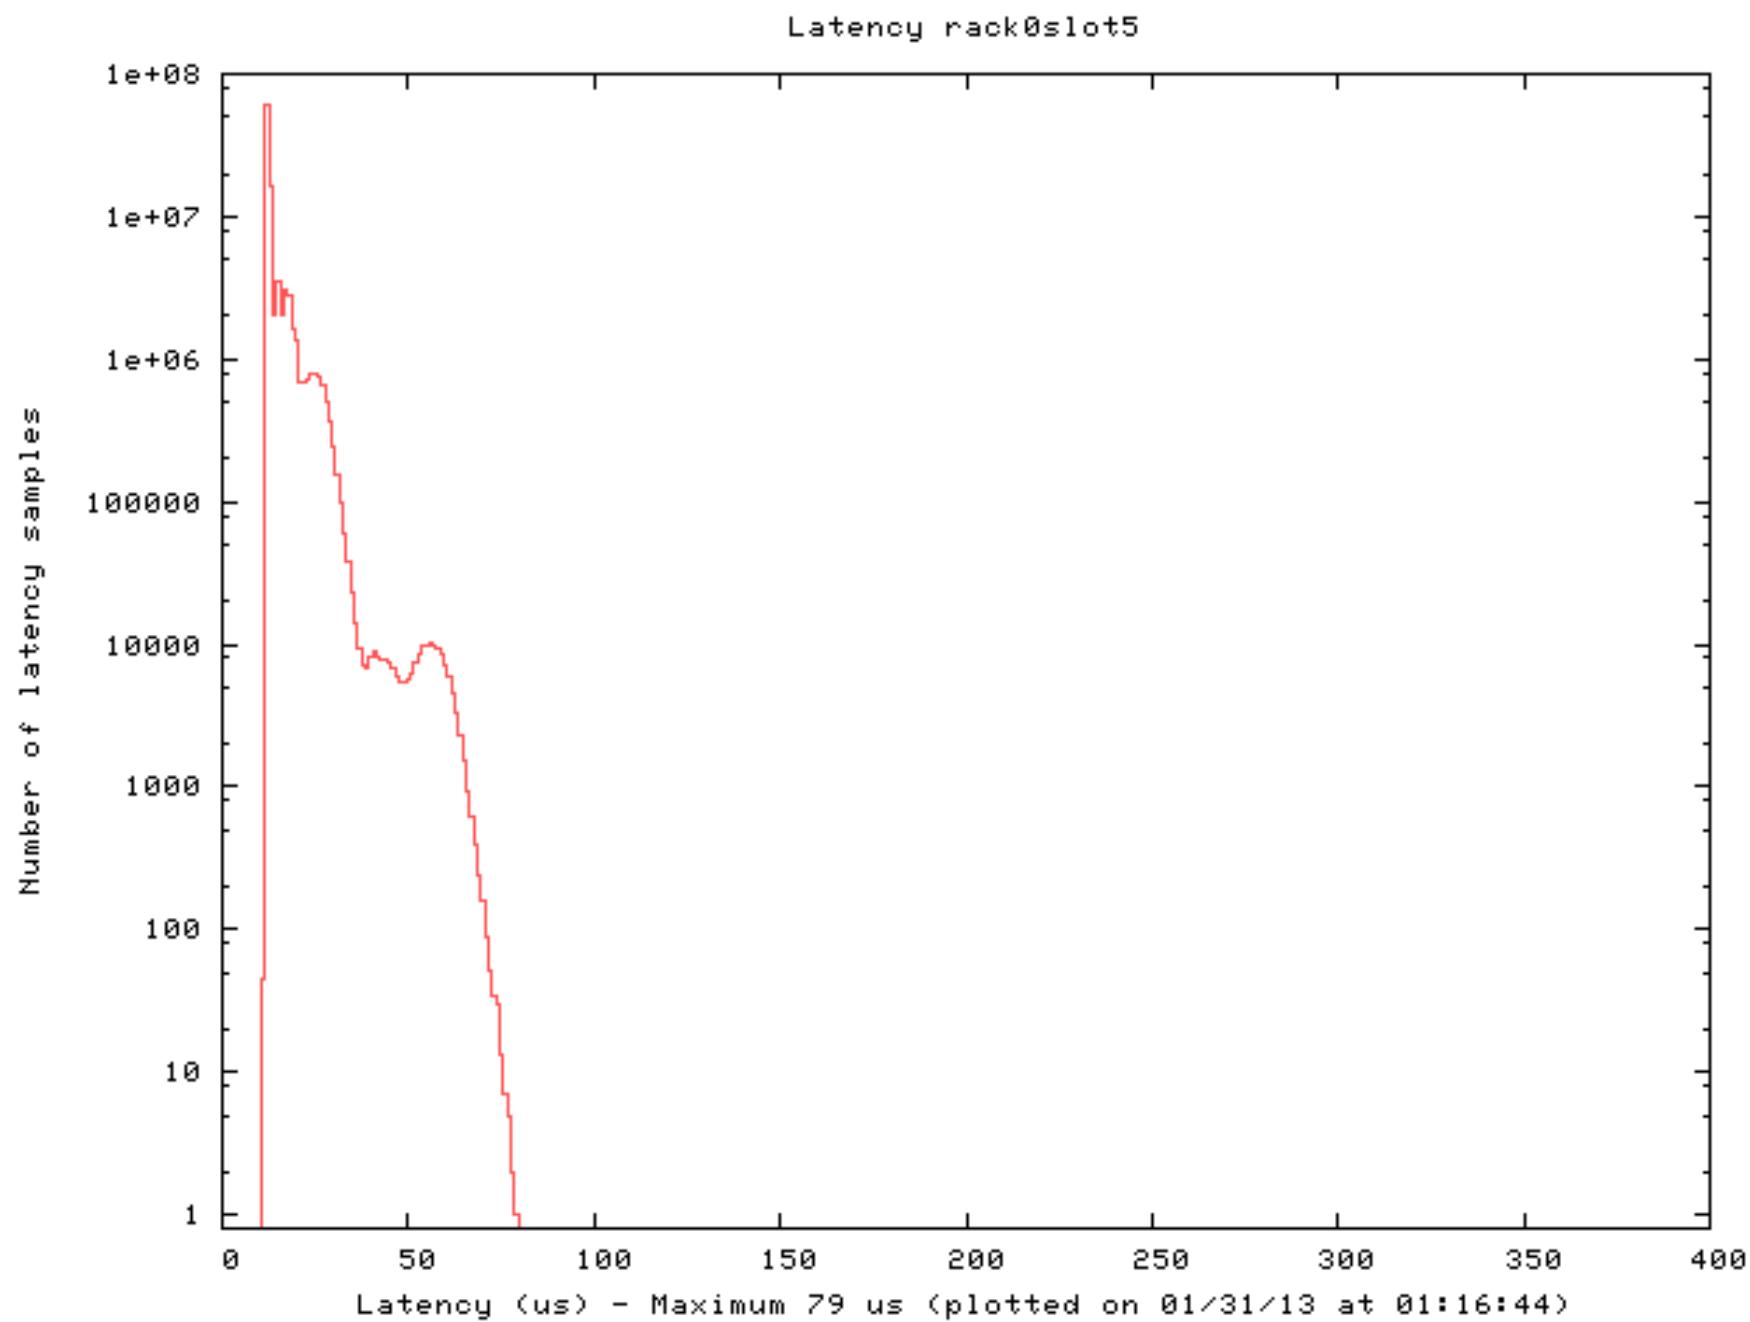
OSADL members have access to additional data, such as

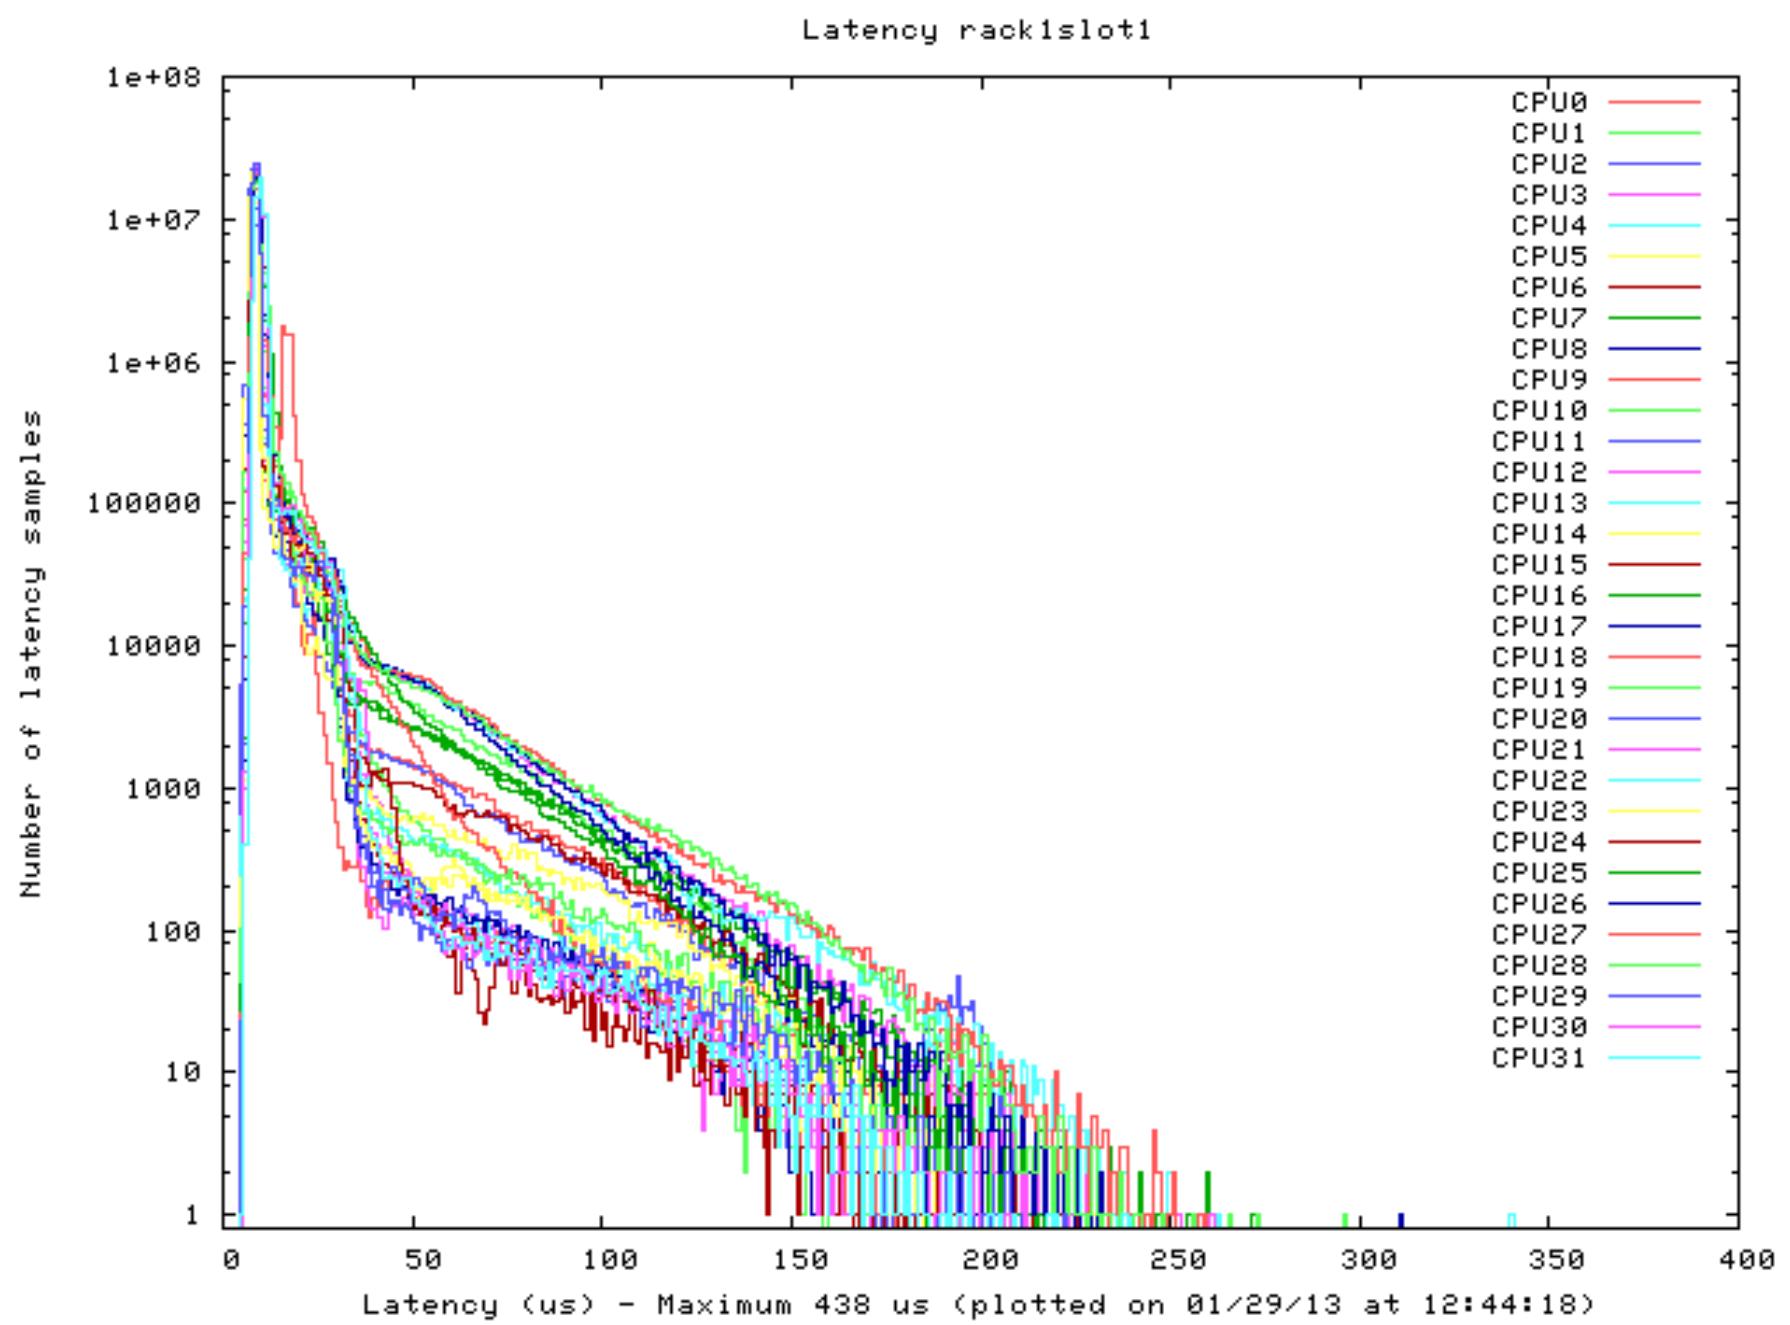
- the data used to create the graphs
- the latency graphs extended in a third dimension, showing all test runs

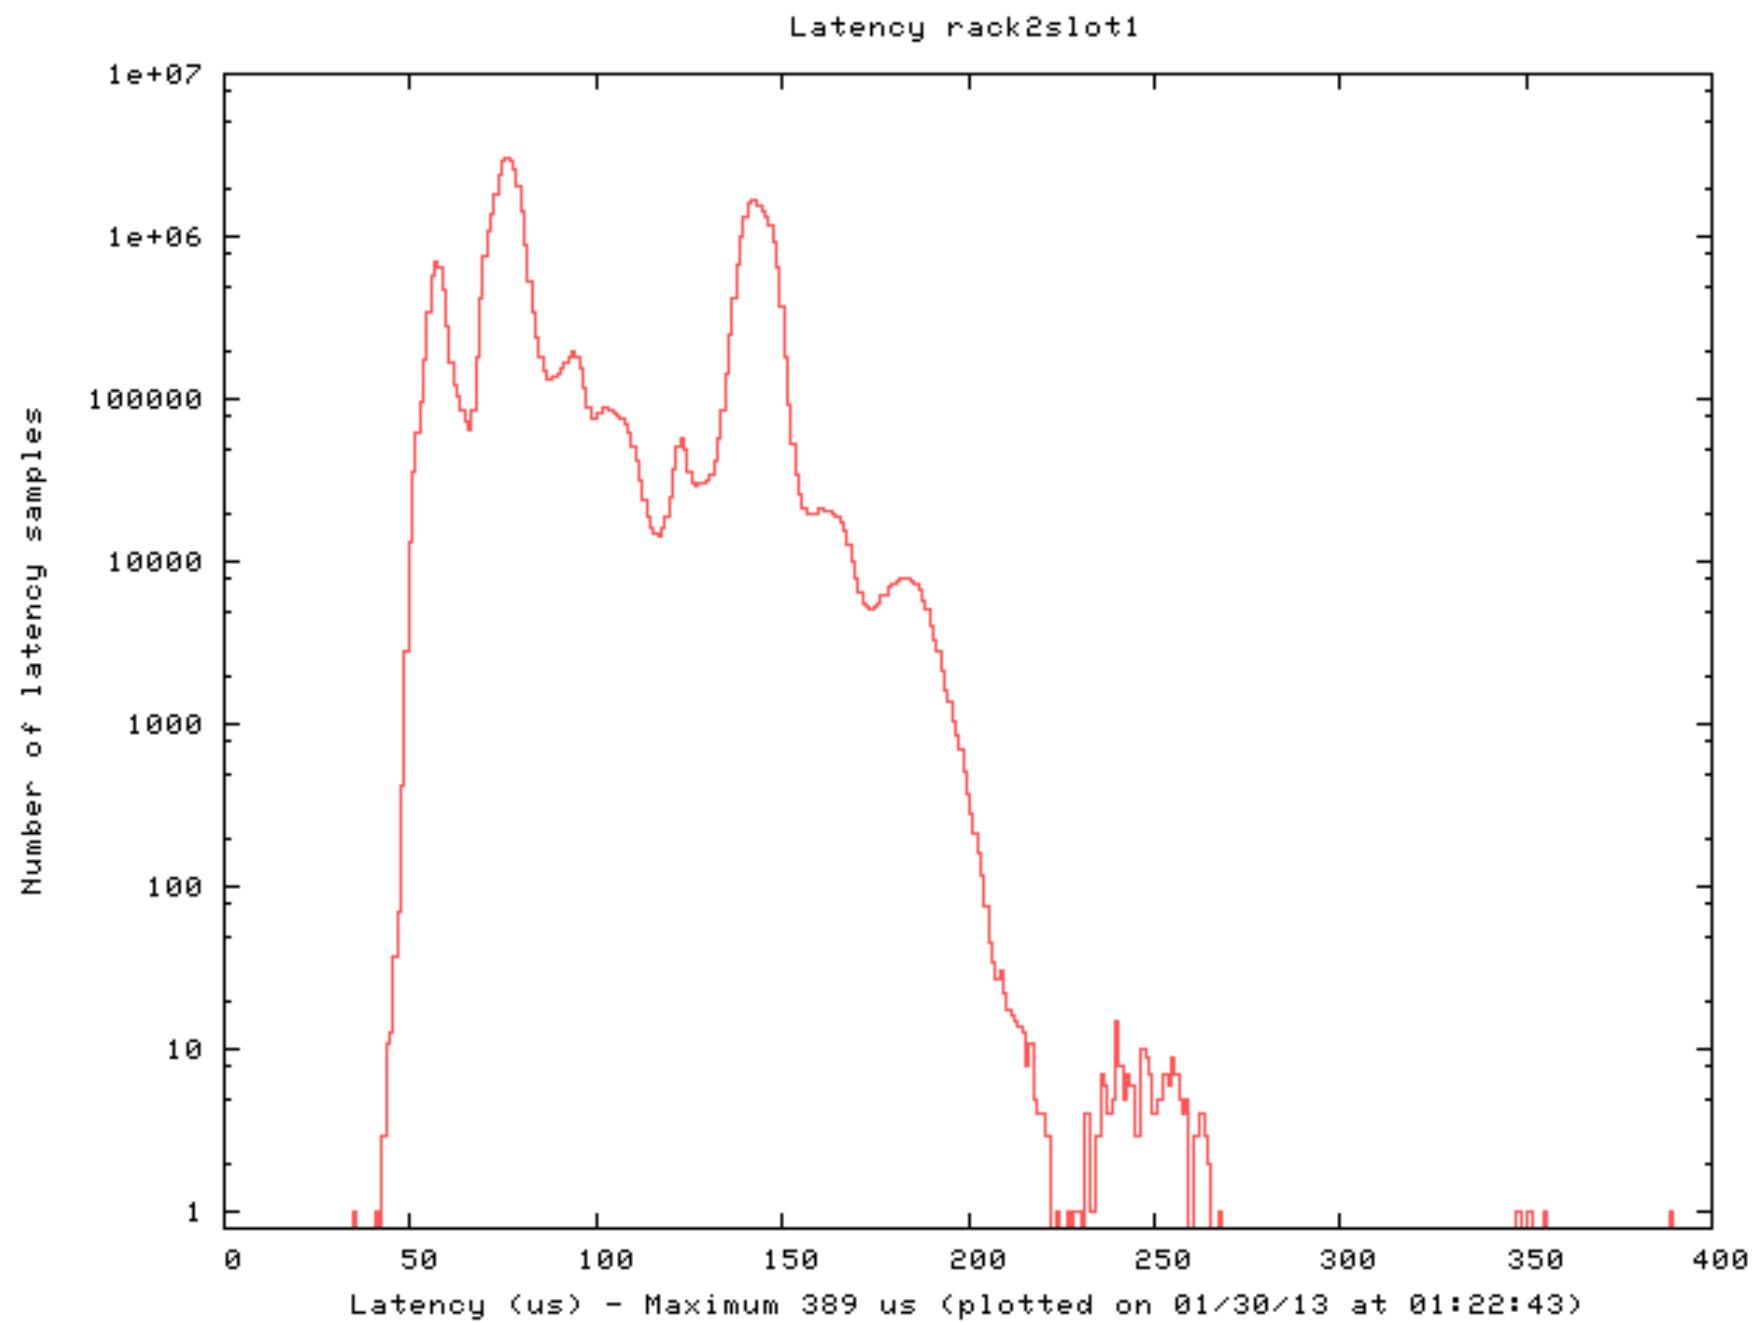
Some Random Individual Systems

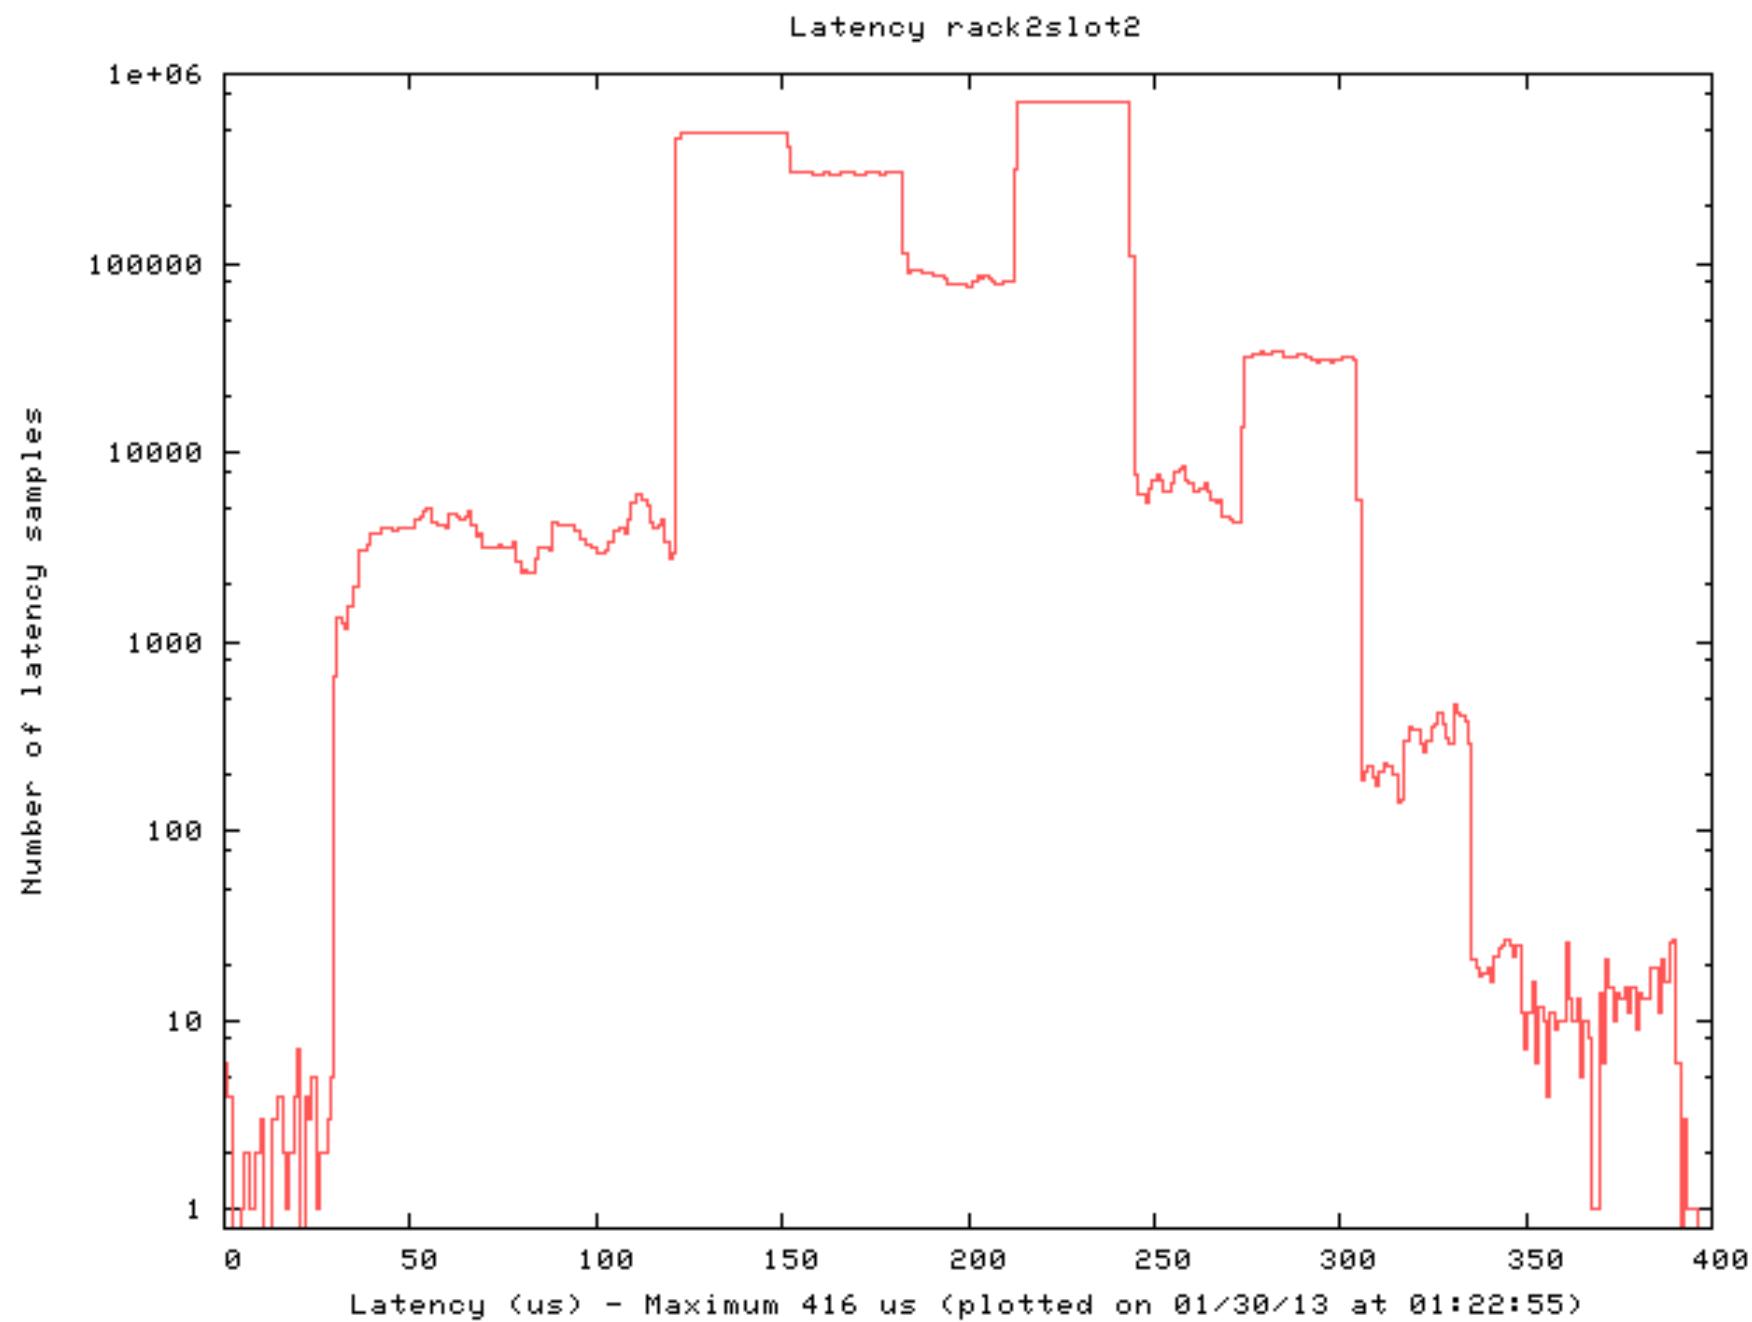
Picked from the OSADL spaghetti graph

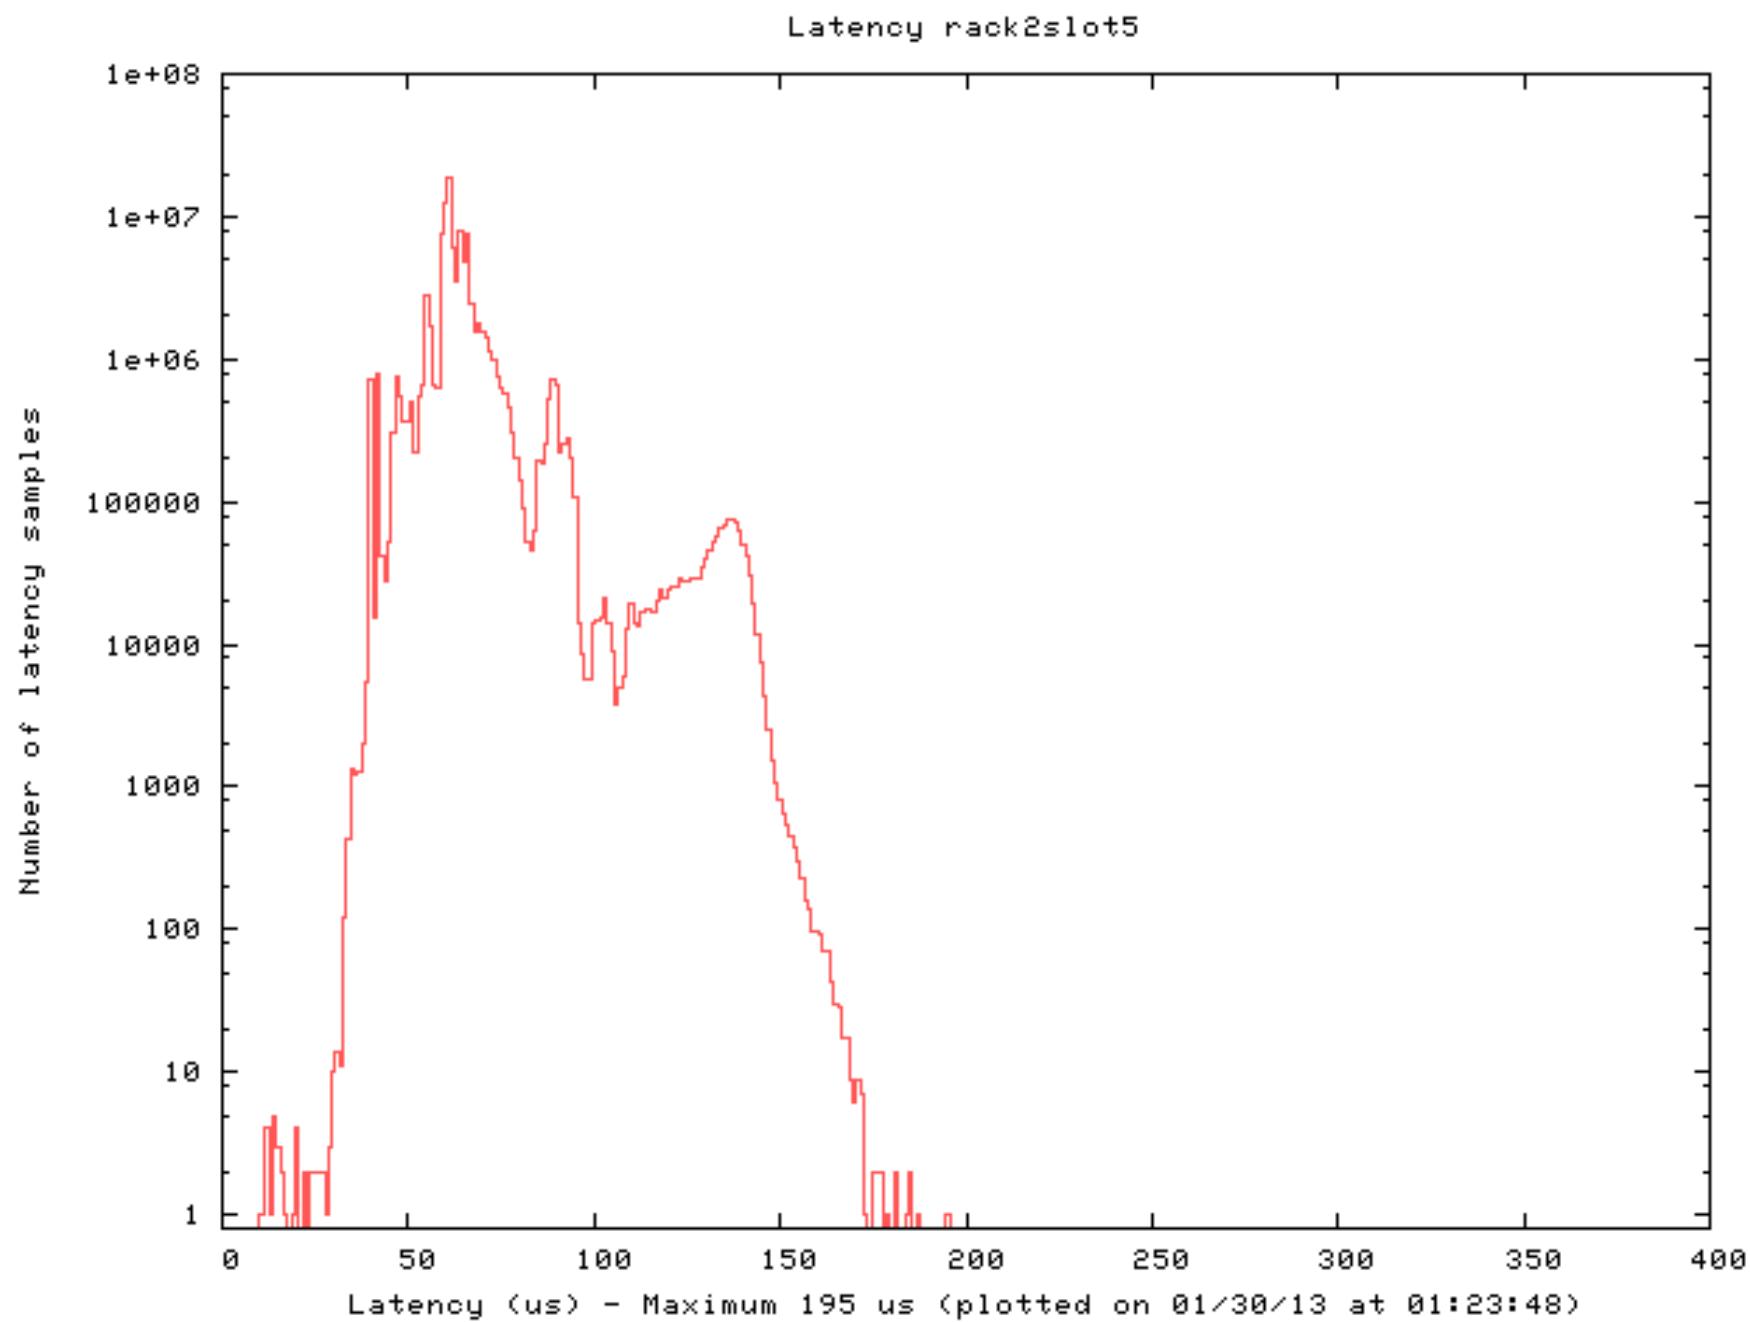


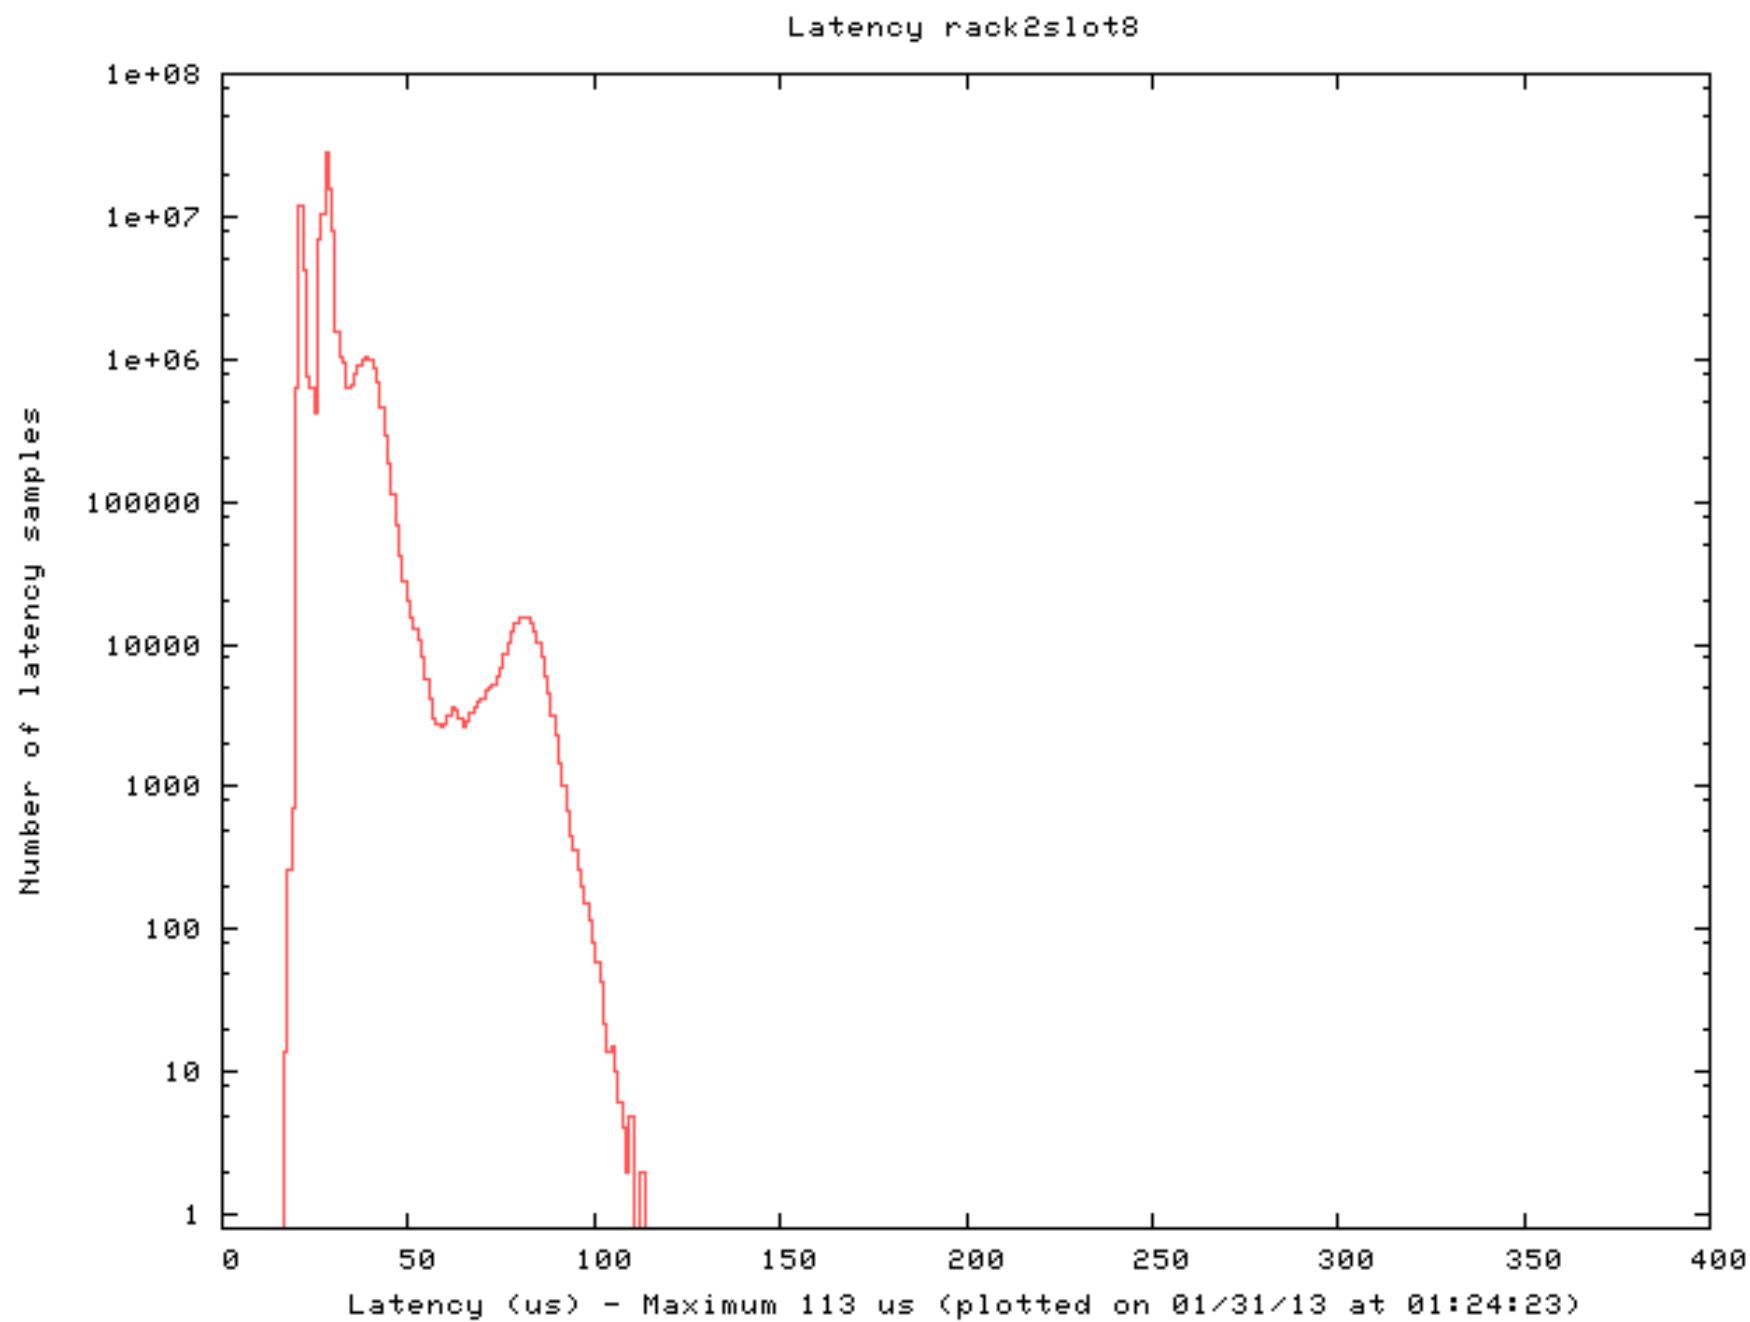


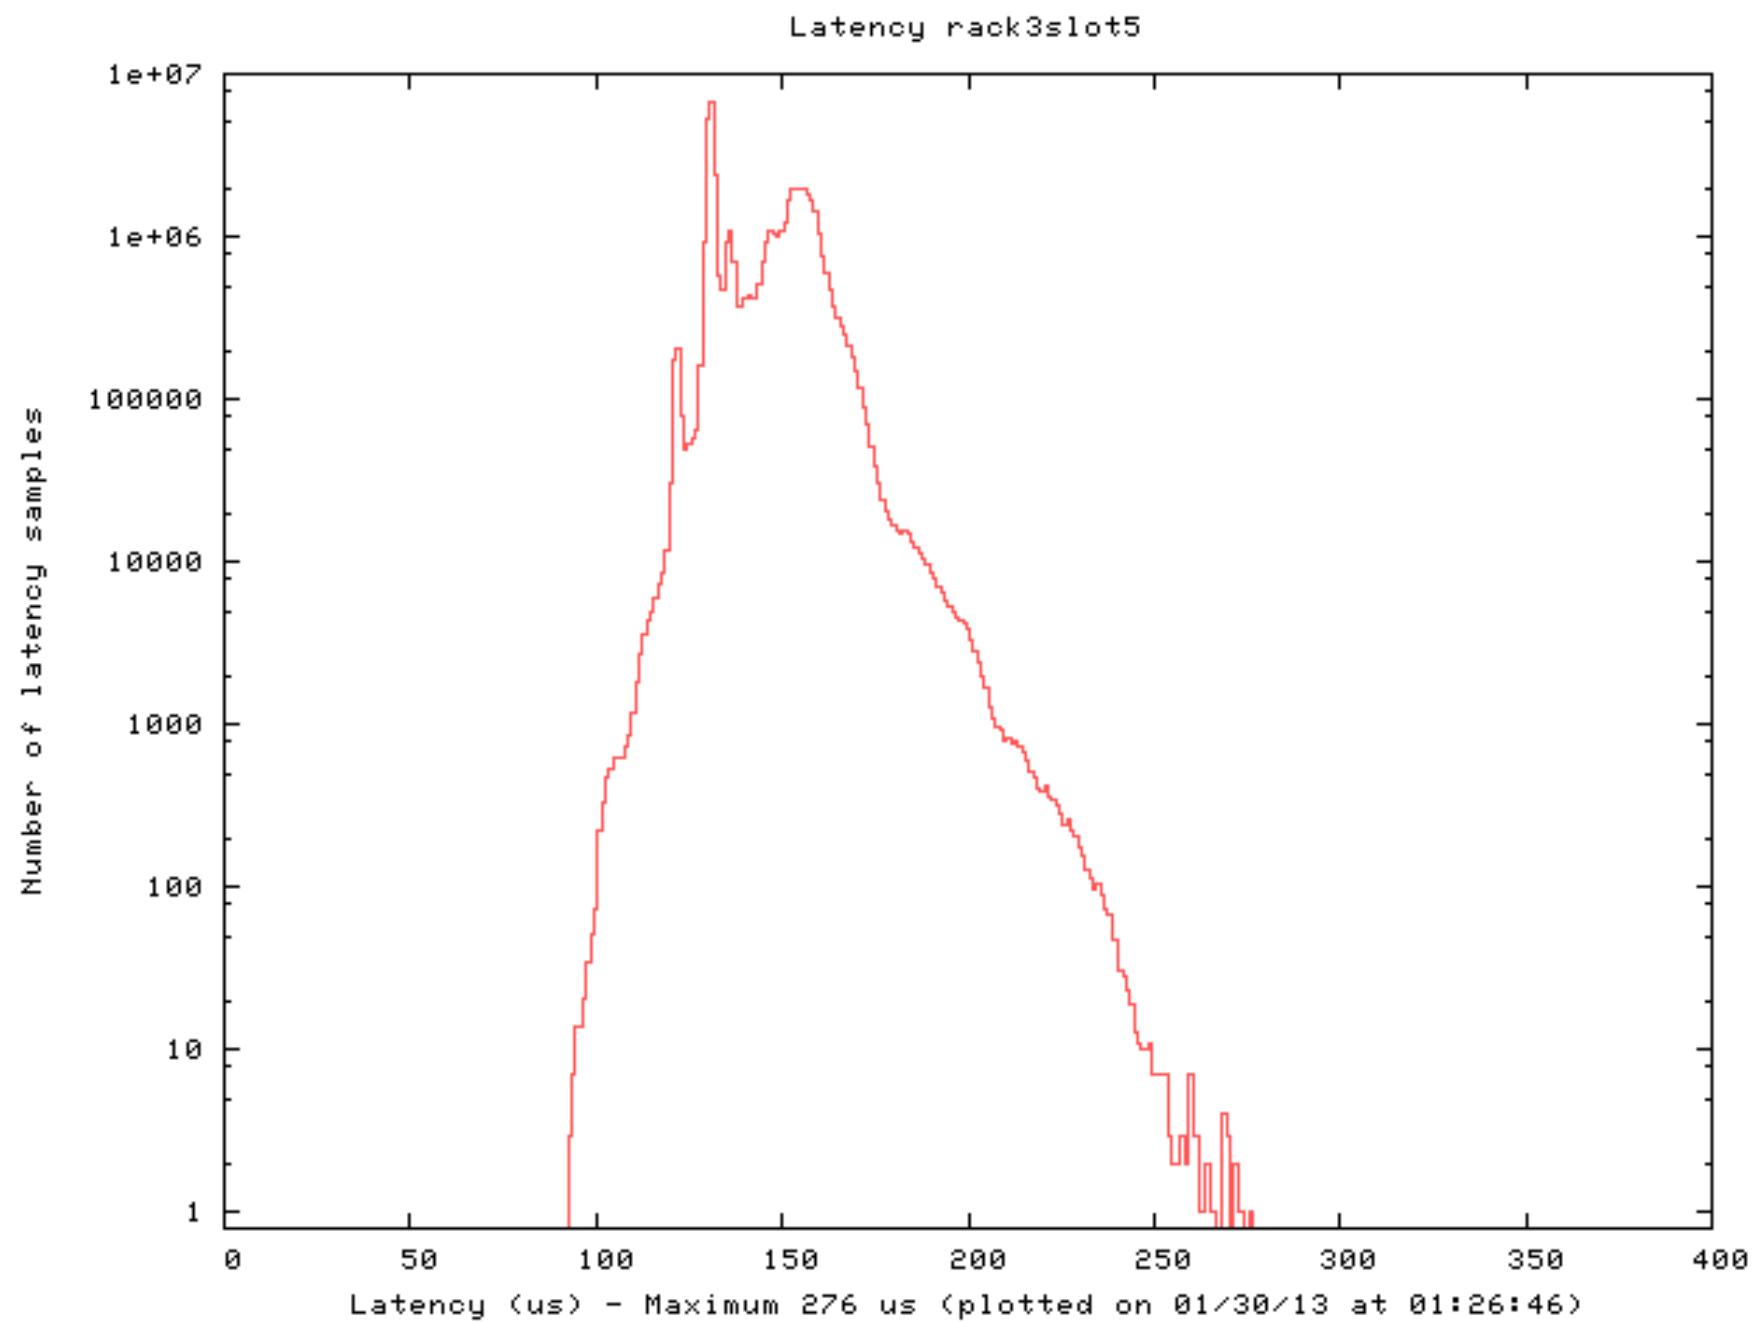


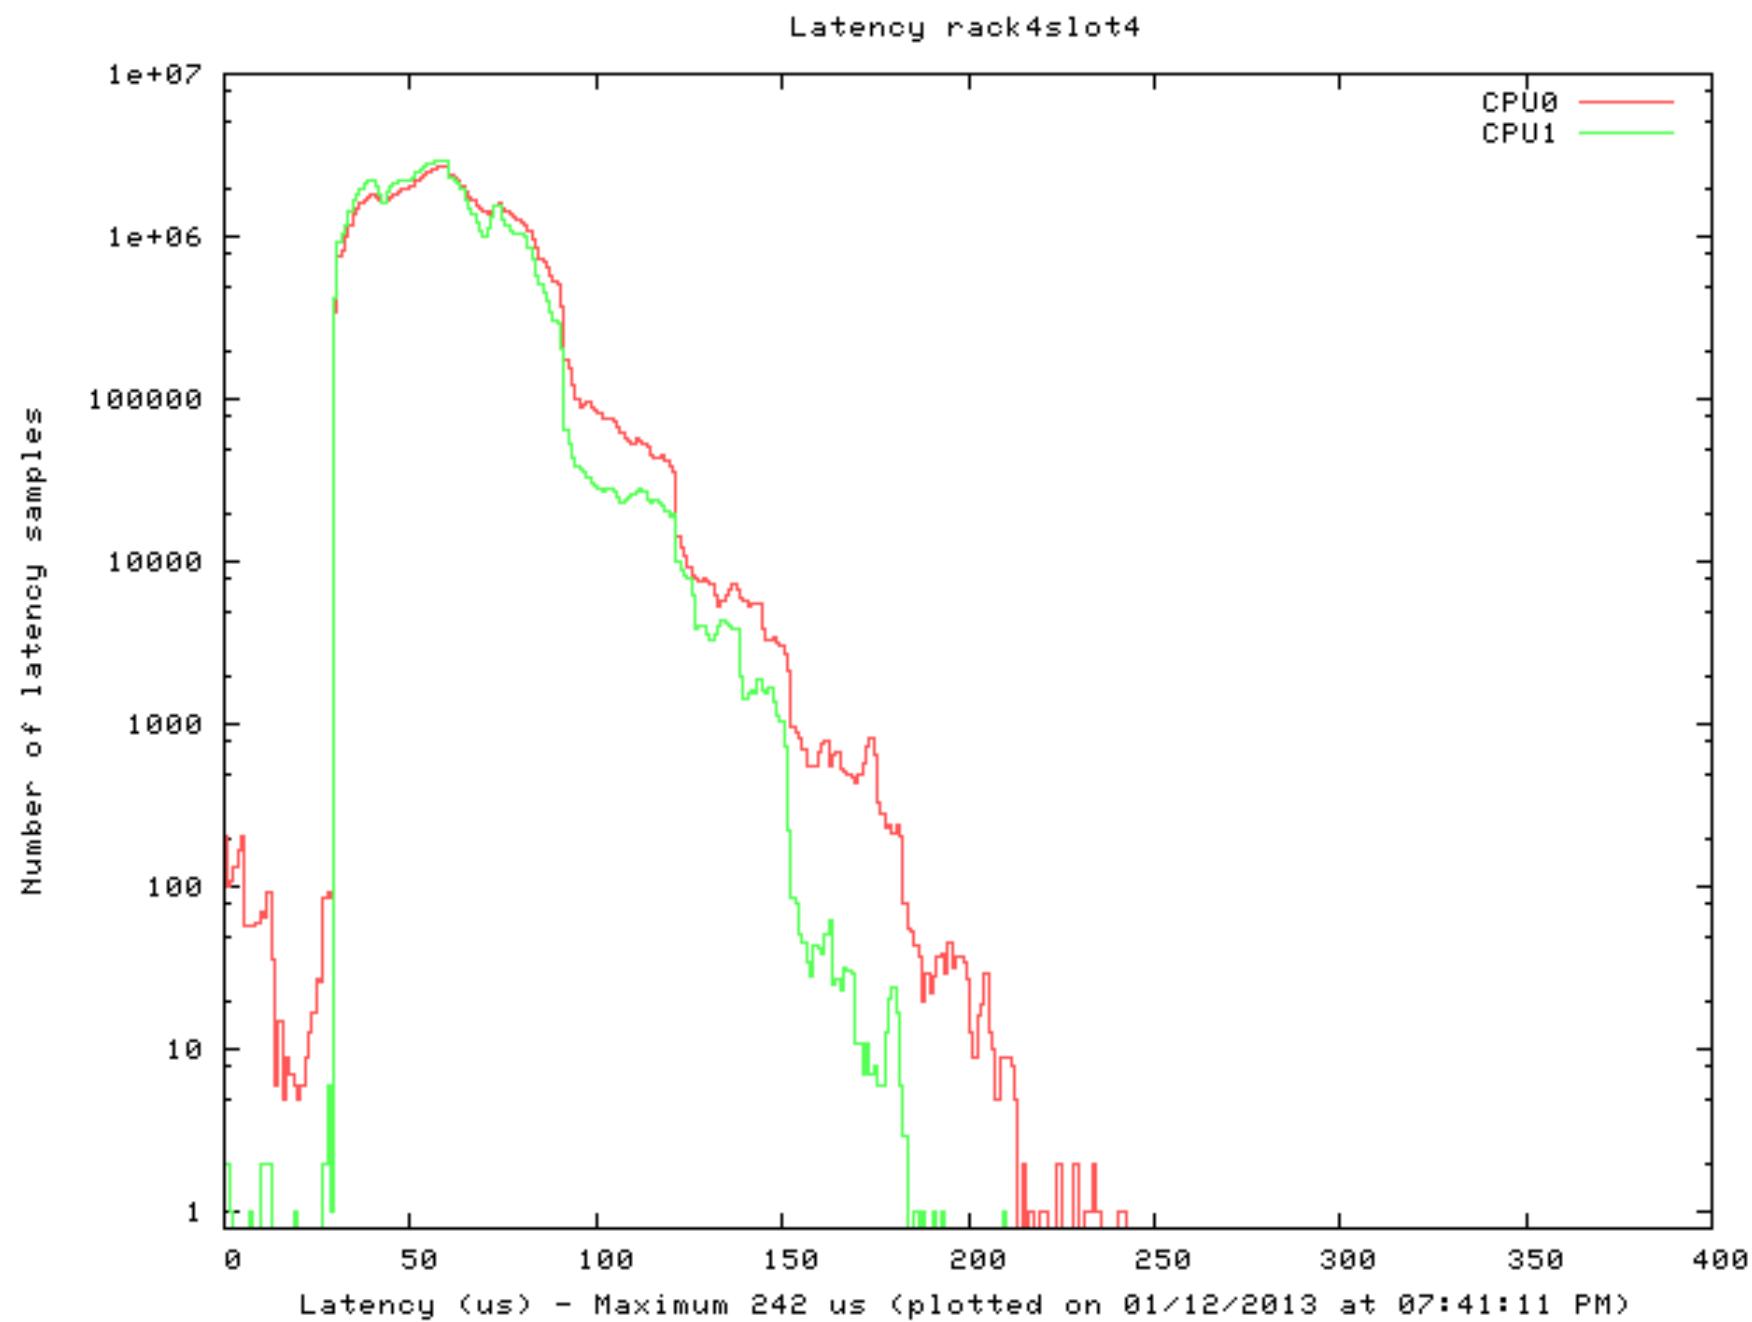


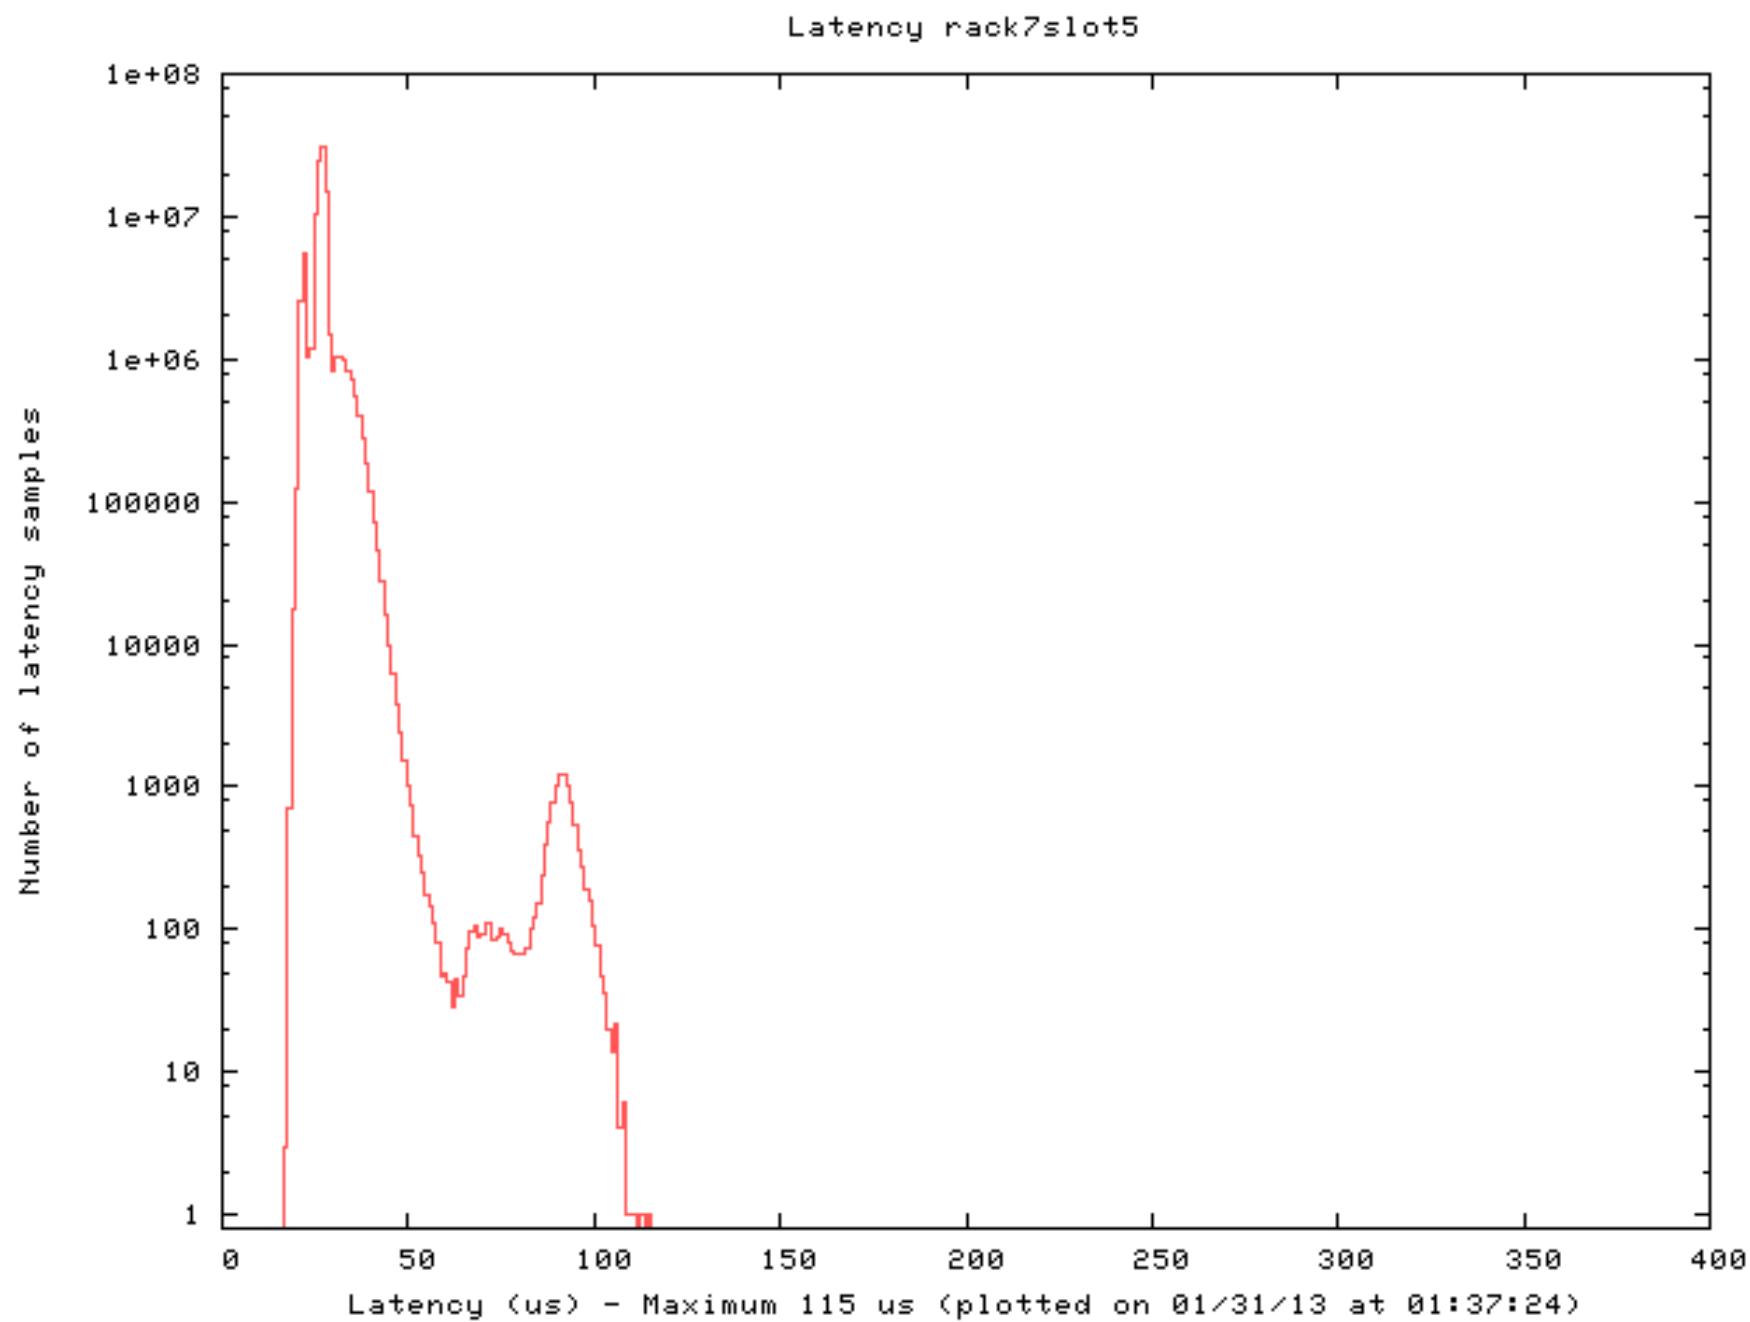




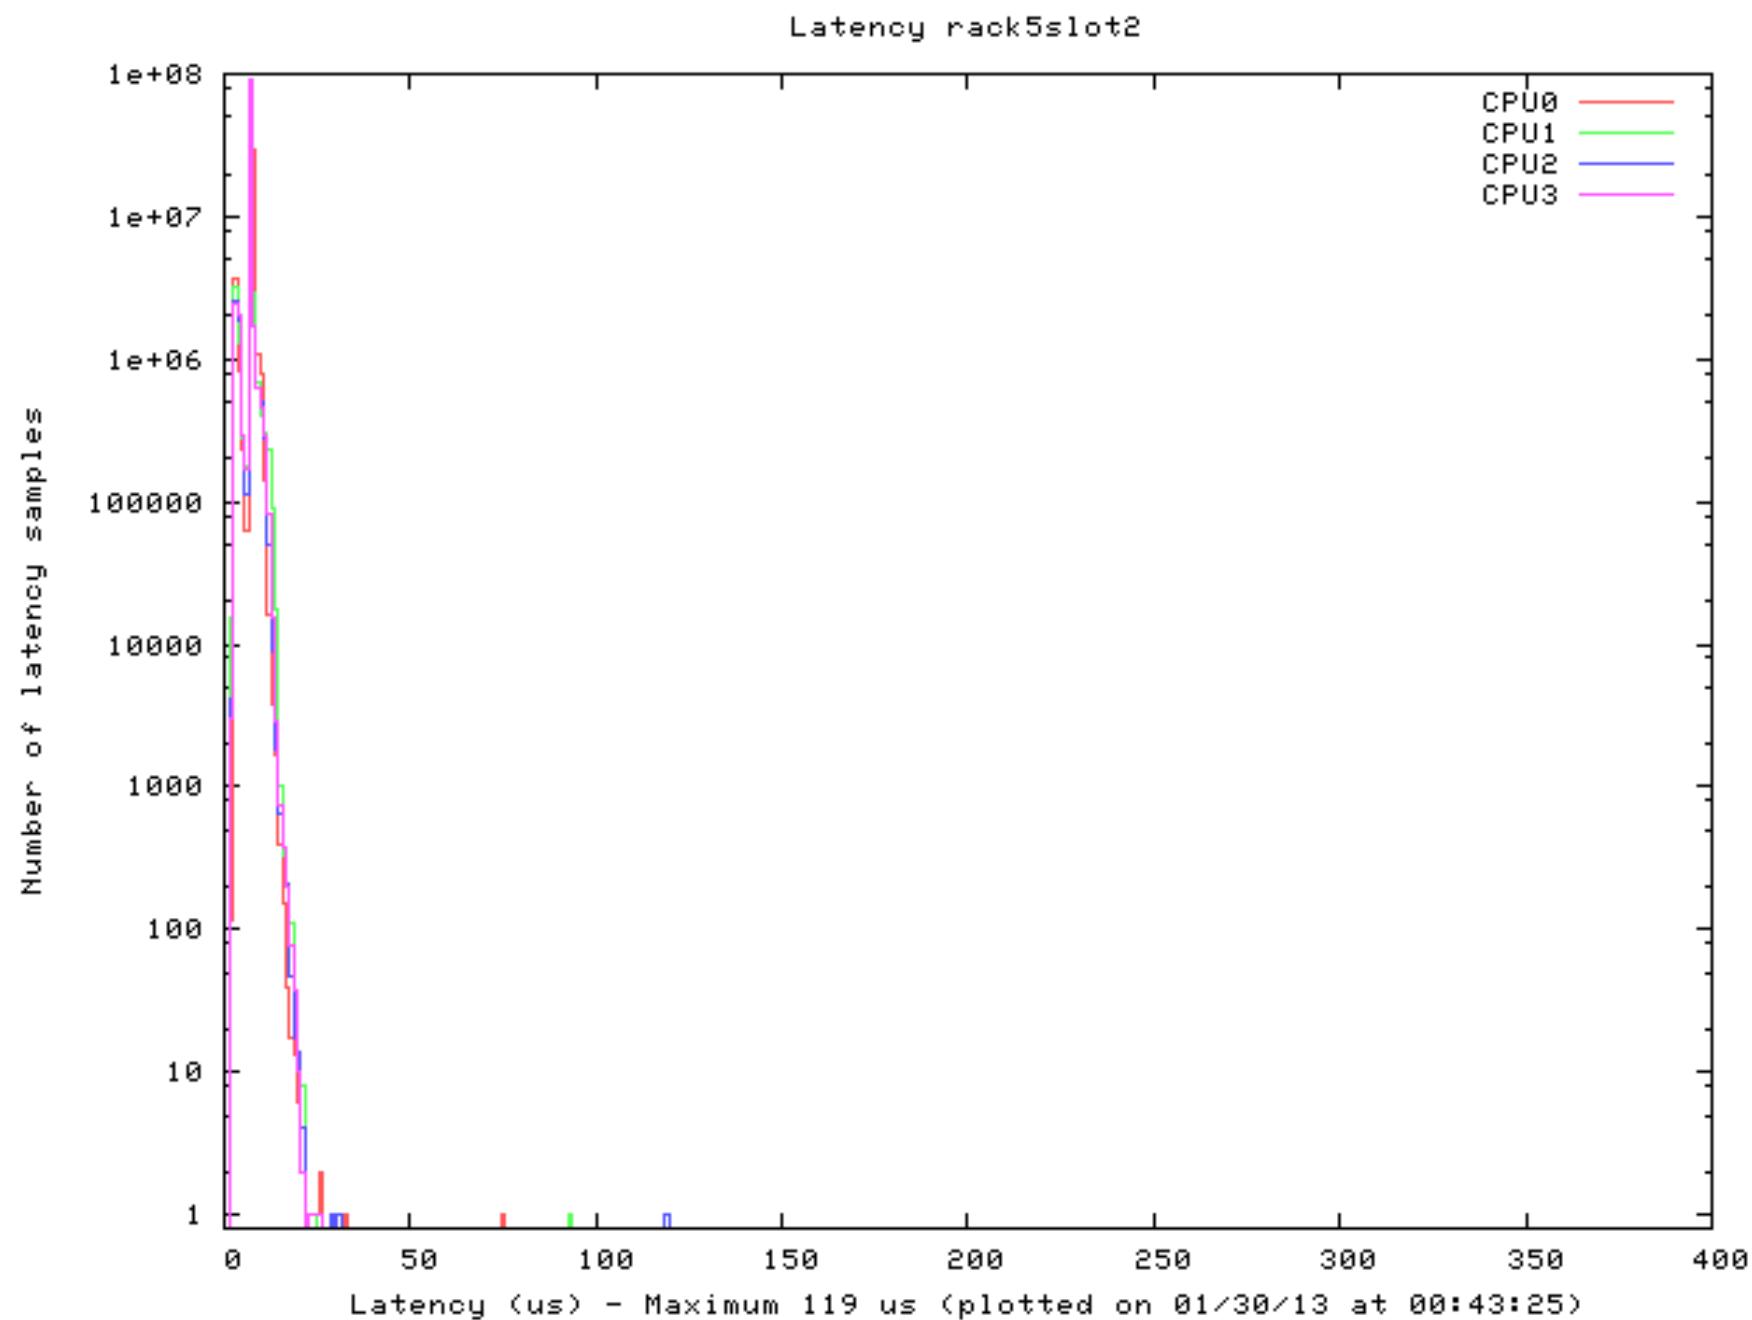




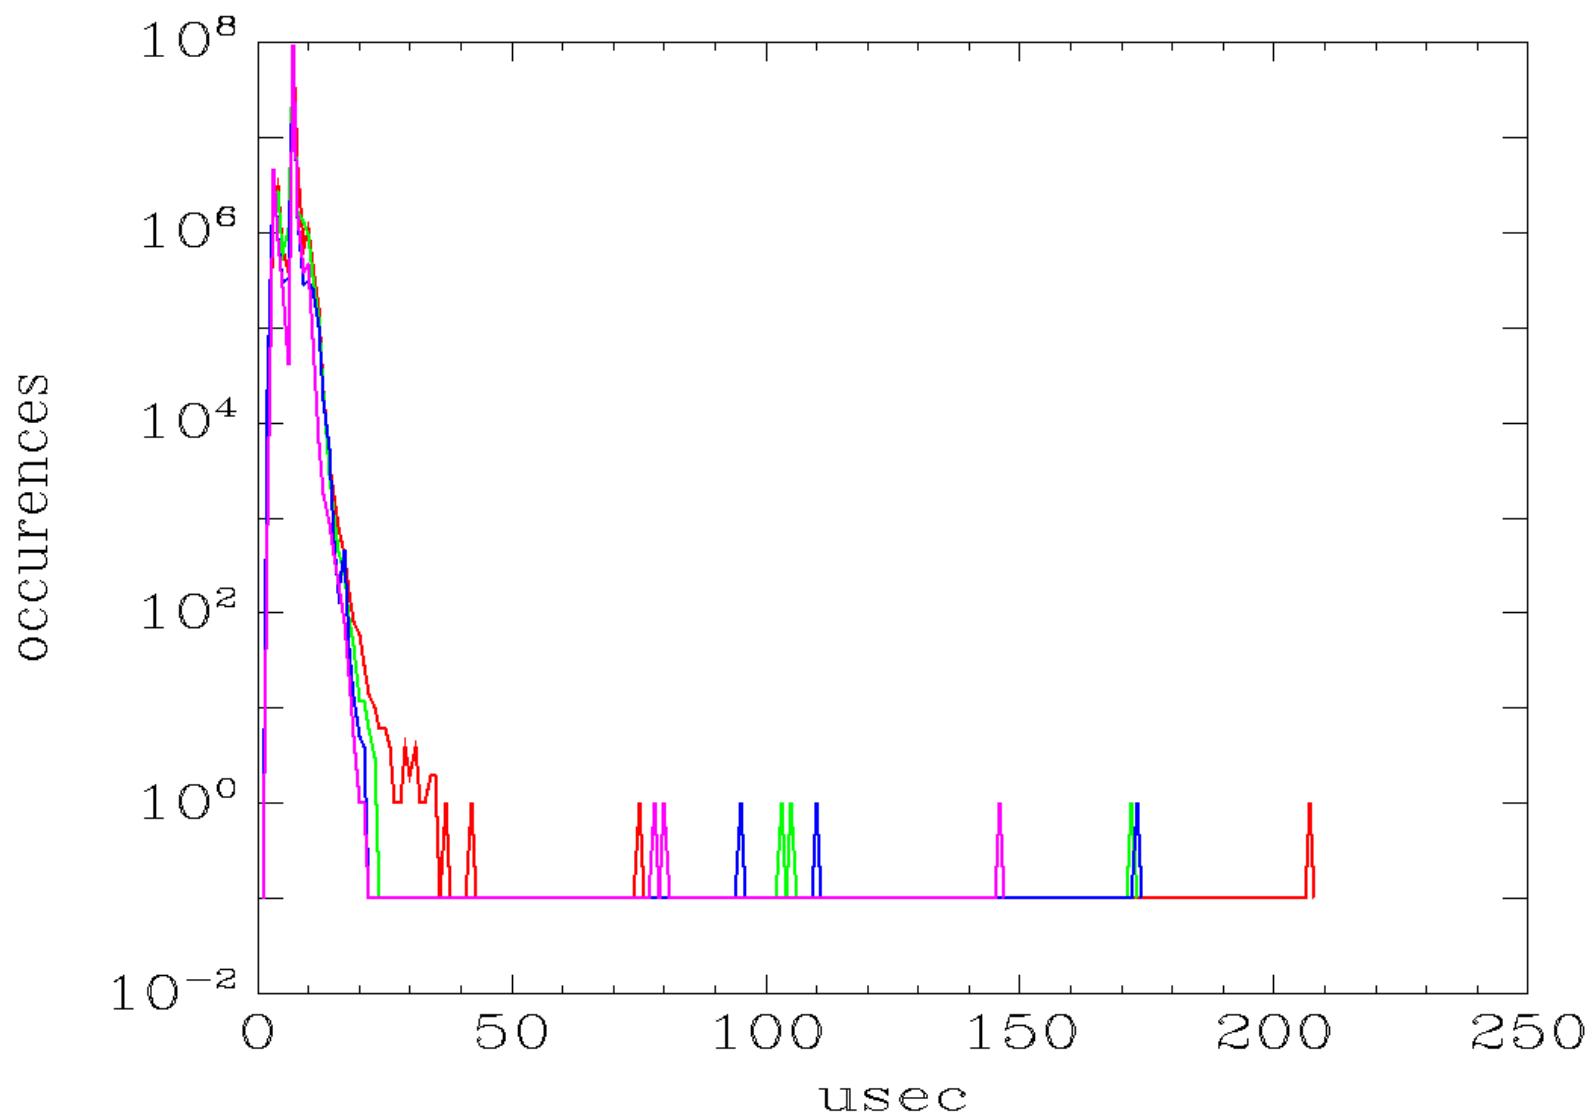




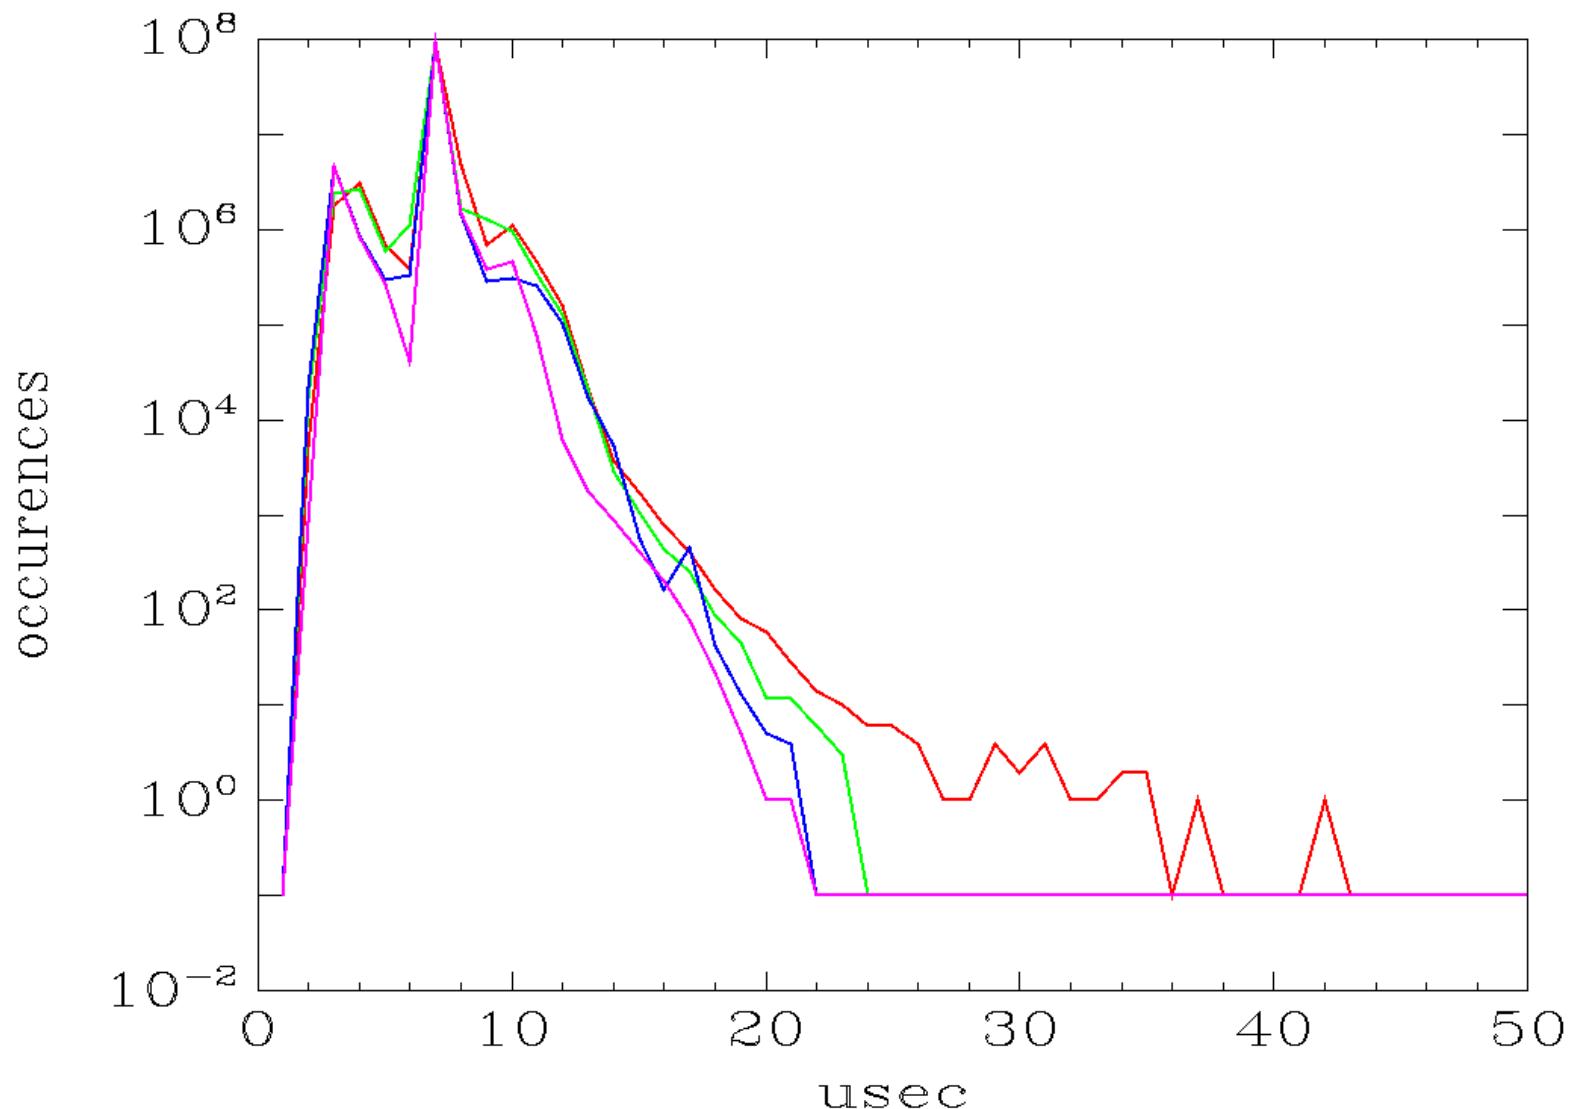
Even “boring” graphs may
contain interesting details



Latency rack5slot2
r5s2-2013-01-30-13



Latency rack5slot2
r5s2-2013-01-30-13



Command Line Options

An unruly, out of control, set of control knobs

```
$ cyclictest --help
cyclictest V 0.85
Usage:
cyclictest <options>

-a [NUM] --affinity          run thread #N on processor #N, if possible
                             with NUM pin all threads to the processor NUM
-b USEC   --breaktrace=USEC  send break trace command when latency > USEC
-B       --preemptirqs        both preempt and irqsoft tracing (used with -b)
-c CLOCK  --clock=CLOCK      select clock
                             0 = CLOCK_MONOTONIC (default)
                             1 = CLOCK_REALTIME
-C       --context            context switch tracing (used with -b)
-d DIST   --distance=DIST    distance of thread intervals in us default=500
-D       --duration=t         specify a length for the test run
                             default is in seconds, but 'm', 'h', or 'd' maybe added
                             to modify value to minutes, hours or days
-e       --latency=PM_QOS     write PM_QOS to /dev/cpu_dma_latency
-E       --event              event tracing (used with -b)
-f       --ftrace             function trace (when -b is active)
-g MAX   --of_max=MAX        Report time in ms (up to MAX) for histogram overflows
-h       --histogram=US       dump a latency histogram to stdout after the run
                             (with same priority about many threads)
                             US is the max time to be tracked in microseconds
-H       --histofall=US       same as -h except with an additional summary column
-i INTV   --interval=INTV    base interval of thread in us default=1000
-I       --irqsoft            Iirqsoft tracing (used with -b)
-l LOOPS  --loops=LOOPS      number of loops: default=0(endless)
-m       --mlockall          lock current and future memory allocations
-M       --refresh_on_max     delay updating the screen until a new max latency is hit
-n       --nanosleep          use clock_nanosleep
-N       --nsecs               print results in ns instead of us (default us)
-o RED    --oscope=RED        oscilloscope mode, reduce verbose output by RED
-O TOPT   --traceopt=TOPT     trace option
-p PRI0   --prio=PRI0        priority of highest prio thread
-P       --preemptoff         Preempt off tracing (used with -b)
-q       --quiet              print only a summary on exit
-Q       --priospread         spread priority levels starting at specified value
-r       --relative            use relative timer instead of absolute
-R       --resolution          check clock resolution, calling clock_gettime() many
                             times. list of clock_gettime() values will be
                             reported with -X
-s       --system              use sys_nanosleep and sys_setitimer
-S       --smp                 Standard SMP testing: options -a -t -n and
                             same priority of all threads
-t       --threads             one thread per available processor
-t [NUM]  --threads=NUM       number of threads:
                             without NUM, threads = max_cpus
                             without -t default = 1
-T TRACE  --tracer=TRACER    set tracing function
                             configured tracers: blk function_graph wakeup_rt wakeup function nop
-u       --unbuffered         force unbuffered output for live processing
-U       --numa                Standard NUMA testing (similar to SMP option)
                             thread data structures allocated from local node
-v       --verbose             output values on stdout for statistics
                             format: n:c:v n=tasknum c=count v=value in us
-w       --wakeup              task wakeup tracing (used with -b)
-W       --wakeuprt            rt task wakeup tracing (used with -b)
-X       --dbg_cyclictest      print info useful for debugging cyclictest
-y POLI   --policy=POLI       policy of realtime thread, POLI may be fifo(default) or rr
                             format: --policy=fifo(default) or --policy=rr
```

Thread Behavior Options

-a [NUM] --affinity	run thread #N on processor #N, if possible with NUM pin all threads to the processor NUM
-c CLOCK --clock=CLOCK	select clock 0 = CLOCK_MONOTONIC (default) 1 = CLOCK_REALTIME
-d DIST --distance=DIST	distance of thread intervals in us default=500
-i INTV --interval=INTV	base interval of thread in us default=1000
-m	lock current and future memory allocations
-n	use clock_nanosleep
-p PRI0 --prio=PRI0	priority of highest prio thread
-Q	spread priority levels starting at specified value
-r	use relative timer instead of absolute
-S	use sys_nanosleep and sys_setitimer
-S	Standard SMP testing: options -a -t -n and same priority of all threads
-t	one thread per available processor
-t [NUM] --threads=NUM	number of threads: without NUM, threads = max_cpus without -t default = 1
-U	Standard NUMA testing (similar to SMP option)
-y POLI --policy=POLI	thread data structures allocated from local node policy of realtime thread, POLI may be fifo(default) or rr format: --policy=fifo(default) or --policy=rr

side effect, sets -d0

-h	--histogram=US	dump a latency histogram to stdout after the run (with same priority about many threads) US is the max time to be tracked in microseconds
-H	--histofall=US	same as -h except with an additional summary column

Benchmark and System Options

-D --duration=t specify a length for the test run
default is in seconds, but 'm', 'h', or 'd' maybe added
to modify value to minutes, hours or days
number of loops: default=0(endless)

-l LOOPS --loops=LOOPS

-e --latency=PM_QOS write PM_QOS to /dev/cpu_dma_latency

Display Options

-g MAX	--of_max=MAX	Report time in ms (up to MAX) for histogram overflows
-h	--histogram=US	dump a latency histogram to stdout after the run (with same priority about many threads) US is the max time to be tracked in microseconds
-H	--histofall=US	same as -h except with an additional summary column
-M	--refresh_on_max	delay updating the screen until a new max latency is hit
-N	--nsecs	print results in ns instead of us (default us)
-o RED	--oscope=RED	oscilloscope mode, reduce verbose output by RED
-q	--quiet	print only a summary on exit
-u	--unbuffered	force unbuffered output for live processing
-v	--verbose	output values on stdout for statistics format: n:c:v n=tasknum c=count v=value in us

Debug Options

-b	USEC	--breaktrace=USEC	send break trace command when latency > USEC
-B		--preemptirqs	both preempt and irqoff tracing (used with -b)
-C		--context	context switch tracing (used with -b)
-E		--event	event tracing (used with -b)
-f		--ftrace	function trace (when -b is active)
-I		--irqsoff	Irqoff tracing (used with -b)
-O	TOPT	--traceopt=TOPT	trace option
-P		--preemptoff	Preempt off tracing (used with -b)
-R		--resolution	check clock resolution, calling <code>clock_gettime()</code> many times. list of <code>clock_gettime()</code> values will be reported with -X
-T	TRACE	--tracer=TRACER	set tracing function
		configured tracers: blk function_graph wakeup_rt wakeup function nop	
-w		--wakeup	task wakeup tracing (used with -b)
-W		--wakeuprt	rt task wakeup tracing (used with -b)
-X		--dbg_cyclictest	print info useful for debugging cyclictest

Debug Options

No time to describe in this talk

Hooks to invoke various tools that can capture
the cause of large latencies

Options Trivia

Options parsing is not robust - example 1

affinity will be 0

```
$ cyclictest -t -l100 -a0
```

```
$ cyclictest -t -l100 -a 0
```

```
$ cyclictest -t -l100 -a7 -a0
```

affinity will be 7, with no error message

```
$ cyclictest -t -l100 -a7 -a 0
```

-a cpu affinity

Options Trivia

Options parsing is not robust - example 2

```
$ cyclictest -ant
```

```
T: 0 (26978) P: 0 I:1000 C: 2091 Min:
```

```
$ cyclictest -an -t
```

```
T: 0 (26980) P: 0 I:1000 C: 1928 Min:
```

```
T: 1 (26981) P: 0 I:1500 C: 1285 Min:
```

-a cpu affinity

-n clock_nanosleep()

-t one thread per cpu

Options Trivia

Options parsing is not robust

Best Practice:

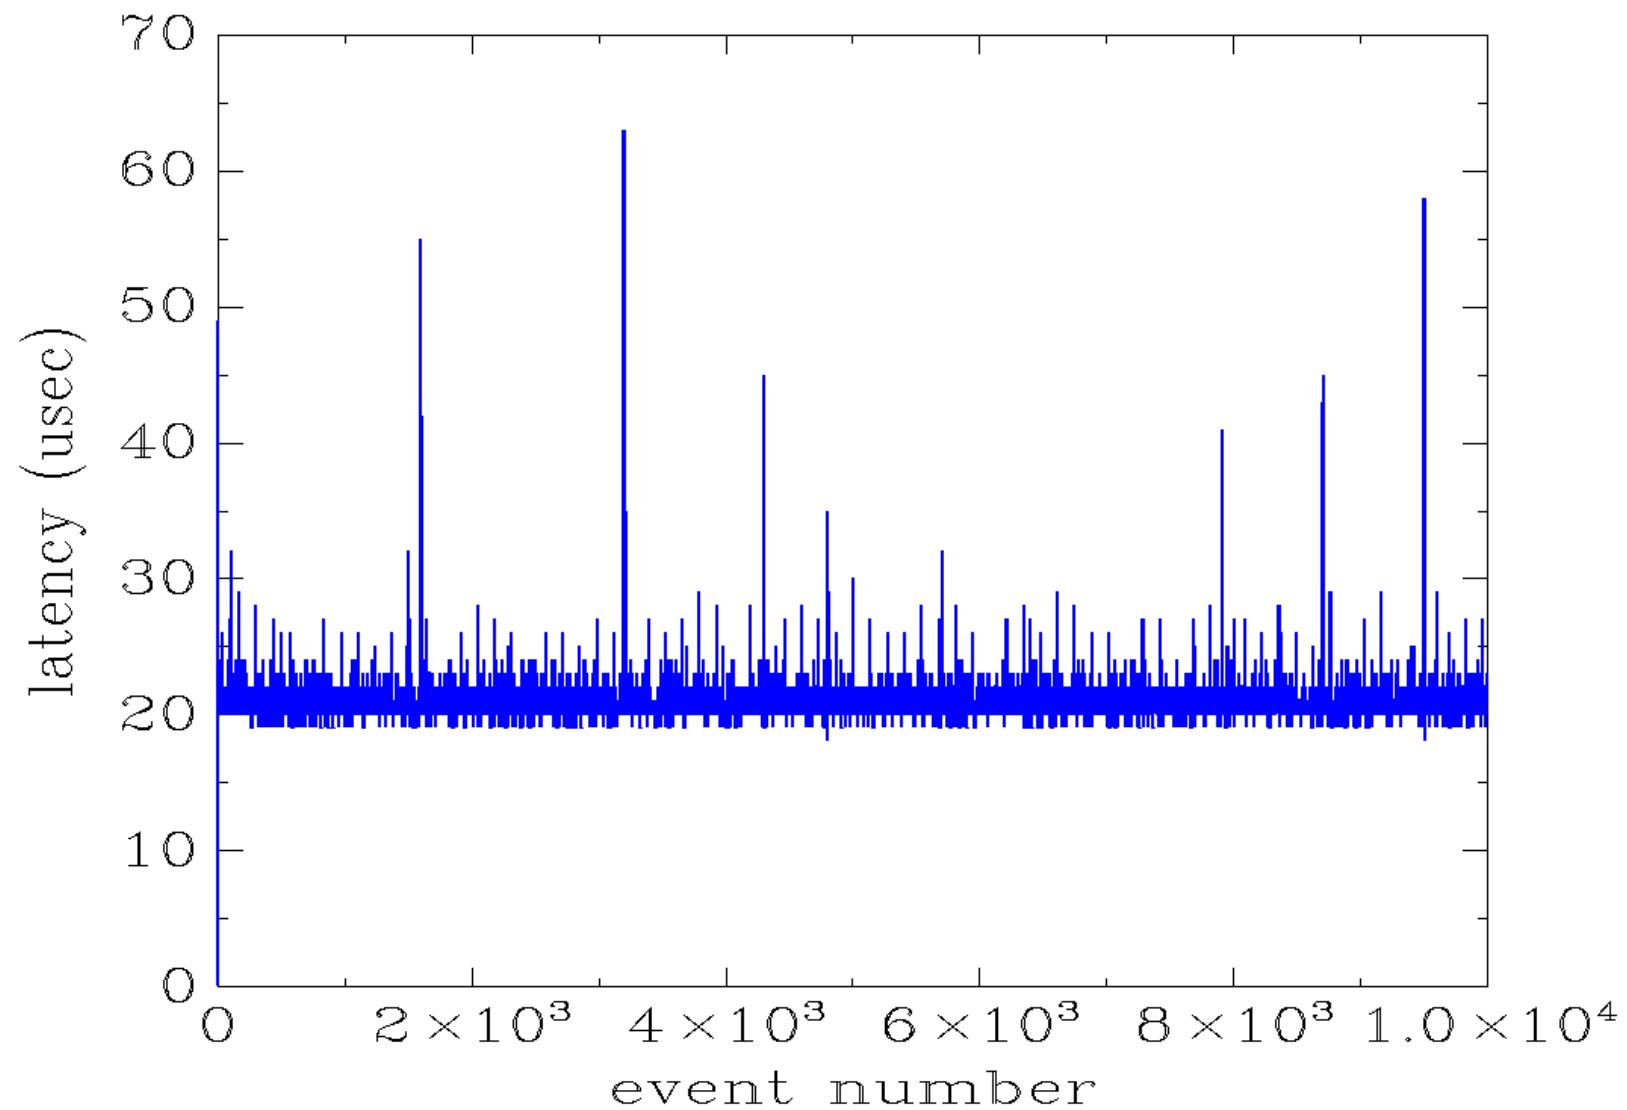
- do not combine options
- specify each separately with a leading "-"

Third Data Format

Report each latency

```
$ cyclictest -q -n -t1 -p 48 -l 10000 -v
```

Cyclictest Latency
pri=51 ping flood



Hitting the RT sched throttle

/proc/sys/kernel/sched_rt_runtime_us

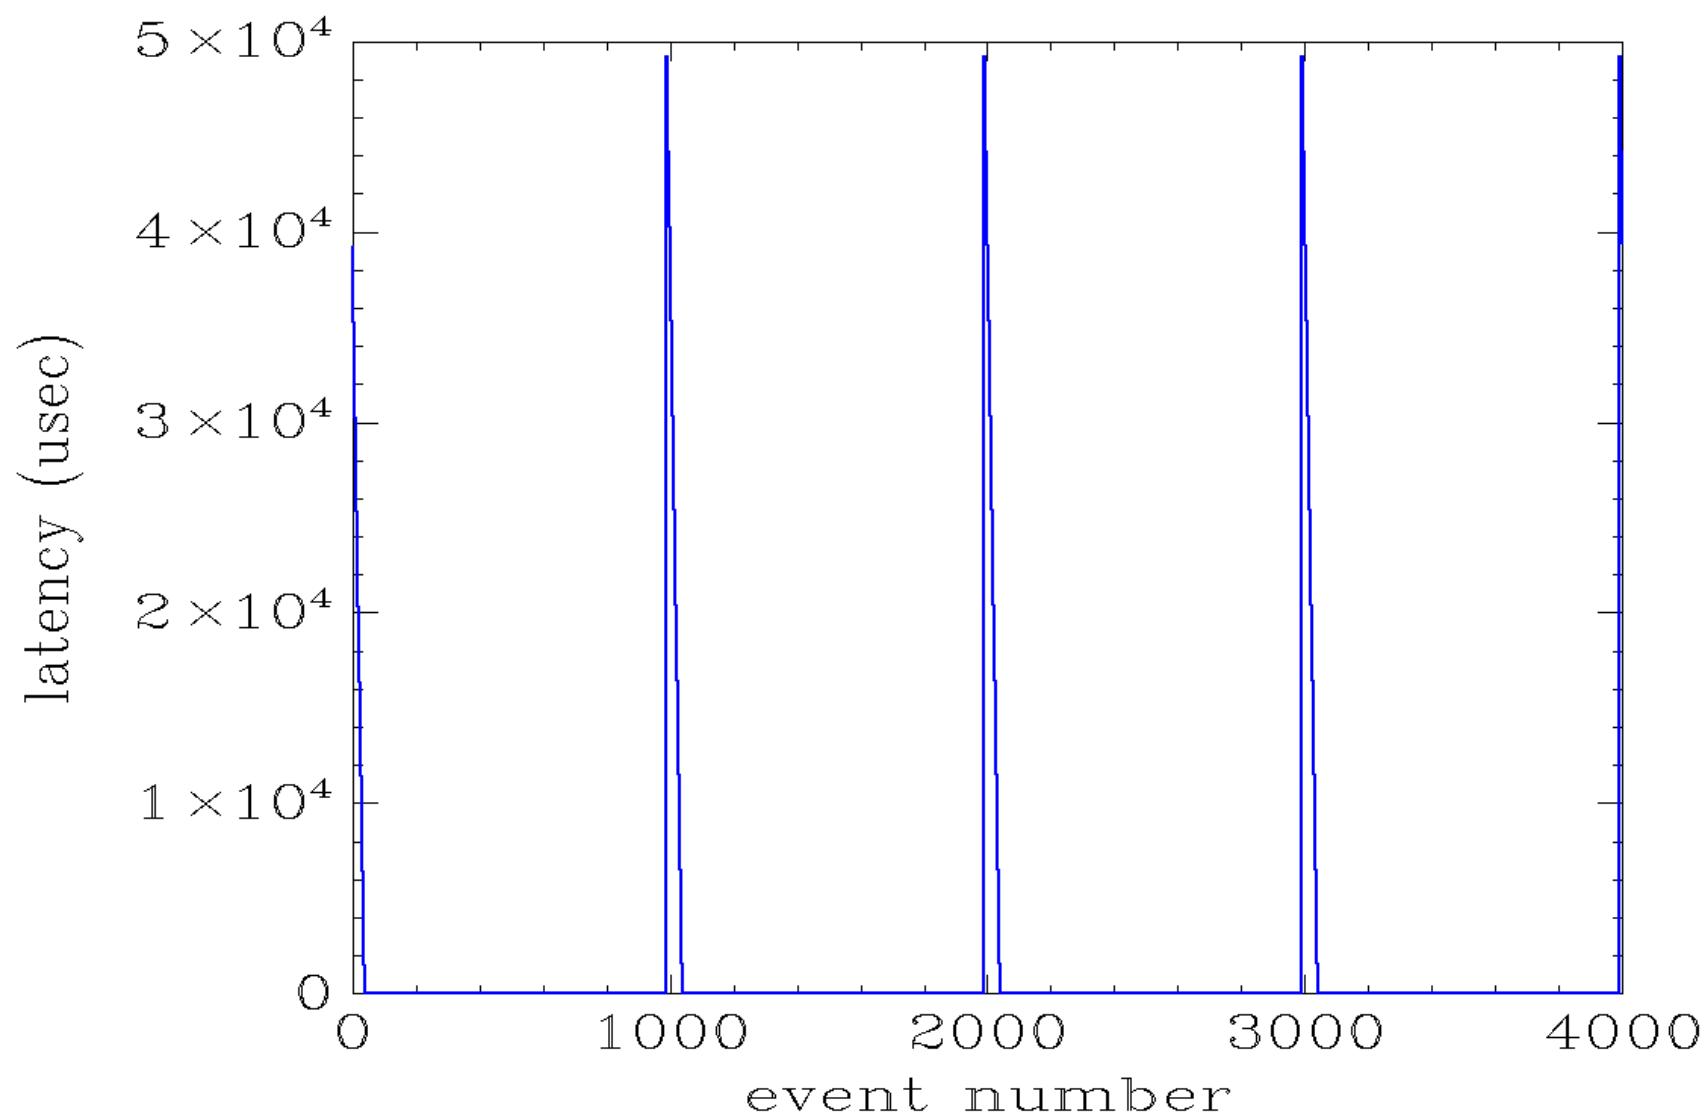
/proc/sys/kernel/sched_rt_period_us

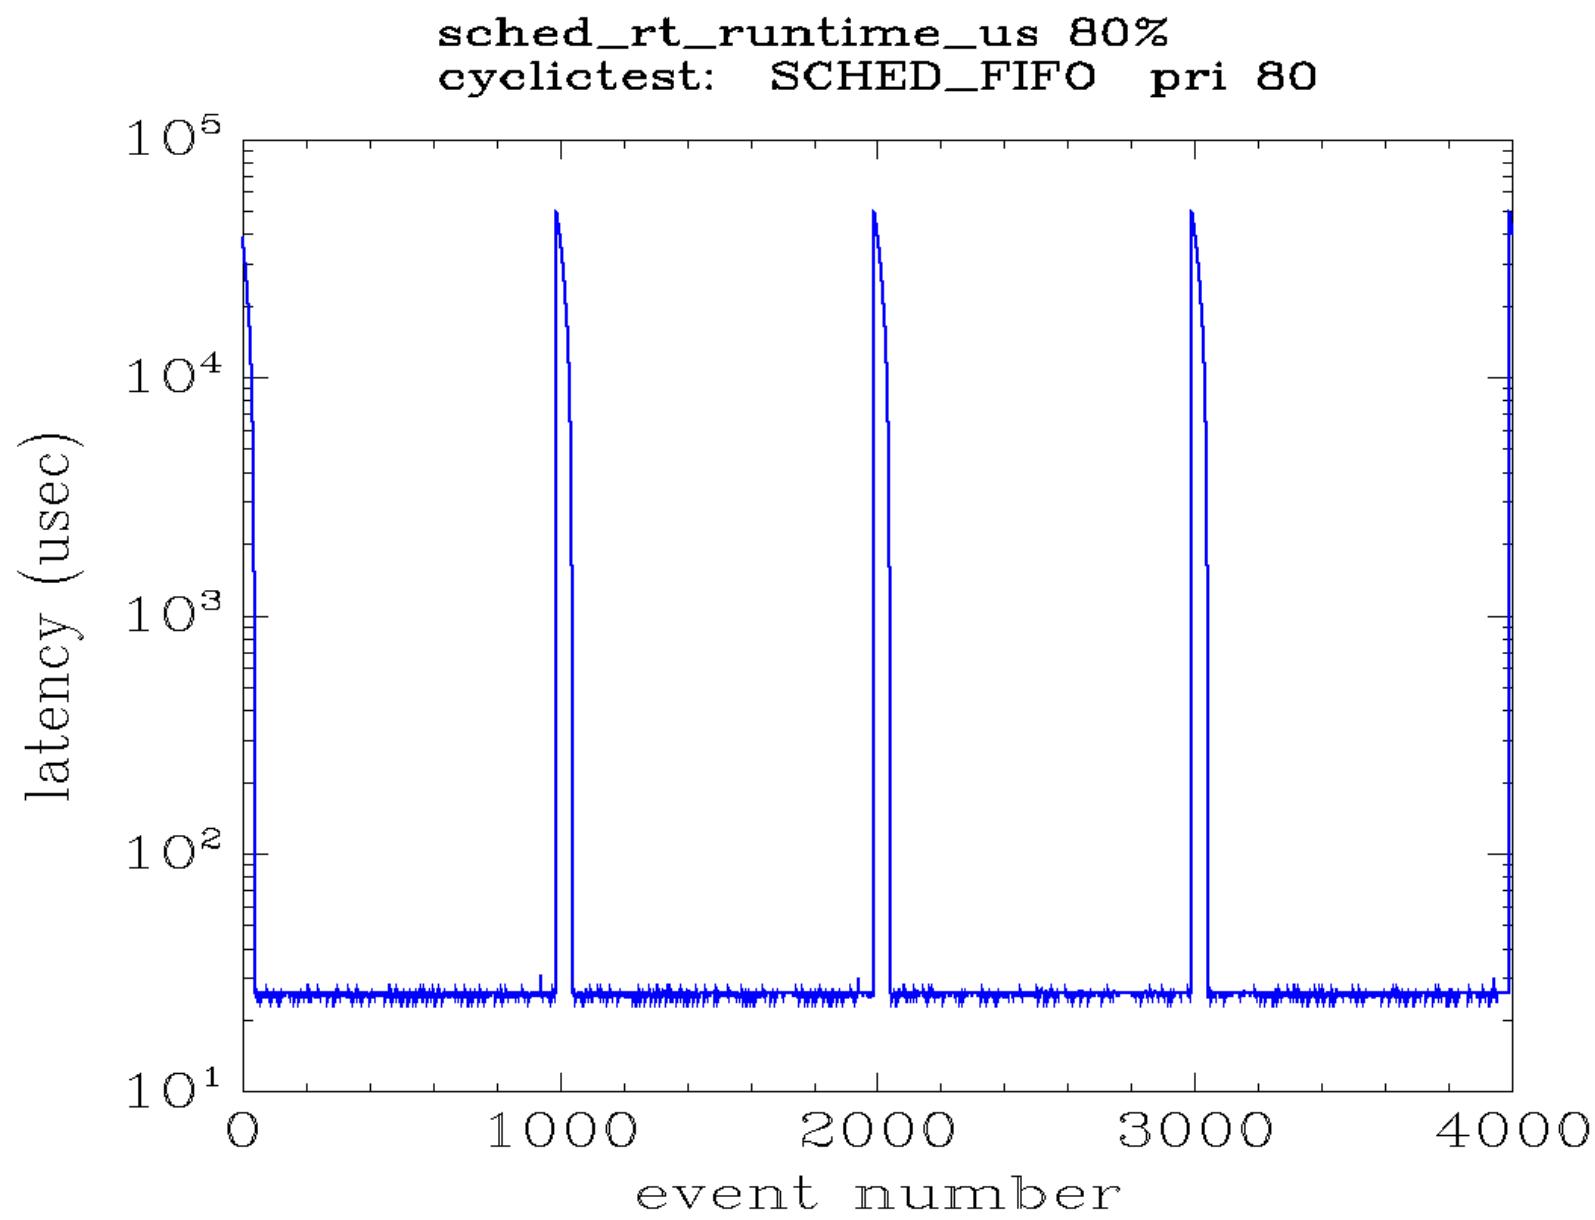
cyclictest: SCHED_FIFO priority=80

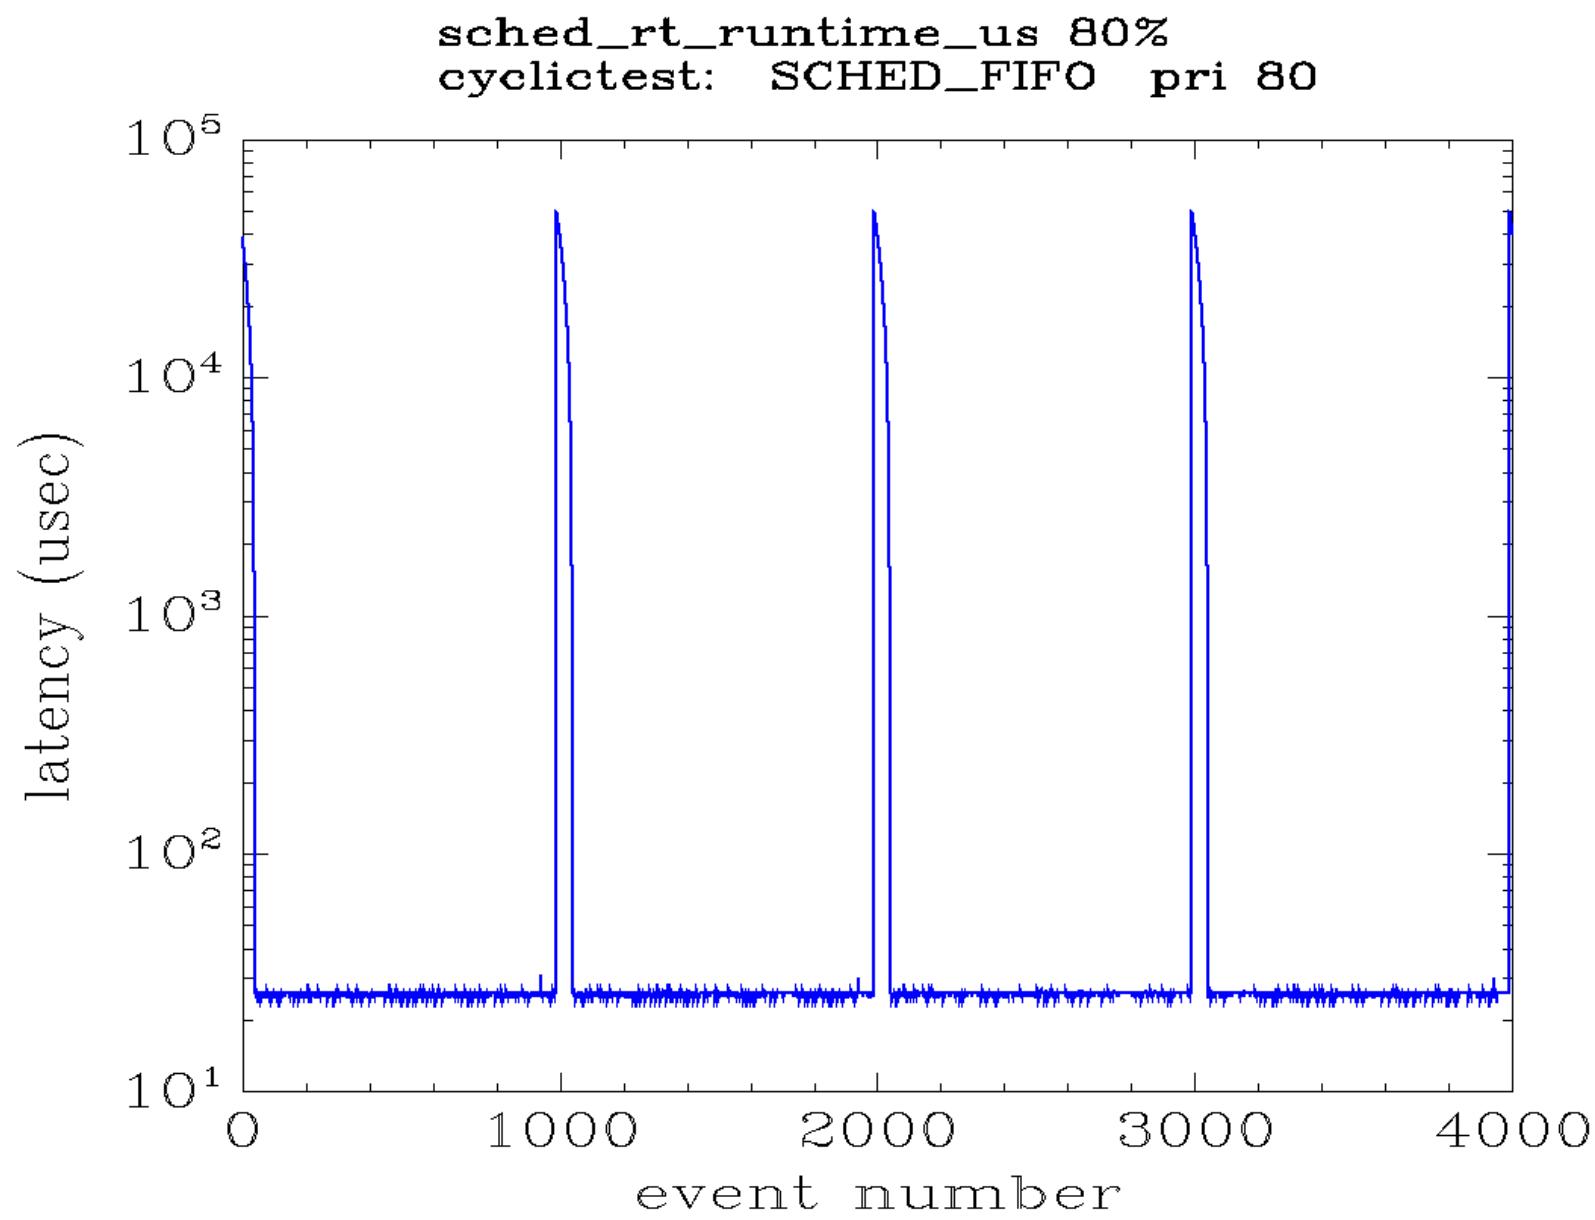
background load:

- continuous
- SCHED_FIFO priority=40

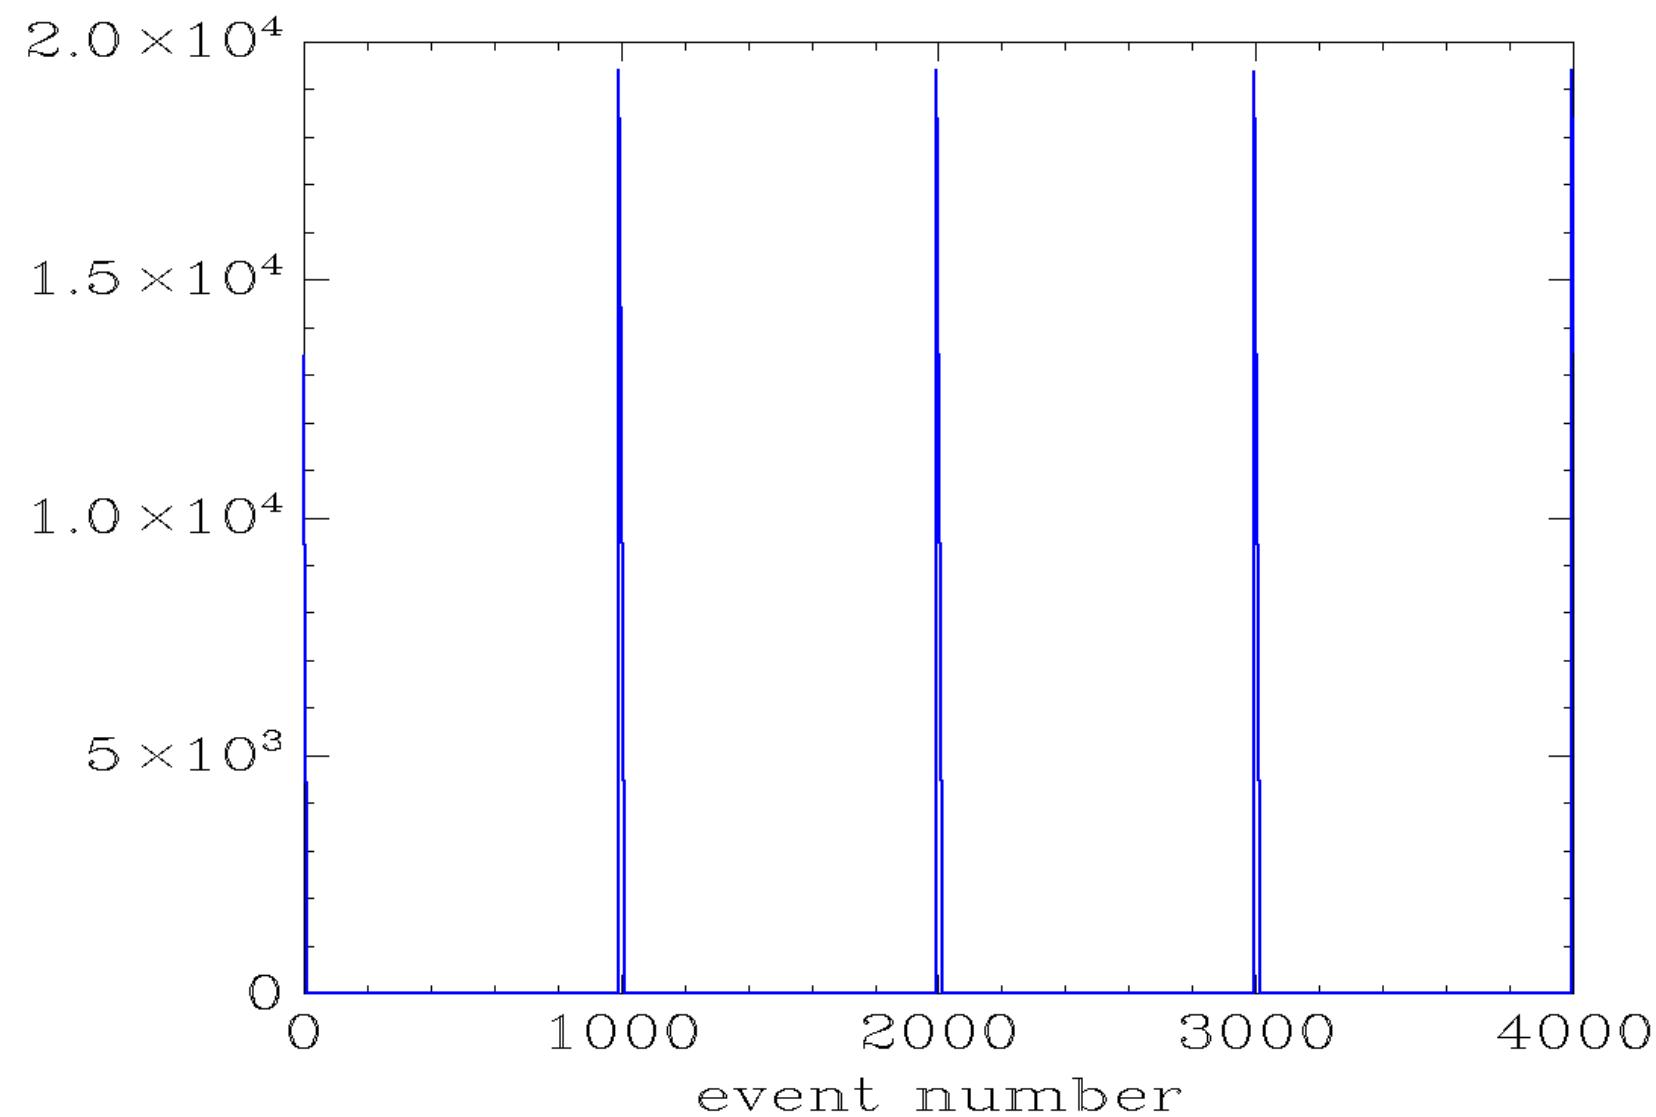
sched_rt_runtime_us 80%
cyclictest: SCHED_FIFO pri 80

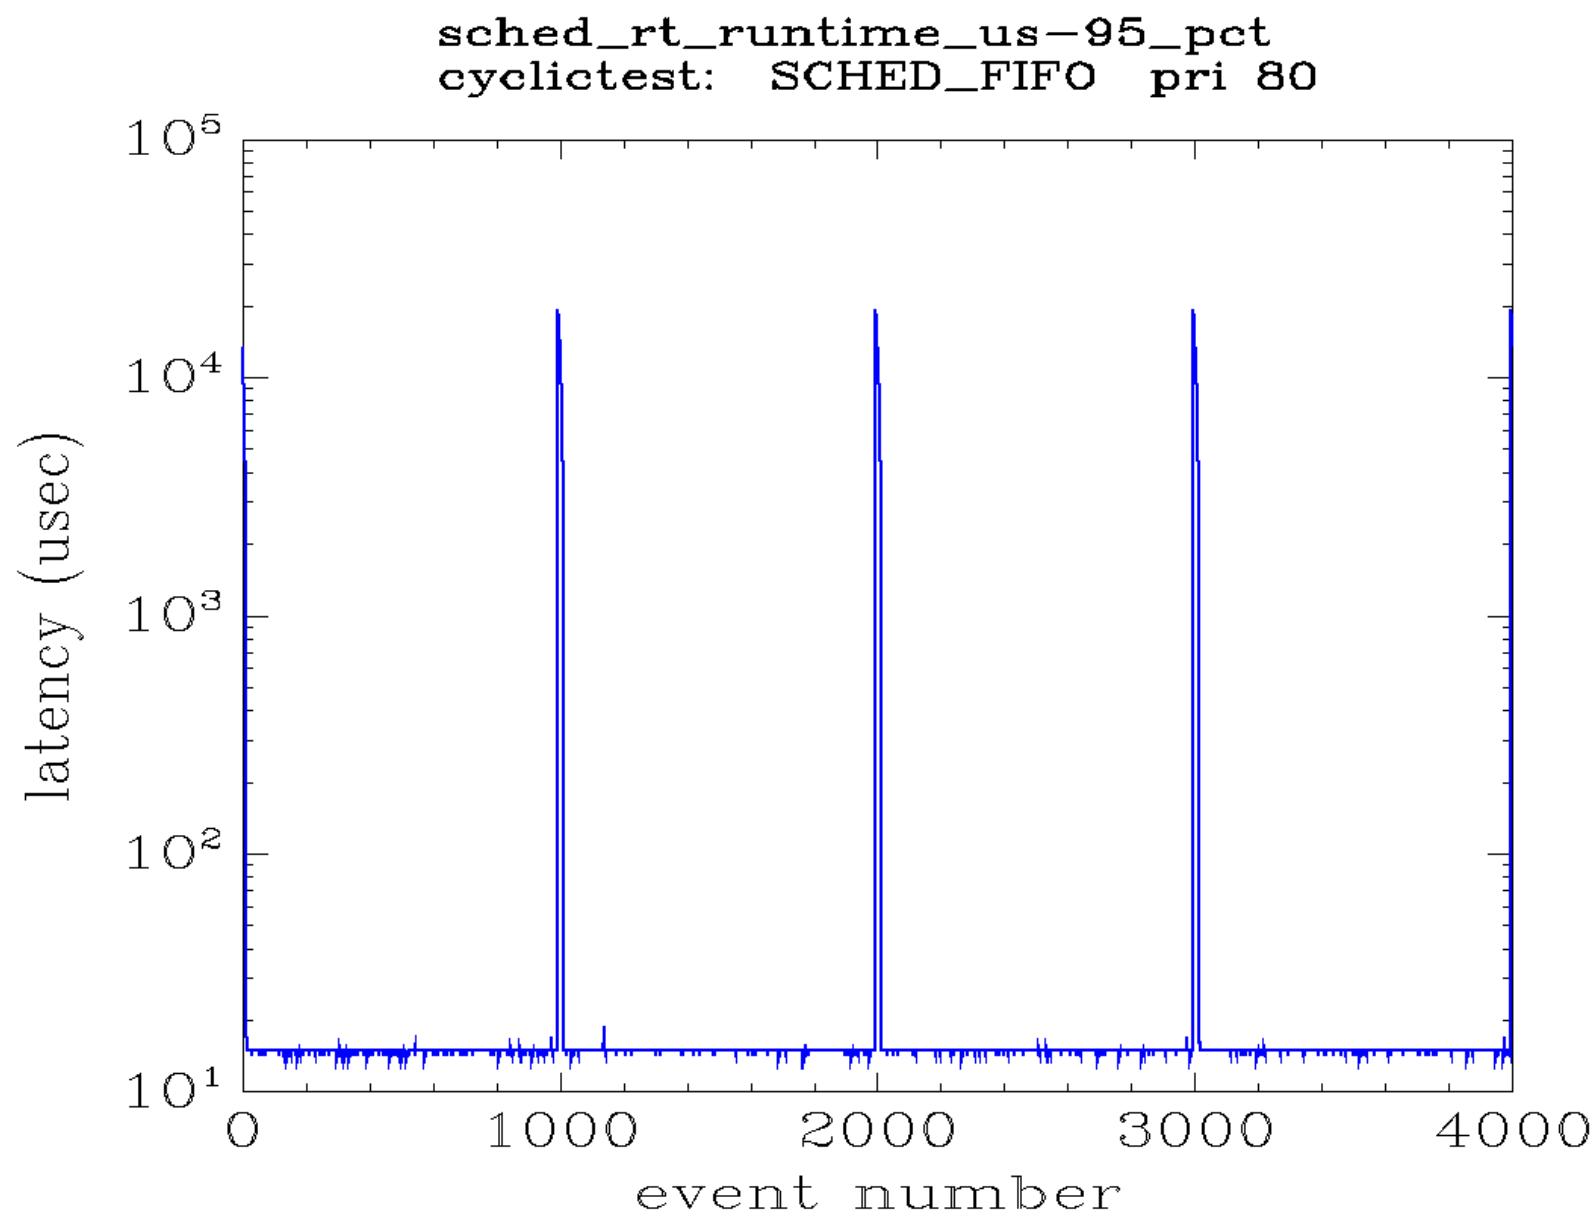






sched_rt_runtime_us-95_pct
cyclictest: SCHED_FIFO pri 80





Hitting the RT sched throttle

/proc/sys/kernel/sched_rt_runtime_us

/proc/sys/kernel/sched_rt_period_us

cyclictest: SCHED_NORMAL

background load:

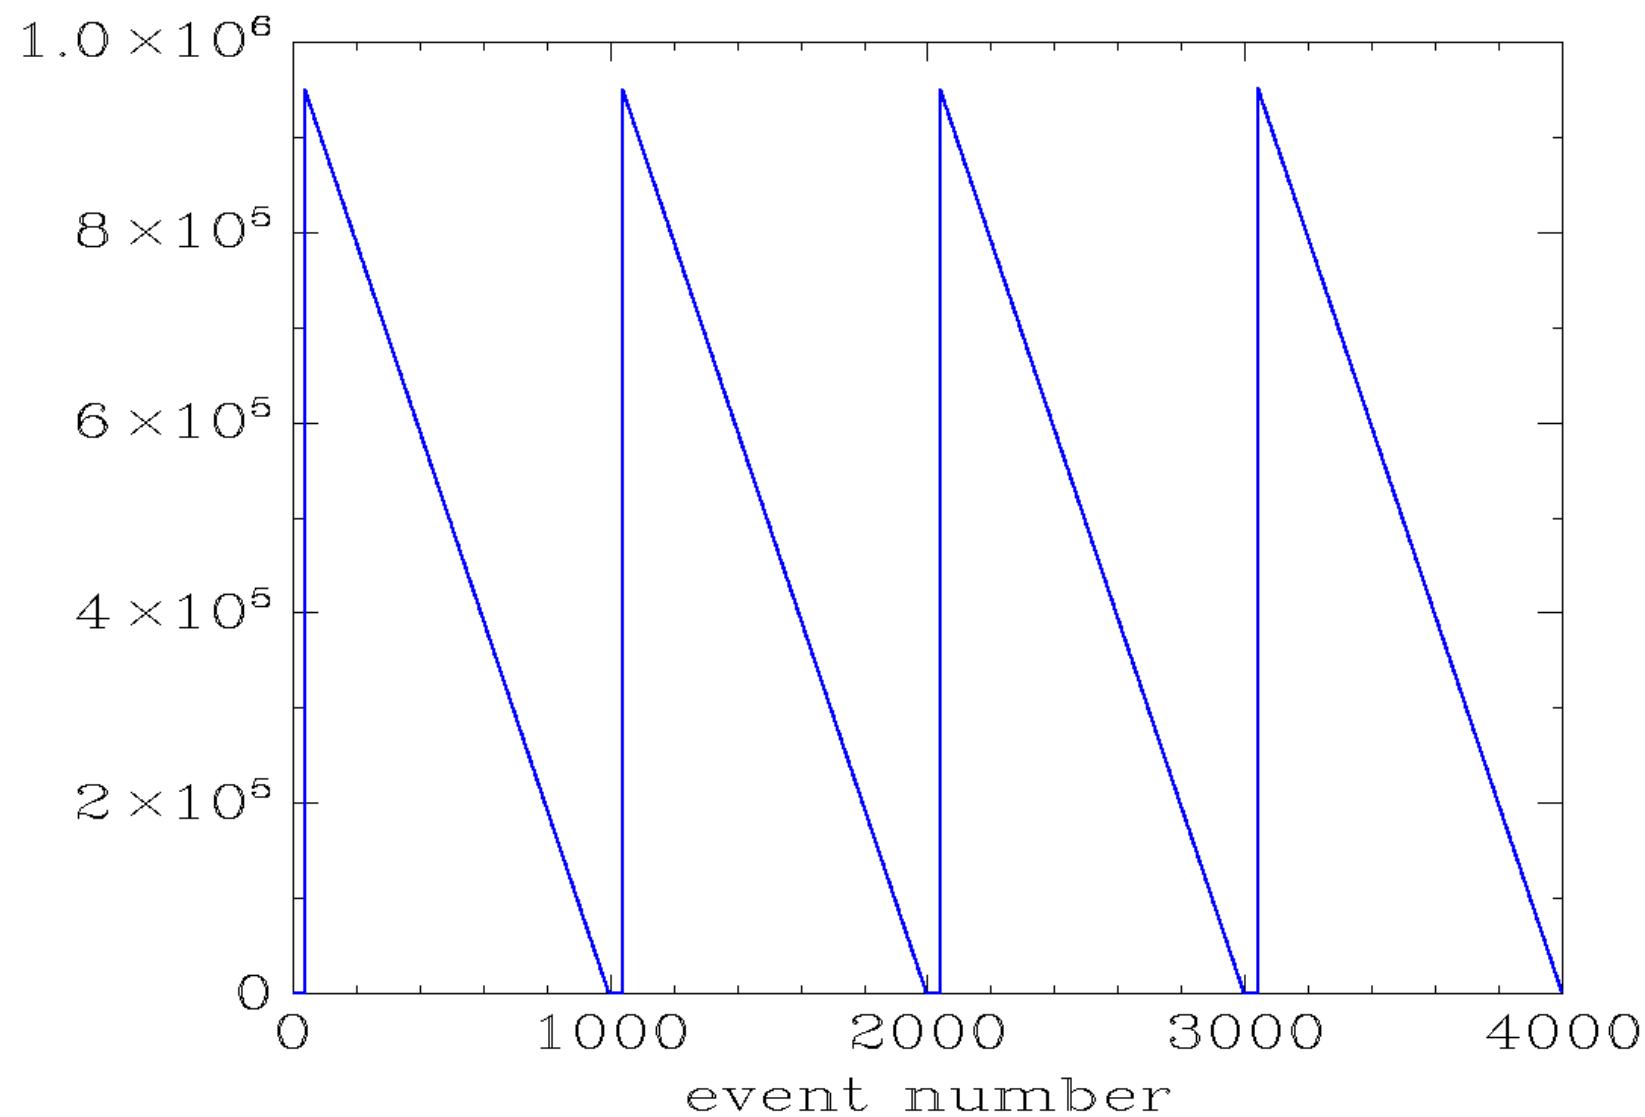
- continuous
- SCHED_FIFO priority=40

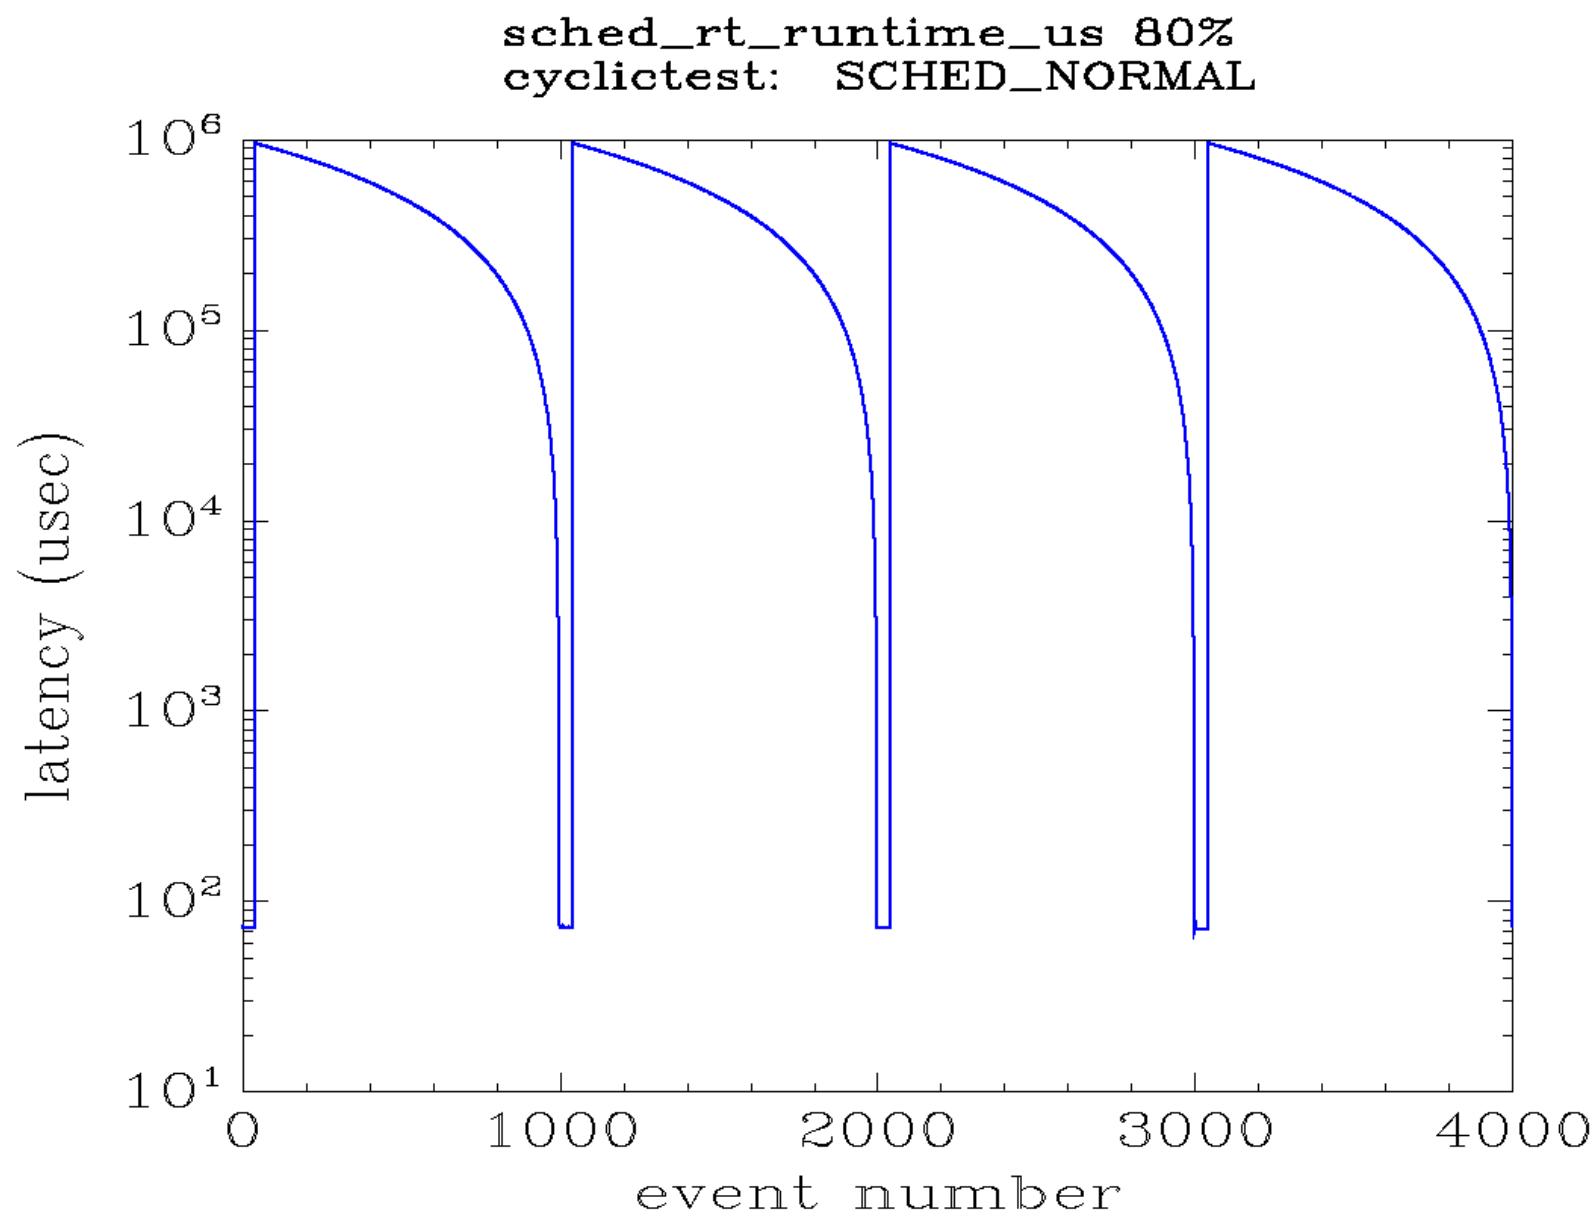
Hitting the RT sched throttle

Why is this measurement interesting???

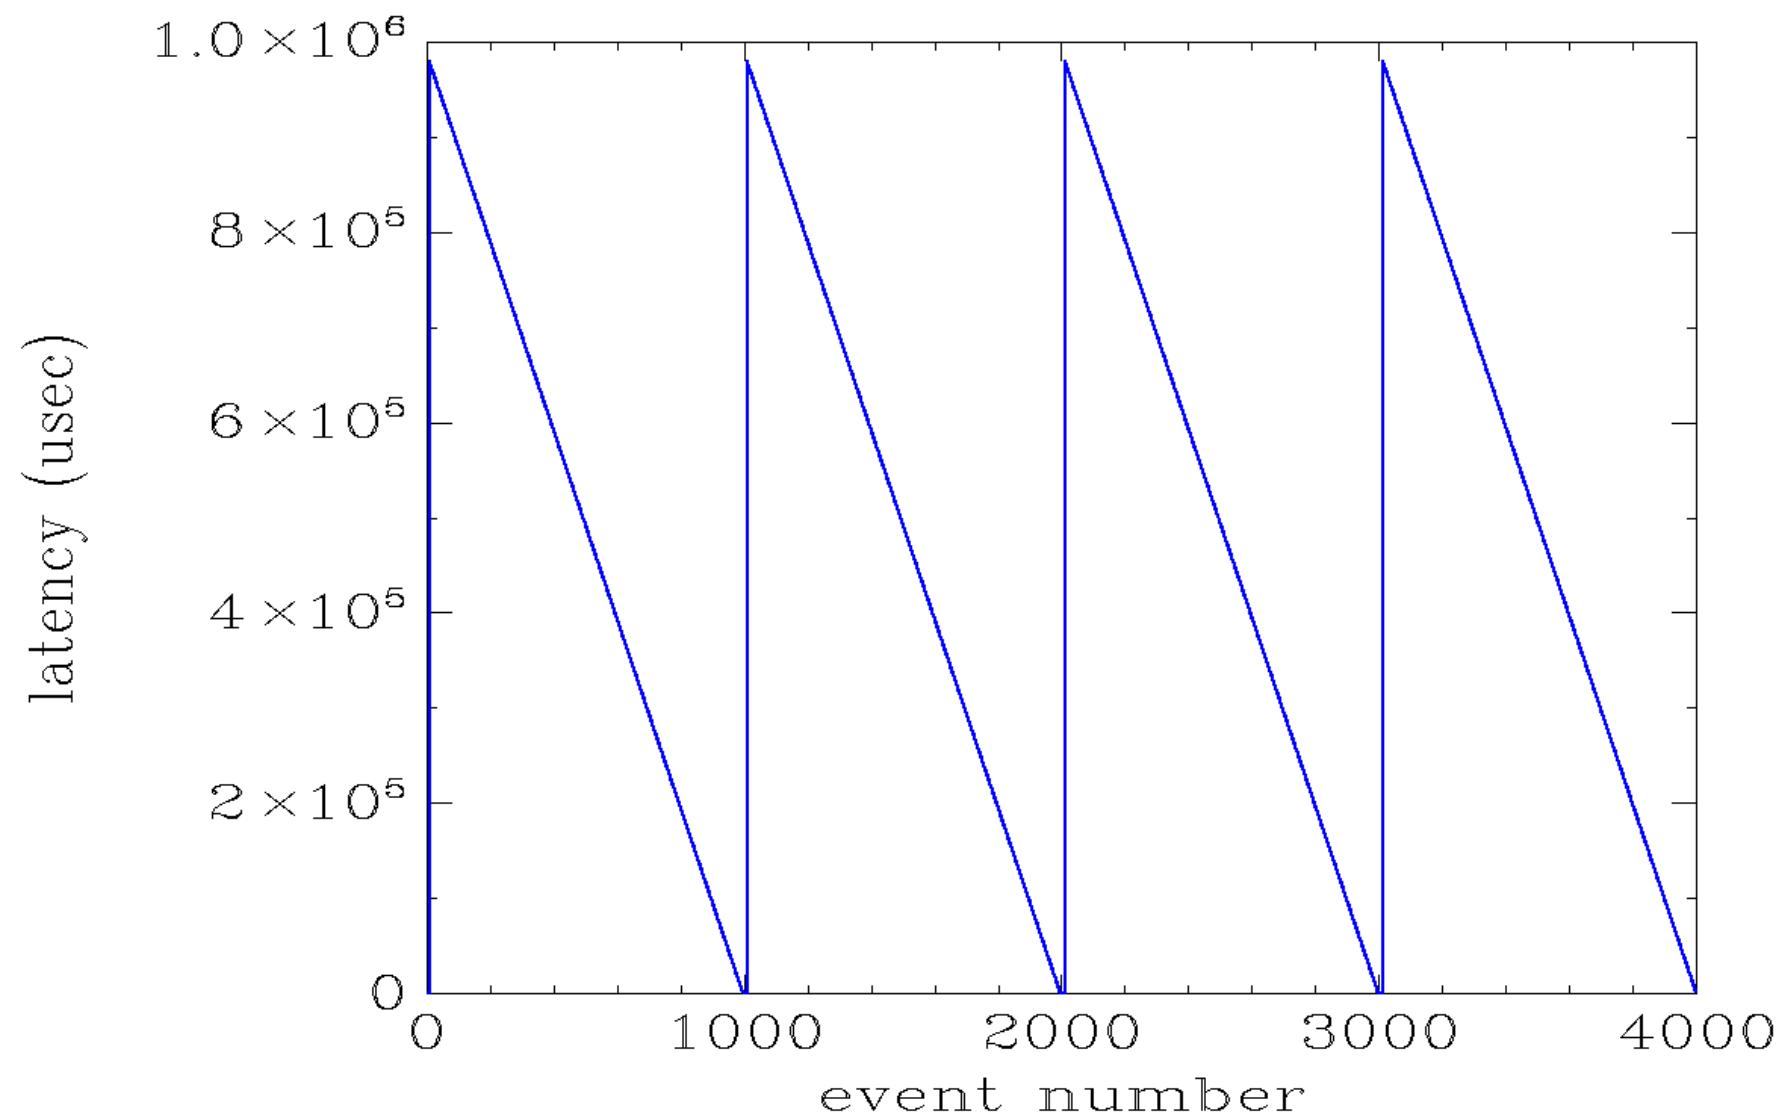
Gives a picture of how much cpu is NOT used by the real time tasks

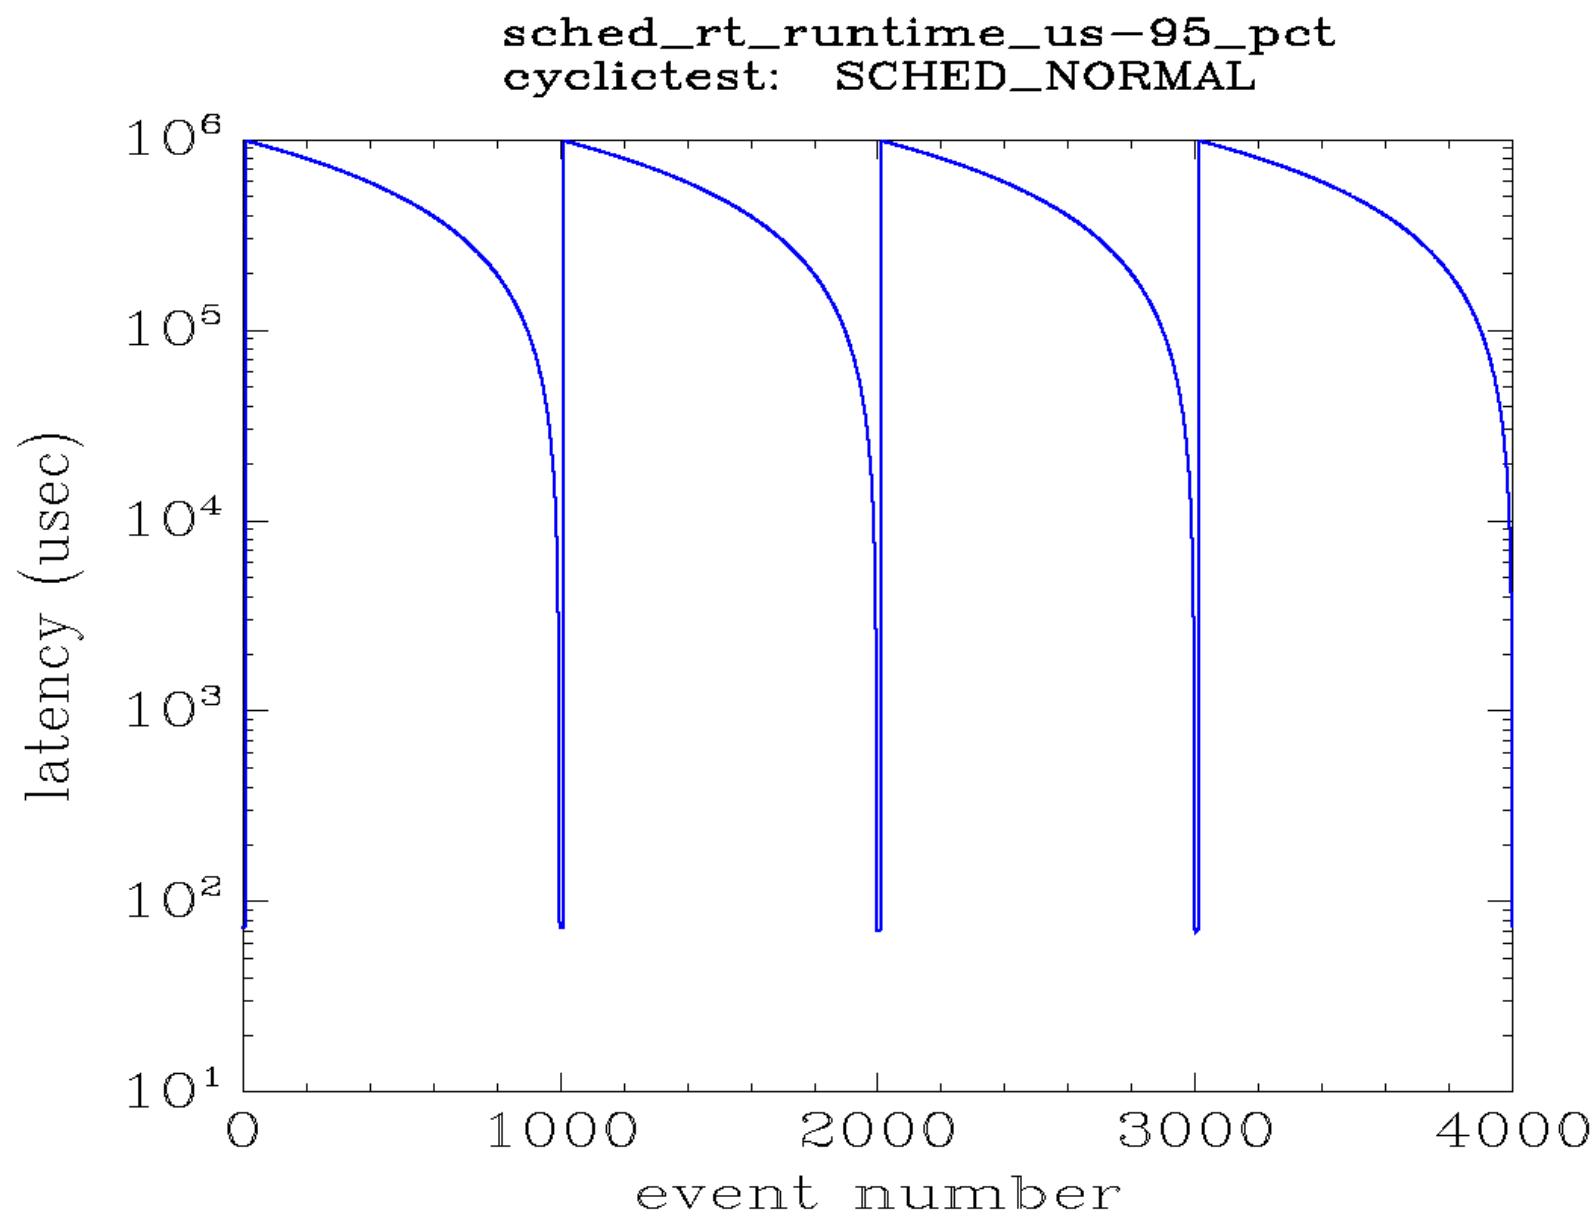
sched_rt_runtime_us 80%
cyclictest: SCHED_NORMAL





sched_rt_runtime_us-95_pct
cyclictest: SCHED_NORMAL





Unusual Uses of Cyclictest

Rough measurement of response time of a task, without instrumenting task.

cyclictest latency =~

task latency + task work duration

This is not an accurate measurement, but it does provide a rough picture.

Response Time of a Task

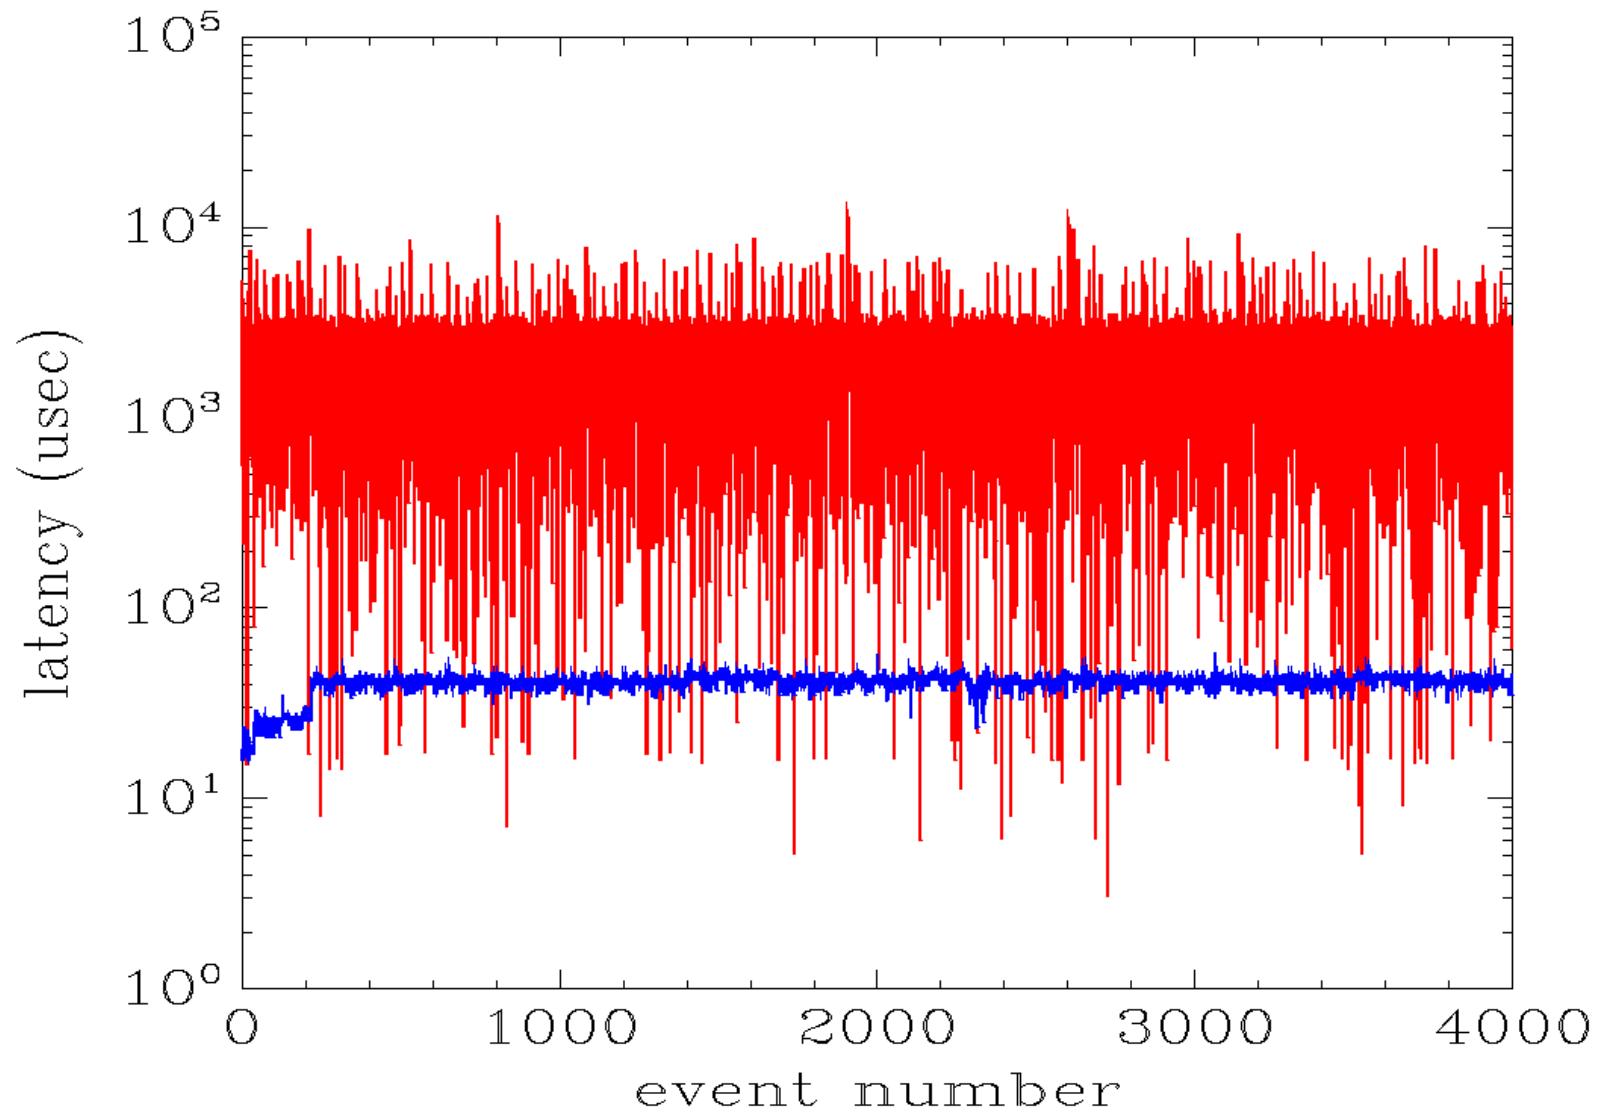
Cyclictest:

- (1) SCED_FIFO priority=30
baseline latency
- (2) SCED_FIFO priority=80
approximation of task response time

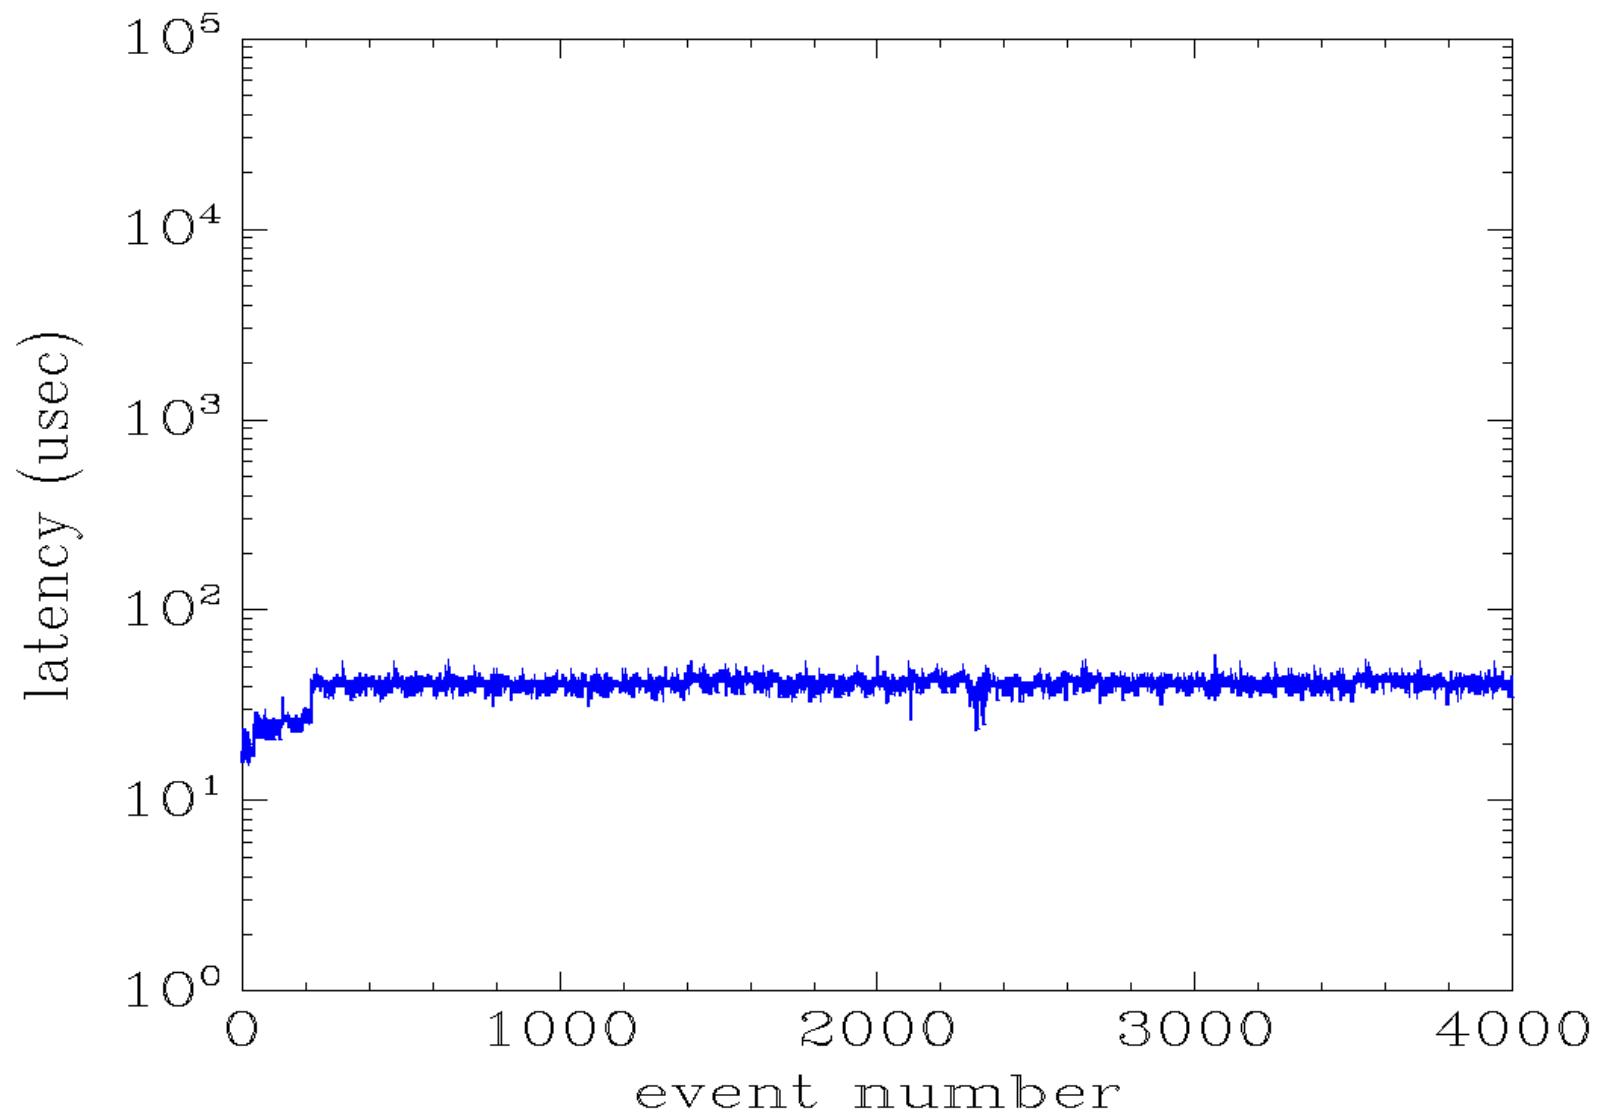
background load:

- busy loop, random number of iterations
- SCED_FIFO priority=40

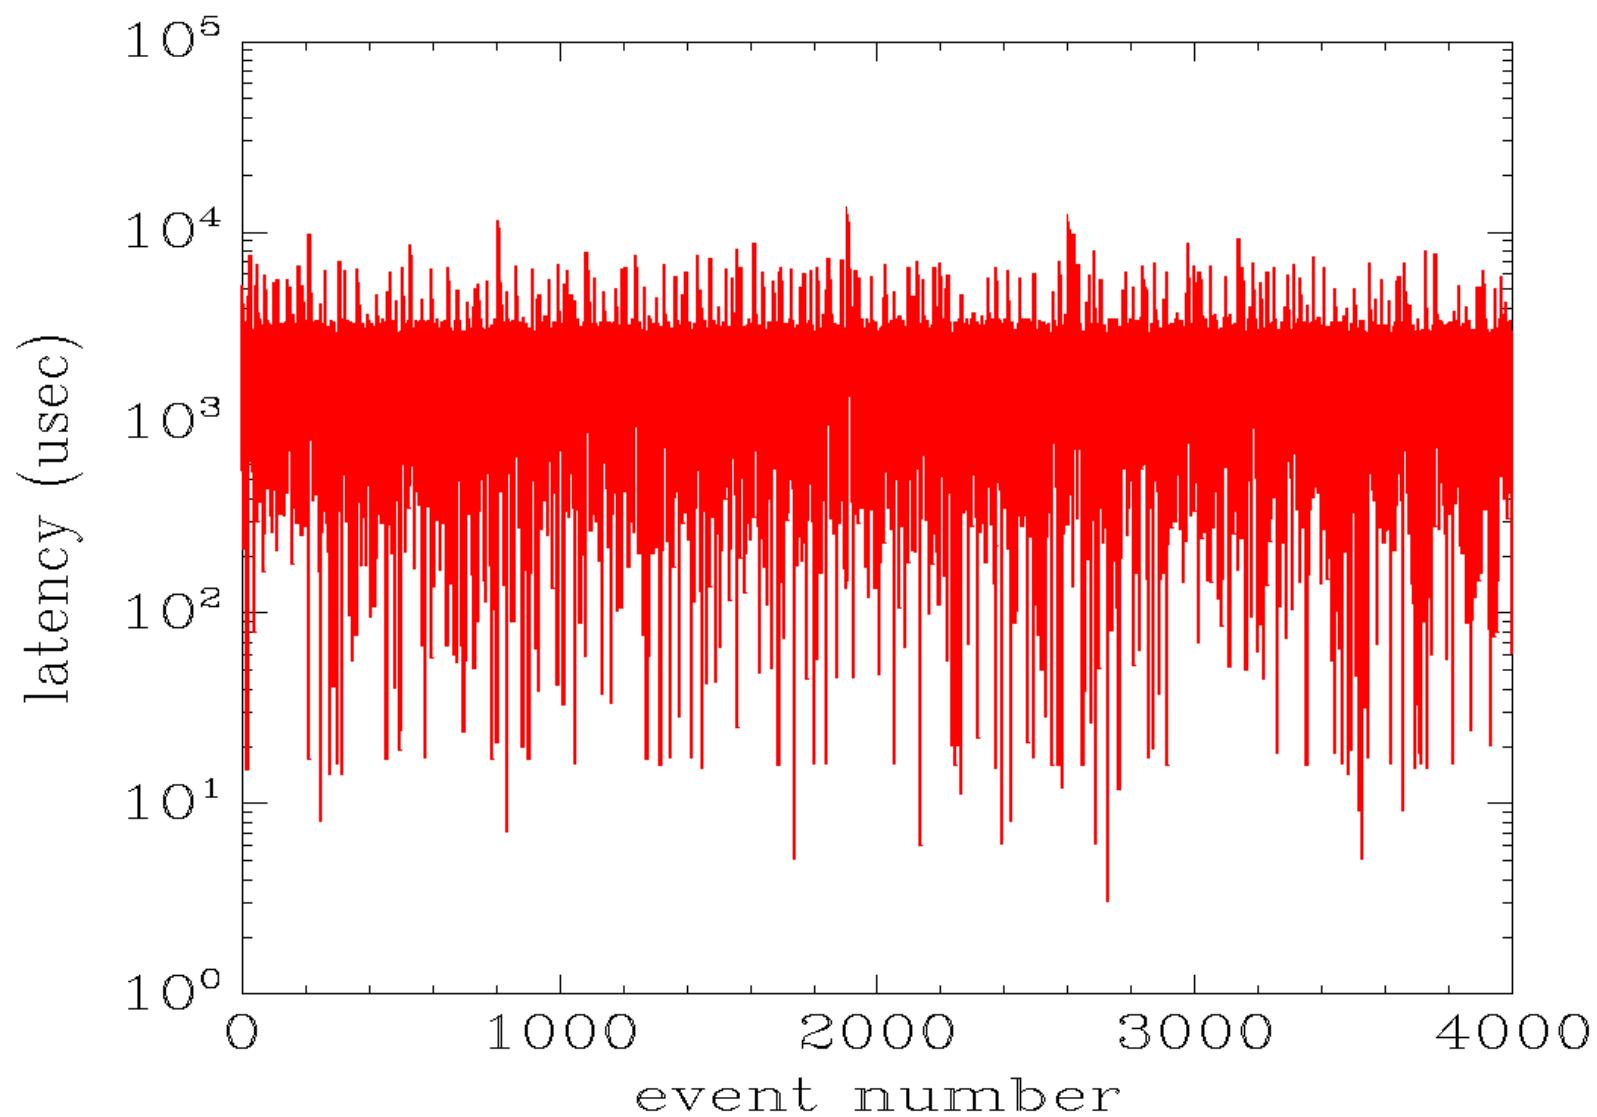
blue: pri better than load red: pri worse than load
load: random busy loop



pri == 80 (better than load)
load: random busy loop



pri == 30 (worse than load)
load: random busy loop



Response Time of a Task

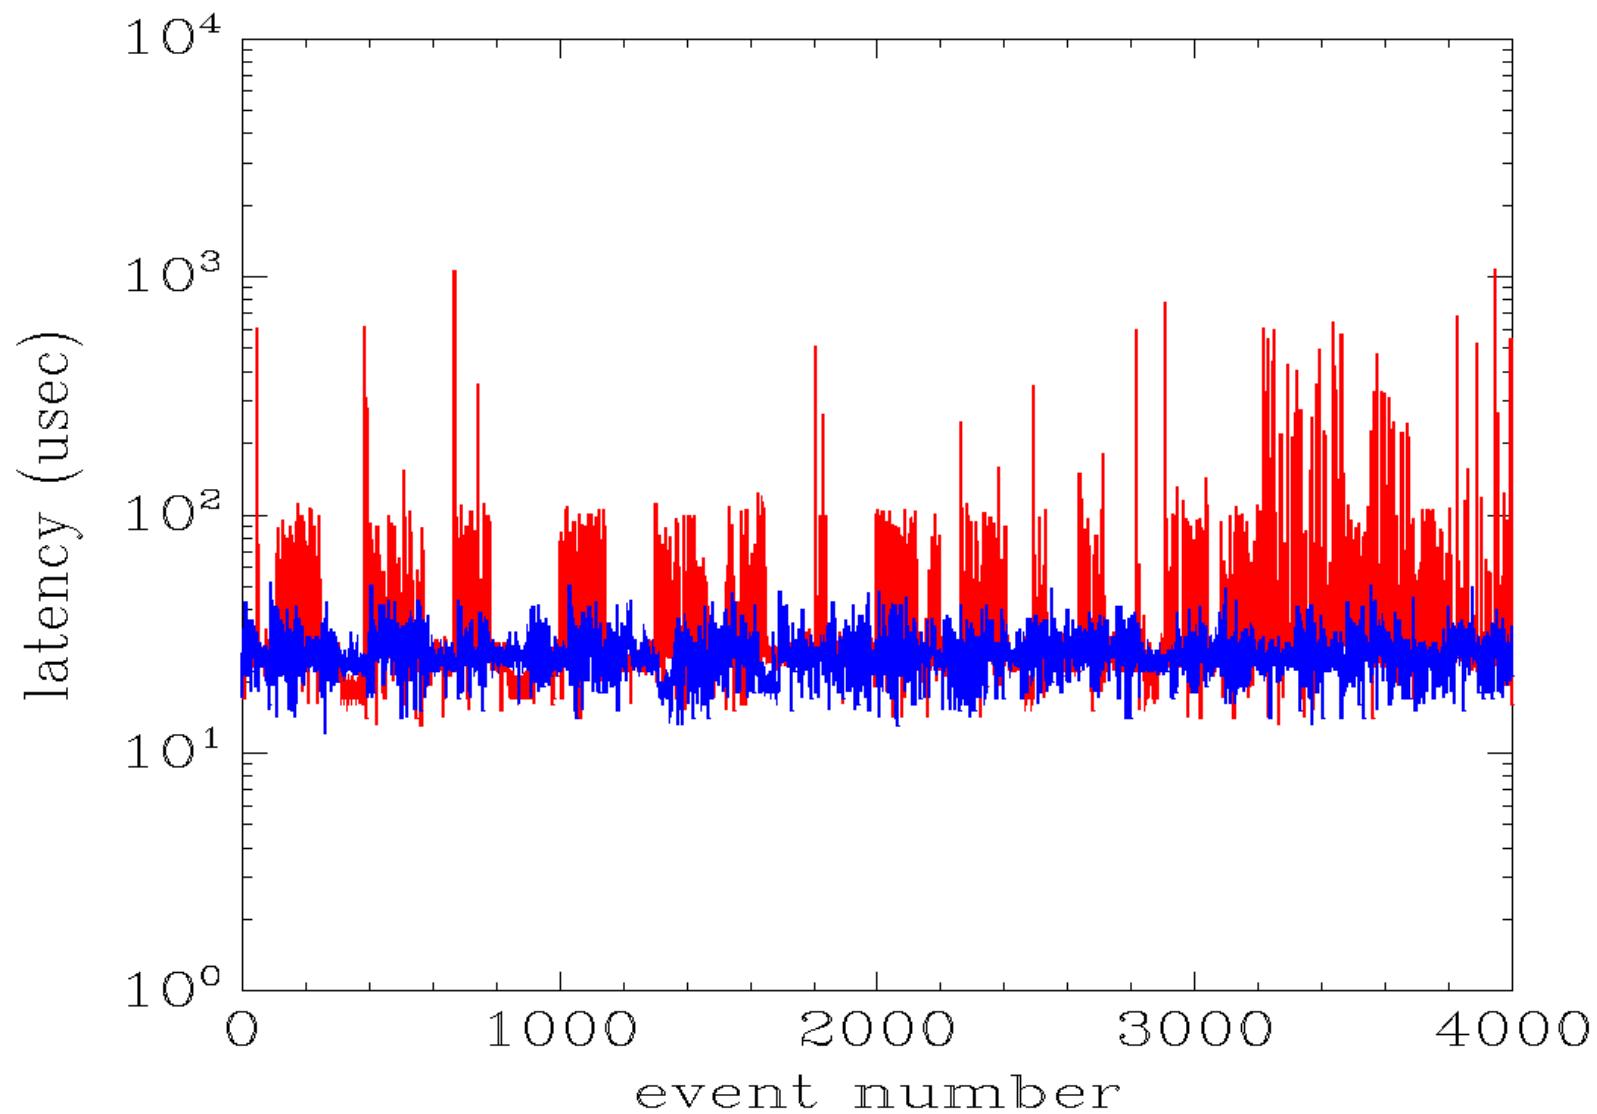
Cyclictest:

- (1) SCED_FIFO priority=30
baseline latency
- (2) SCED_FIFO priority=80
approximation of task response time

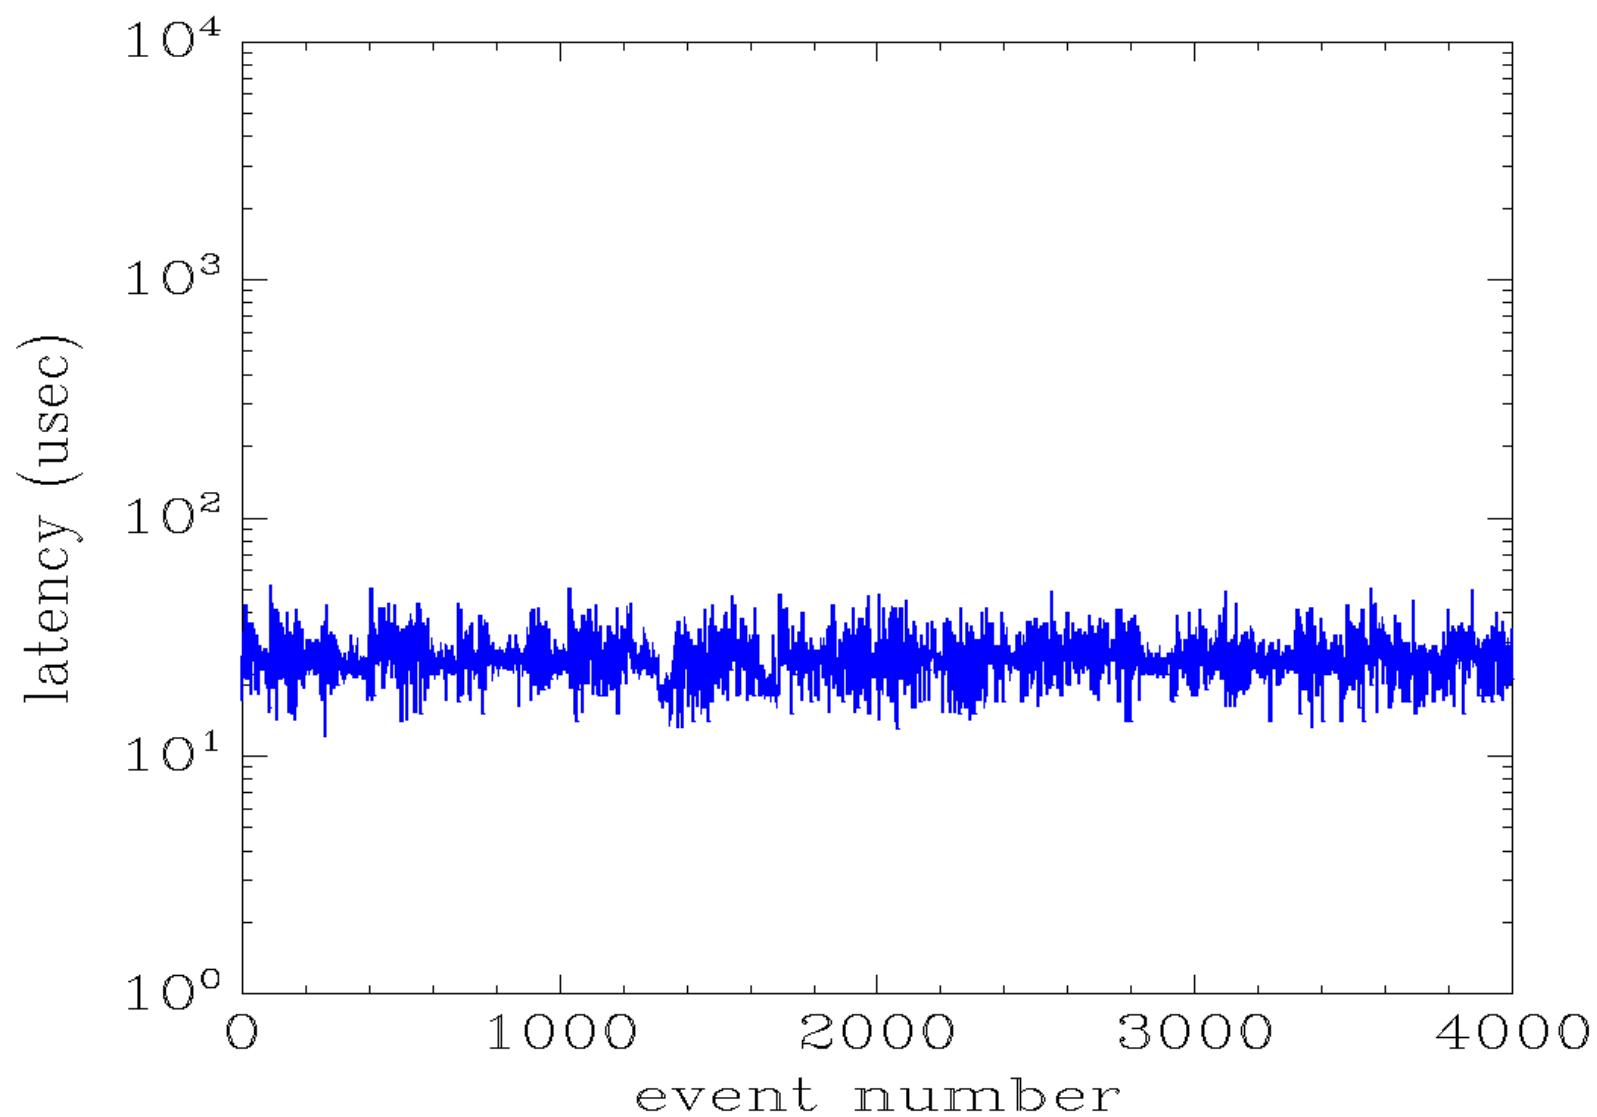
background load:

- recursive scp of a panda root file system
- SCED_FIFO priority=40

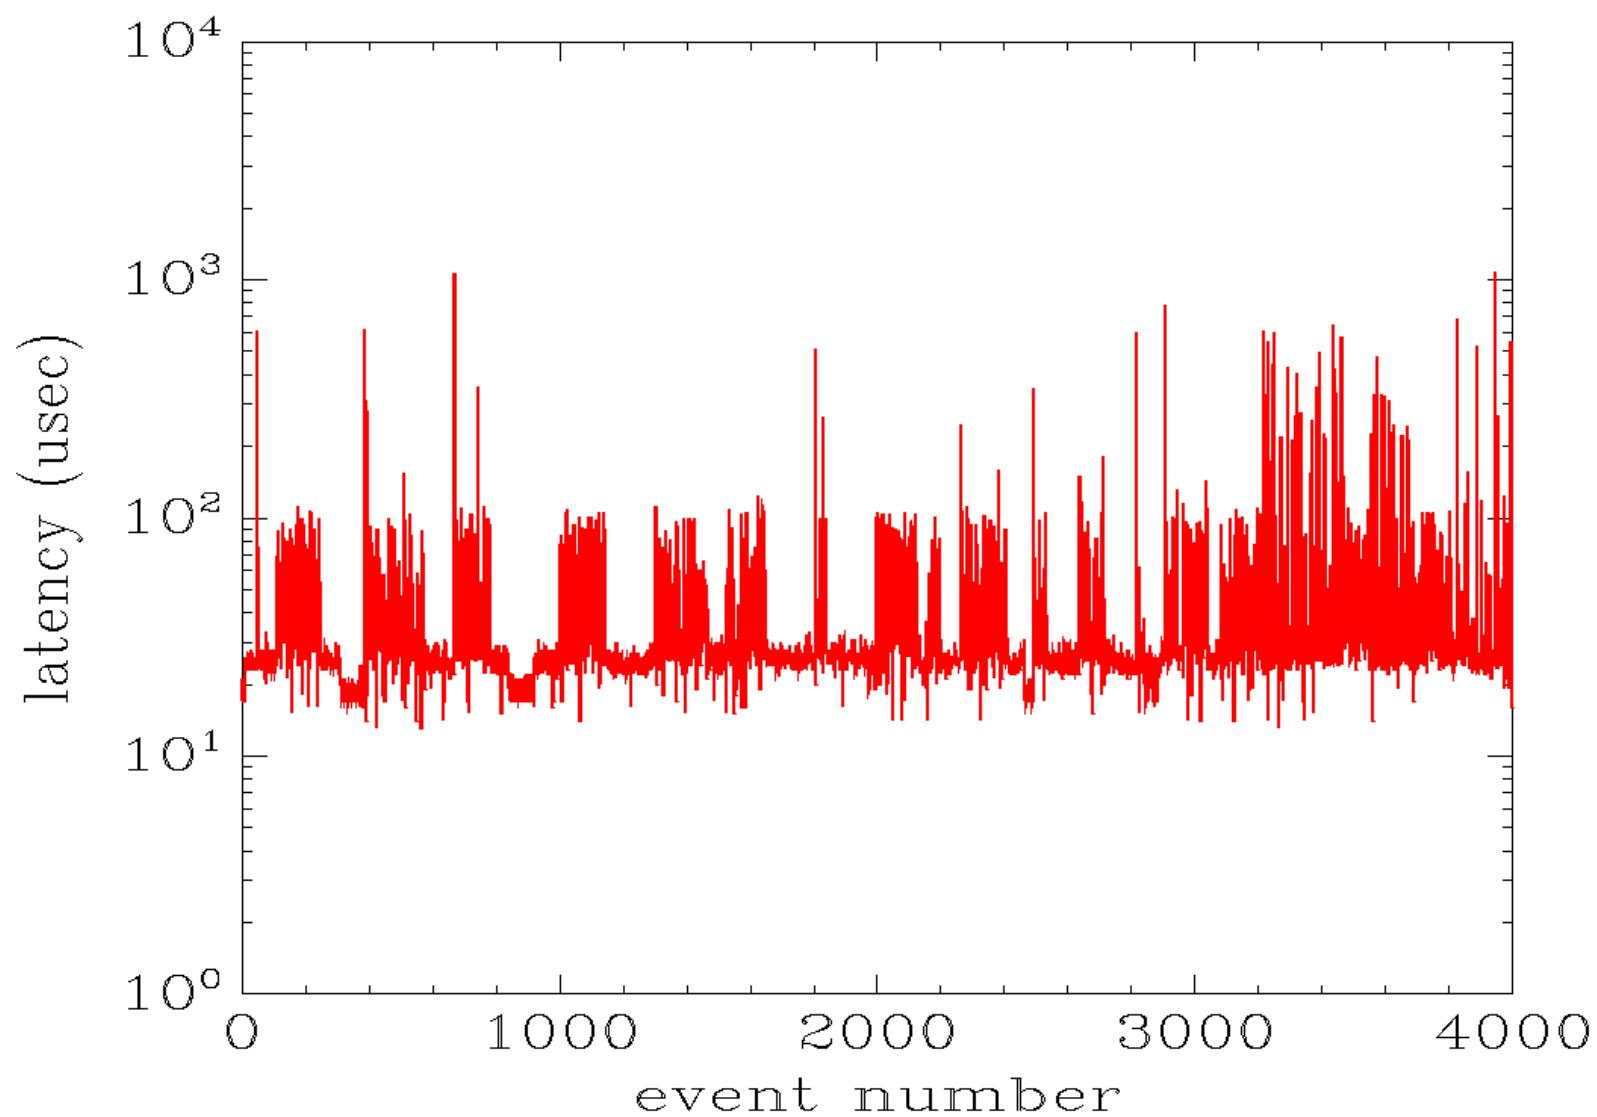
blue: pri better than load red: pri worse than load
load: scp



pri == 80 (better than load)
load: scp



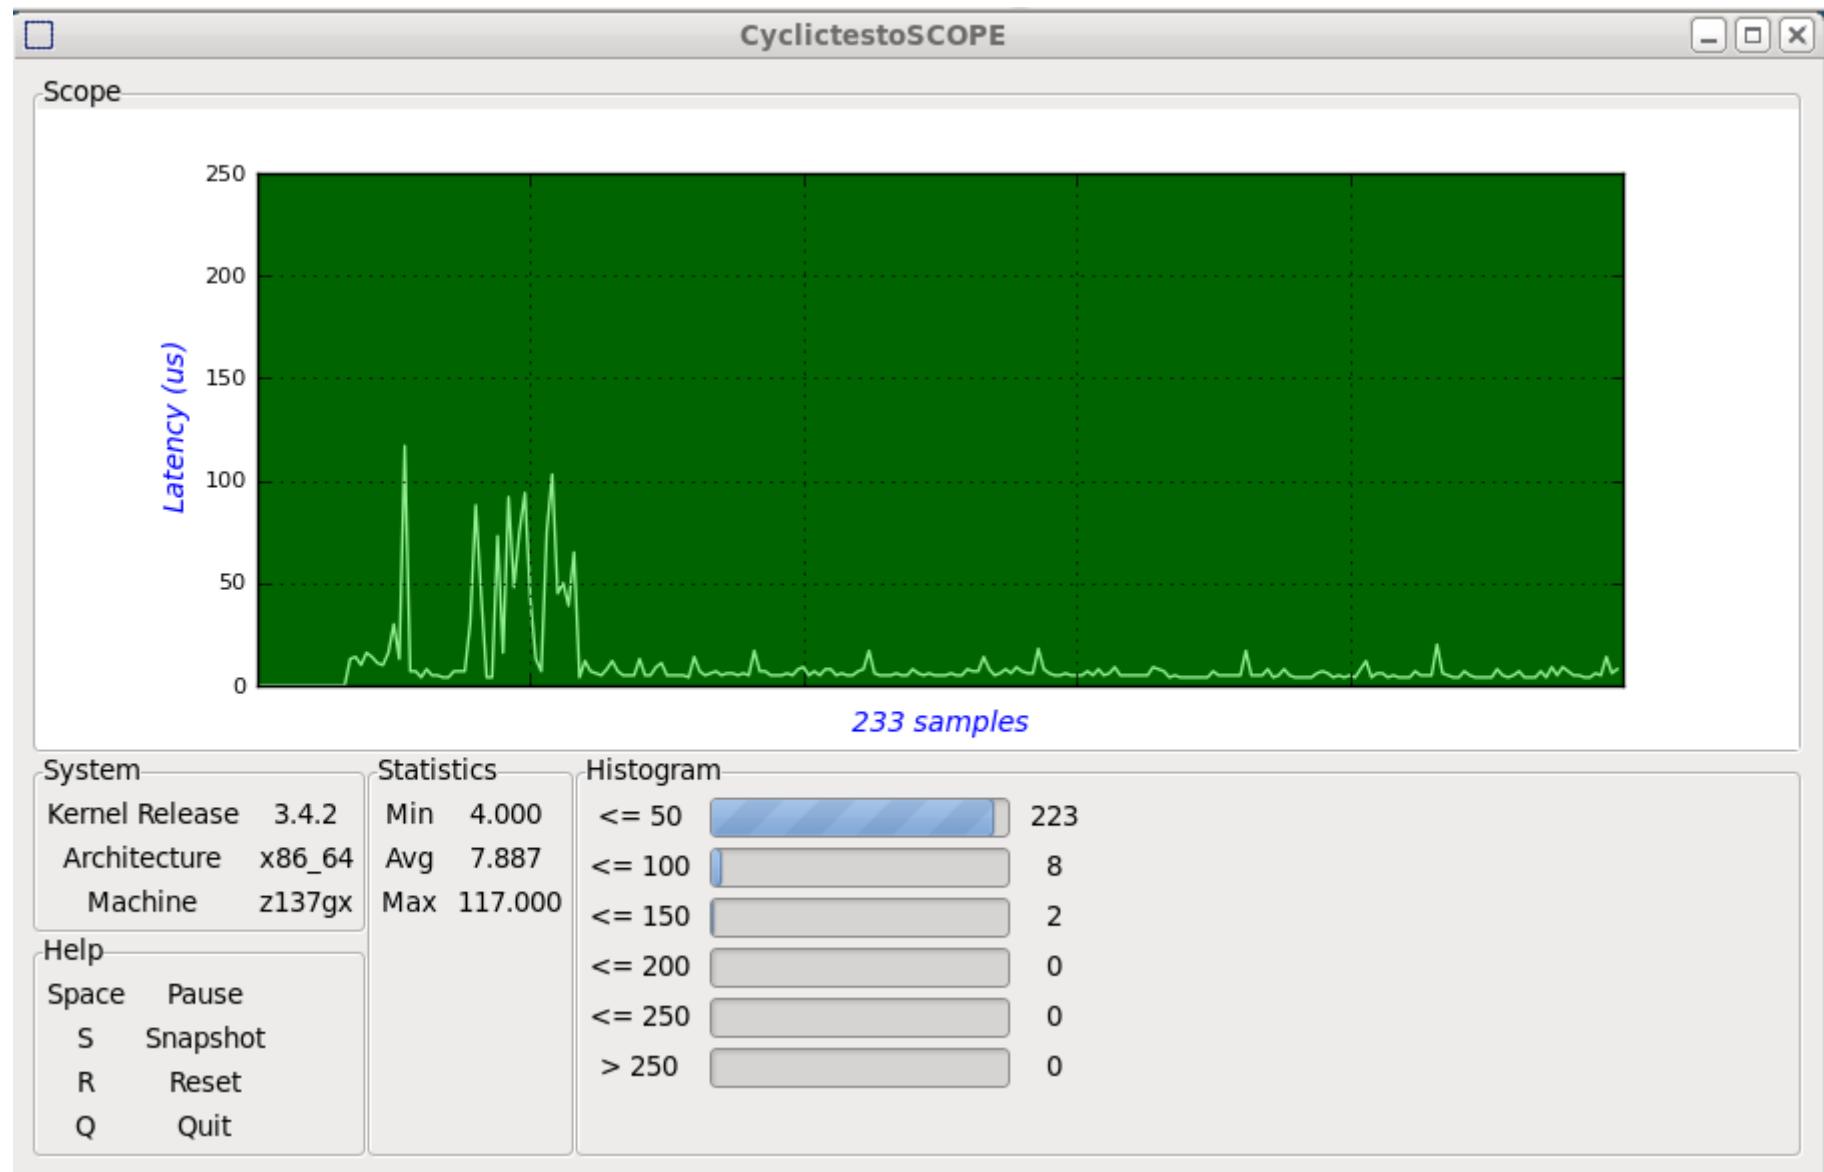
pri == 30 (worse than load)
load: scp



Demo - oscilloscope

```
cyclictest_0.85 -t1 -n -p80 -i100 -o10 -v \
| oscilloscope >/dev/null
```

oscilloscope screen shot



Fourth Data Format

Report time of each histogram overflow

Should be in next version of cyclictest (0.86?)

```
$ cyclictest -q -h 400 -g 1000
```

The same information can be extracted from the third data format (-v), but this method is lower overhead.

Finding and Building

```
git clone \
  git://git.kernel.org/pub/scm/linux/kernel/git/clrkwillms/rt-tests.git
source: src/cyclictest/cyclictest.c
```

self-hosted:
make

self-hosted without NUMA:
make NUMA=0

cross-build without NUMA:
make NUMA=0 CC="\${CROSS_COMPILE}gcc"

Review

- Simple methodology captures all sources of latency fairly well
- Options must be used with care
- Options are powerful
- Different data formats are each useful
- Debug features can capture the cause of large latencies

THE END

Thank you for your attention...

Questions?

How to get a copy of the slides

- 1) leave a business card with me
- 2) frank.rowand@sonymobile.com

