



8:25

A close-up photograph of a digital clock's display panel. The time is shown in large, yellow, seven-segment digits. The display reads "8:25". The background is dark, and the digits have a slight glow.

Siemens Corporate Technology | May 2013

# Virtualizing Real-Time Systems with Linux

# Virtualizing Real-Time Systems with Linux

## Agenda

### Motivation & Background

Real-Time Virtualization Basics

Linux/KVM-based Setup

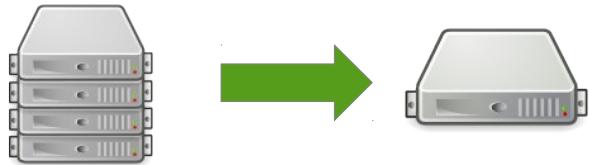
A New Approach for Lowest Latency

Summary

# 4 Main Reasons to Use Virtualization

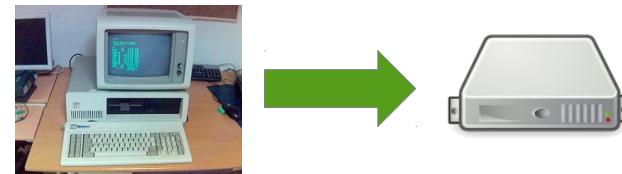
## Consolidation

- Energy saving
- Lower bill of material
- Maintenance reduction
- Integrate multiple levels of safety & security

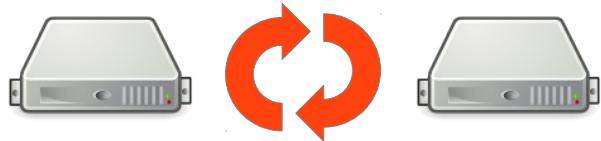


## Legacy system migration

- Stable hardware interfaces
- Multicore exploitation



## High availability



## Development & test environment



# Real-Time Systems Can Benefit from Virtualization

Virtualizable real-time systems

- **Possible scenarios**

- Control systems (industry, healthcare, etc.)
- Communication systems  
(media streaming & switching, etc.)
- Trading systems (stocks, goods, etc.)
- ...

- **Primary drivers**

- Consolidation
- Legacy system migration
- [Development & test]



Images: Ethernet switch by Ben Stanfield, licensed under CC BY-SA 2.0, stock market by Katrina.Tuliao, licensed under CC BY 2.0

# Virtualizing Real-Time Systems with Linux

## Agenda

Motivation & Background

### **Real-Time Virtualization Basics**

Linux/KVM-based Setup

A New Approach for Lowest Latency

Summary

# RT Guests – From Very Simple to Ordinary OS

## Characteristics of Real-Time Guests

Guest types
<ul style="list-style-type: none"><li>• Bare-metal (no OS)</li><li>• Classic real-time operating system (RTOS)</li><li>• Asymmetric multi-processing setup (RTOS + GPOS)</li><li>• General purpose OS with RT requirements</li></ul>



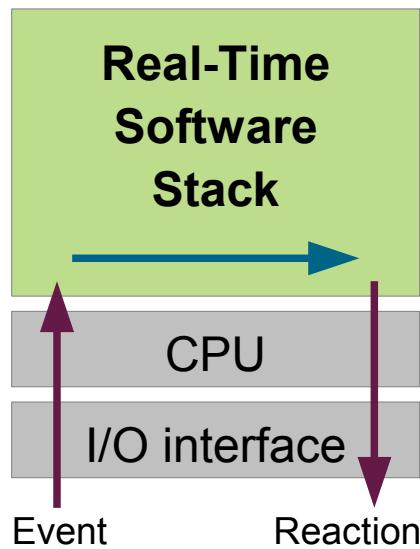
Hard RT guest resources
<ul style="list-style-type: none"><li>• CPU (bandwidth, event latency)</li><li>• Timers &amp; clocks</li><li>• RT networks (RT Ethernet, fieldbuses, ...)</li><li>• Digital/analogue I/O interfaces</li><li>• Custom hardware</li></ul>

Soft RT guest resources
<ul style="list-style-type: none"><li>• Anything on the left</li><li>• Standard networks</li><li>• Mass storage</li><li>• Human-machine interfaces</li><li>• ...</li></ul>

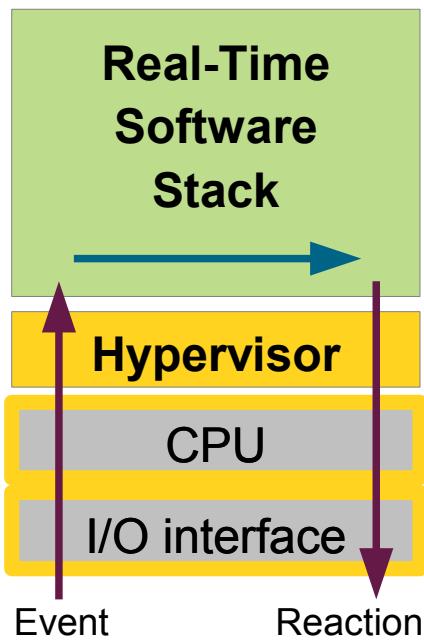
Image: by Daniel Leininger, licensed under CC BY 2.0

# Virtualization ≠ Acceleration

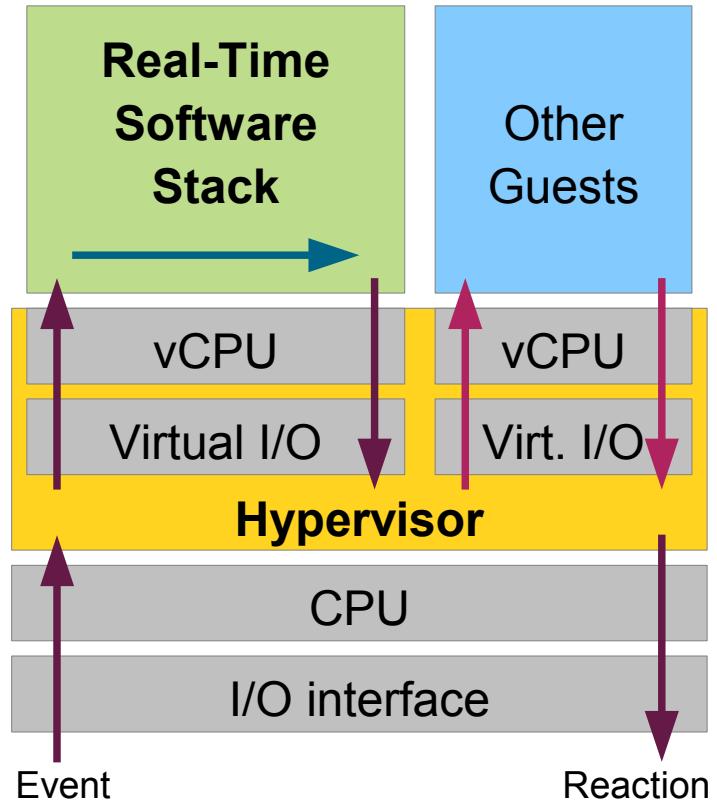
The Critical Data Path with and without Virtualization



**Native real-time setup**

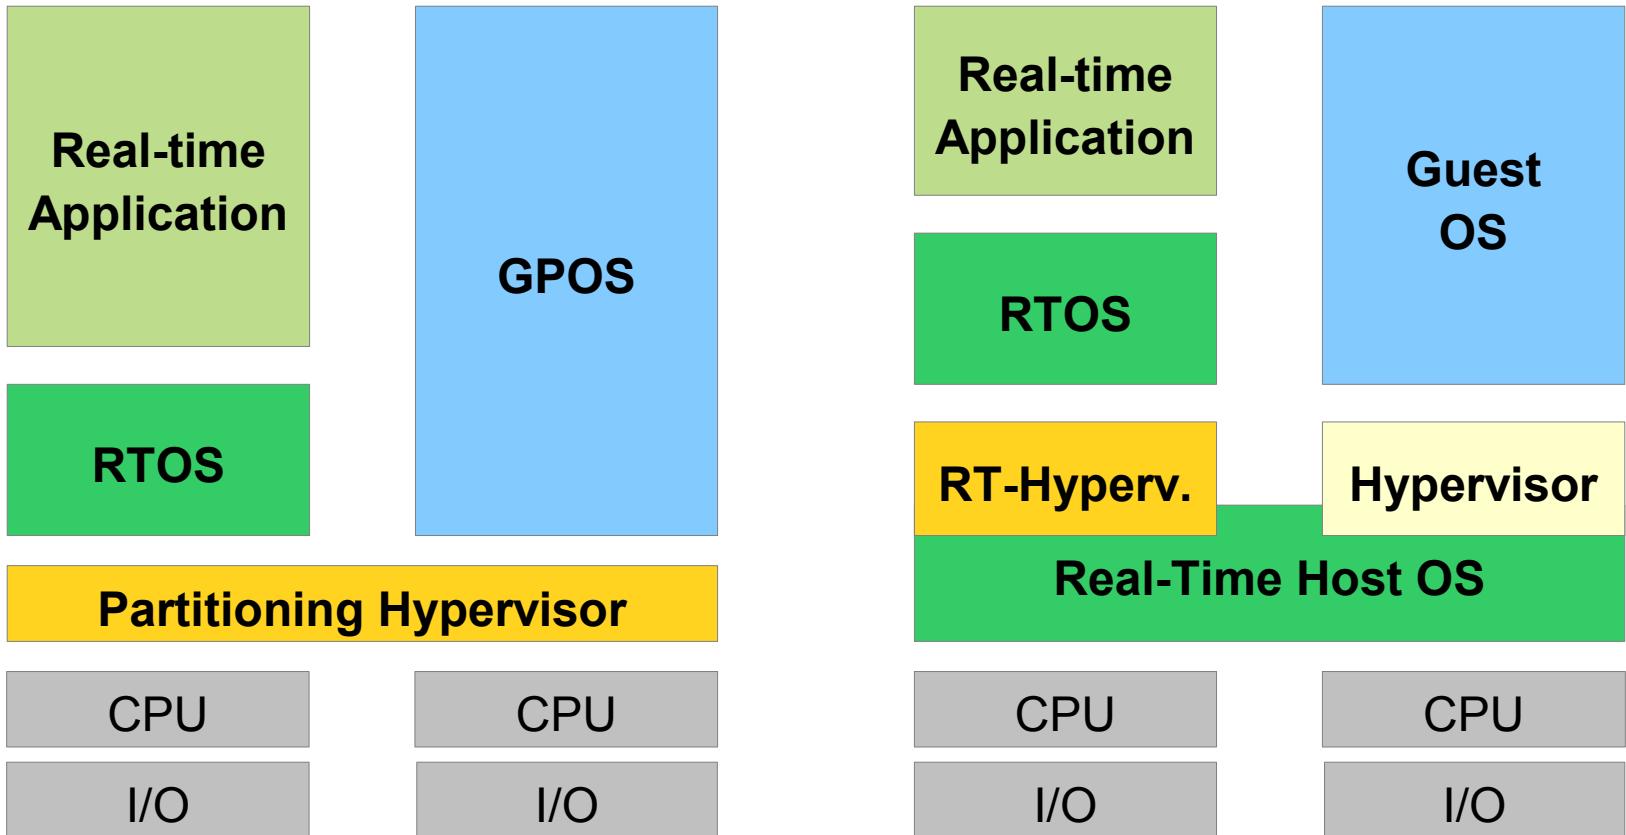


**Hardware-assisted virtualization**



**Resource sharing & abstraction via emulation**

# RT Virtualization – Two Architectural Options



# No “One Size Fits it All”

## Comparison of RT-Virtualization Approaches

Partitioning hypervisor	RT-enhanced hosted hypervisor
<ul style="list-style-type: none"><li>• <b>Reduced hypervisor complexity</b><ul style="list-style-type: none"><li>• Simplifies safety/security validation</li><li>• Can provide better real-time properties</li></ul></li><li>• <b>Mostly static assignments</b></li><li>• <b>No or limited overcommitment</b></li><li>• <b>Duplication of system boot-up logic</b></li><li>• <b>Often requires paravirtualization</b></li><li>• <b>Non-standard management interfaces</b></li><li>• <b>No industry-grade free solutions available</b></li></ul>	<ul style="list-style-type: none"><li>• <b>Reuse of existing solutions possible</b><ul style="list-style-type: none"><li>• Works well with common guests</li><li>• Up-to-date virtualization features</li><li>• Large user base</li><li>• Free Linux-based solution available</li></ul></li><li>• <b>Features vs. complexity</b></li><li>• <b>Adding RT support can be non-trivial task</b></li></ul>

# Virtualizing Real-Time Systems with Linux

## Agenda

Motivation & Background

Real-Time Virtualization Basics

**Linux/KVM-based Setup**

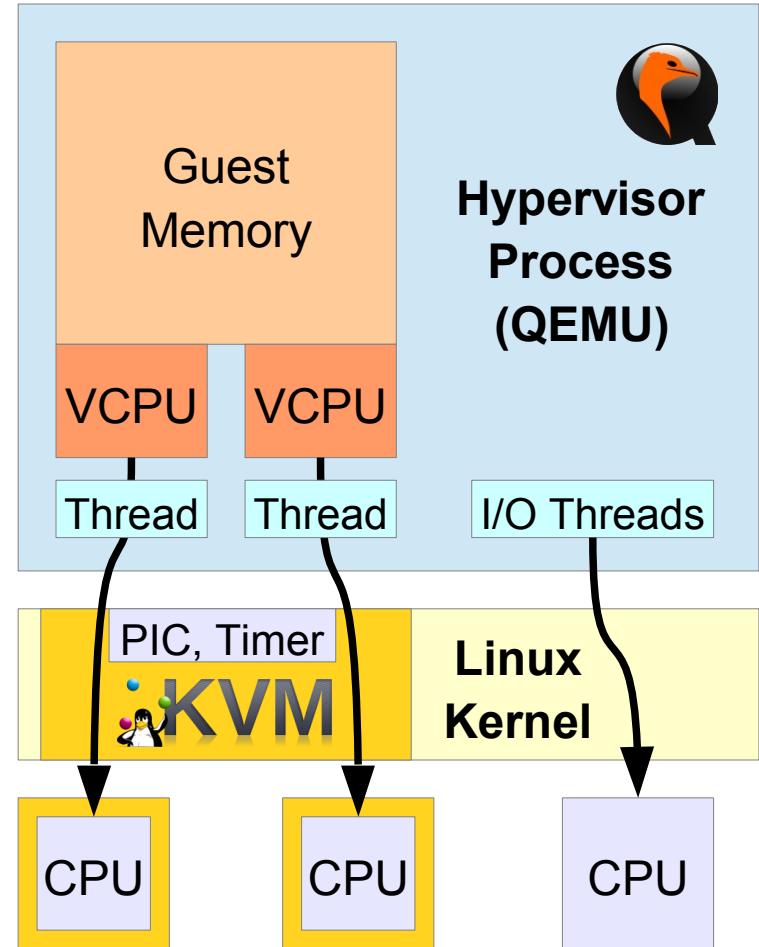
A New Approach for Lowest Latency

Summary

# KVM – the Swiss Knife among Hypervisors

## Main Components of the QEM/KVM Hypervisor Stack

- **KVM subsystem**
  - Gatekeeper for HW- and kernel-assisted virtualization
  - Virtual CPUs
  - Fast device models
- **Use Linux for other hypervisor tasks**
  - Scheduling
  - Memory management
  - I/O stacks
  - Device pass-through (VFIO)
  - ...
- **User-space hypervisor QEMU**
  - Multi-architecture machine emulator
  - Can use KVM for virtual CPUs
  - Large set of device models

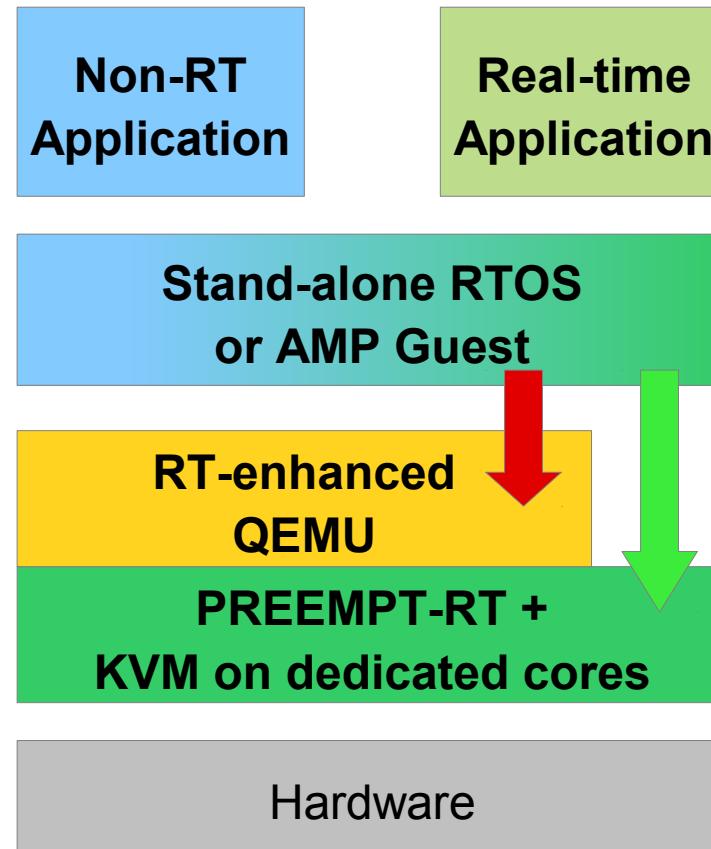


# PREEMPT-RT enables RT-Virtualization

Role of Linux extension PREEMPT-RT

- **PREEMPT-RT: fully preemptible Linux kernel**
- **Typical worst-case event delivery latencies: few 10 microseconds**
- **Integrates KVM support**
  - Original use-case: virtualization + native RT applications
- **Allows to prioritize virtualization workload over uncritical tasks**
- **Can be combined with CPU isolation**
  - 1:1 assignment: host CPU – RT guest CPU
  - Off-load all non-RT tasks  
(including low-priority QEMU threads)
  - Warning: No 100% guest CPU load feasible!
    - NO\_HZ\_FULL extensions work toward enabling this

# Architecture of a KVM-based RT-Hypervisor

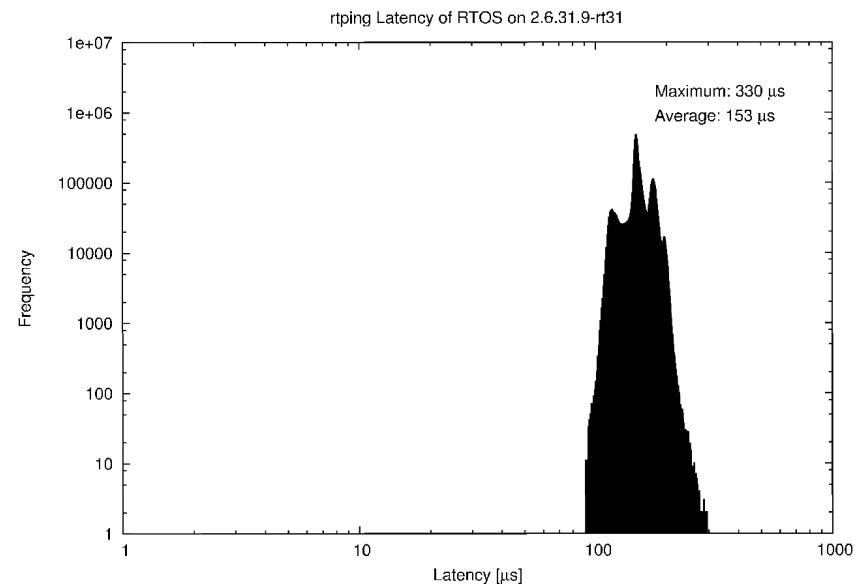
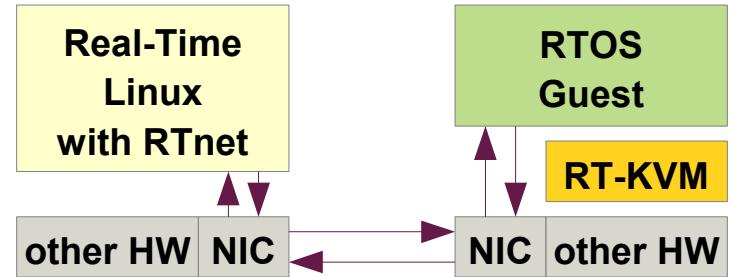


# Decent Latencies Achievable in KVM-only Setups

Measuring I/O latency of an RT Guest

- **Host setup**
  - KVM on x86 PREEMPT-RT Linux
  - Virtual machine on dedicated core
  - Intel NIC (E1000 family) as I/O device, directly assigned to guest
  - Permanent disk I/O load
- **Guest setup**
  - Proprietary RTOS
  - Real-time network stack
- **Measurement setup**
  - Linux/Xenomai (native installation)
  - Real-time network stack RTnet
  - Periodic ICMP ping messages sent to target
  - Recorded round-trip latency (error <50 µs)

**=> Worst-case latency after 16h: 330 µs**



# RT-QEMU is Required for Emulating in Real-Time

QEMU as a Real-Time Device Emulator

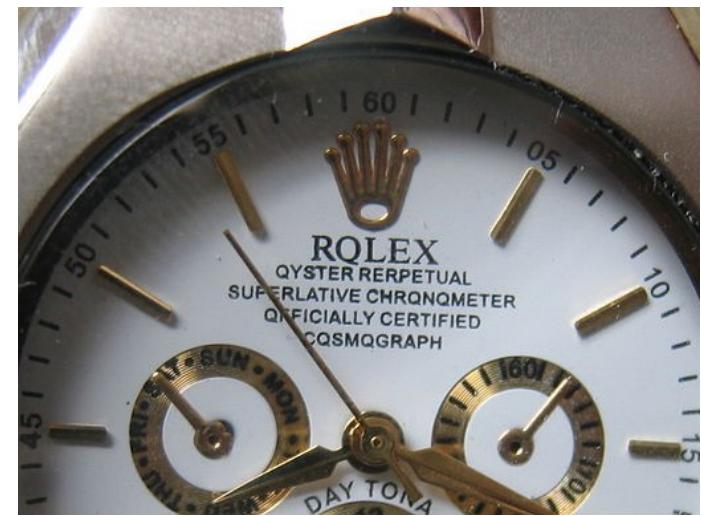
- **Scenarios**

- Guest uses NIC A, host has NIC B attached
- Legacy devices are no longer available on a modern host
- Multiple guests share one I/O interface for talking to different devices (e.g. on a CAN bus)

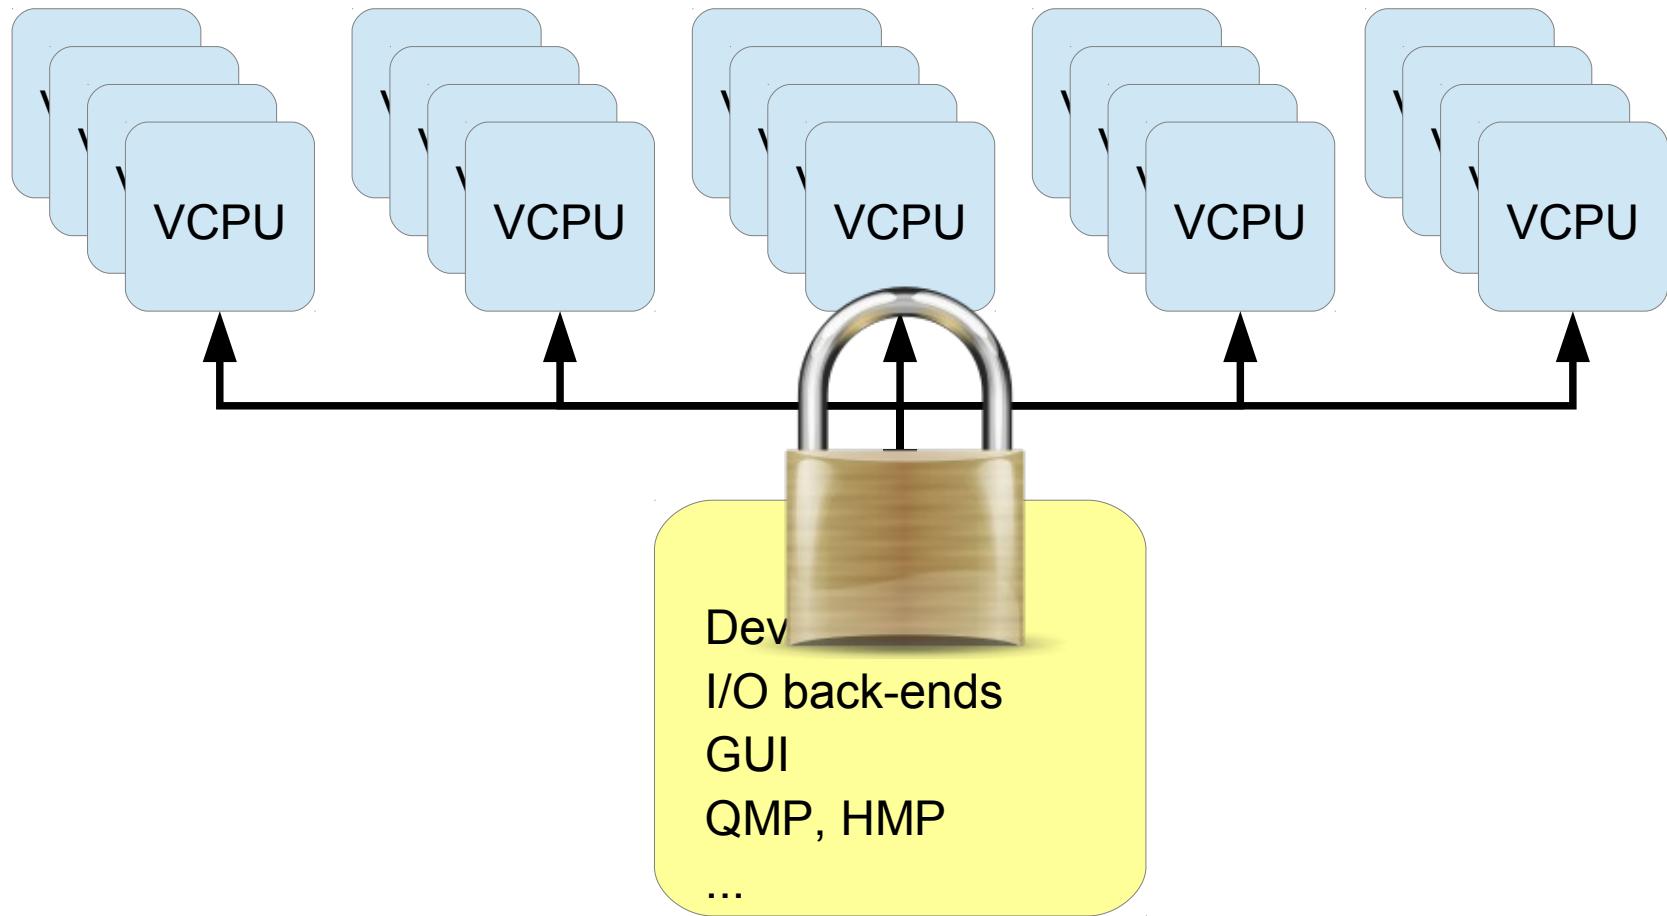
- **QEMU can handle such scenarios via emulation**

- **Requirements on emulation**

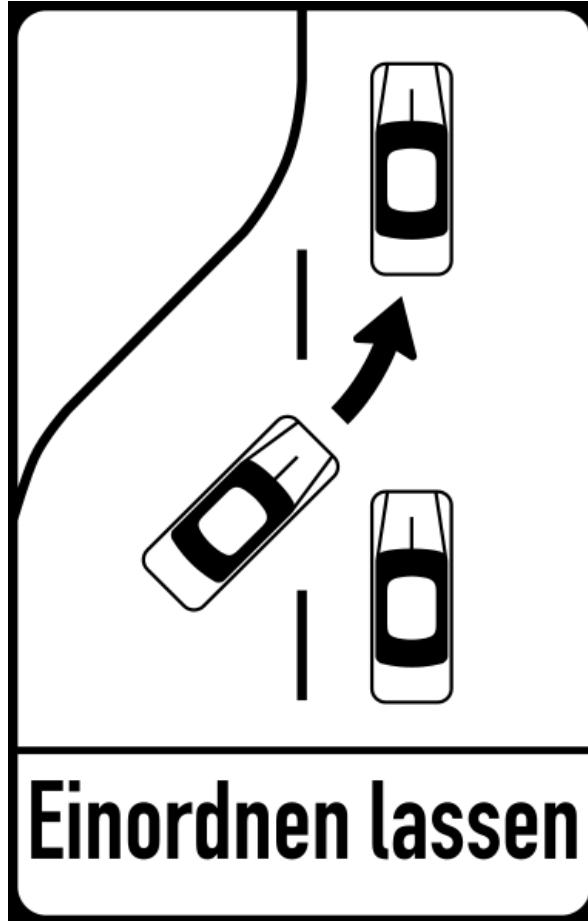
- Equivalent functional behavior
- Devices models need to react in time on guest requests
- Devices models need to deliver external events to the guest timely



# Concurrency in QEMU/KVM – The Big QEMU Lock (BQL)



## BQL – One after the other



# Breaking the BQL is Challenging but Feasible

Challenges in Reworking QEMU's Concurrency Scheme

- Establish lock ordering rules
- Prevent lock recursions
- Design efficient object locking and event dispatching scheme
  - Which objects to track?
  - Reference counting? Mutexes? RCU?
- Enable smooth transition from old to new model
  - Existing code base cannot be converted in a single step
  - Both models must coexists
- Incrementally convert device models and backends
- Teach new model to developers



# Critical BQL Zones

## CPUState

- Read/write access
- cpu\_single\_env



## Coalesced MMIO flushing

## PIO/MMIO request-to-device dispatching

## Back-end access

- TX on network layer
- Write to character device
- Timer setup, etc.

## Back-end events (iothread jobs)

- Network RX, read from chardev, timer signals, ...

## IRQ delivery

- Raising/lowering from device model to IRQ chip
- Injection into VCPU (if user space IRQ chips)

# Virtualizing Real-Time Systems with Linux

## Agenda

Motivation & Background

Real-Time Virtualization Basics

Linux/KVM-based Setup

**A New Approach for Lowest Latency**

Summary

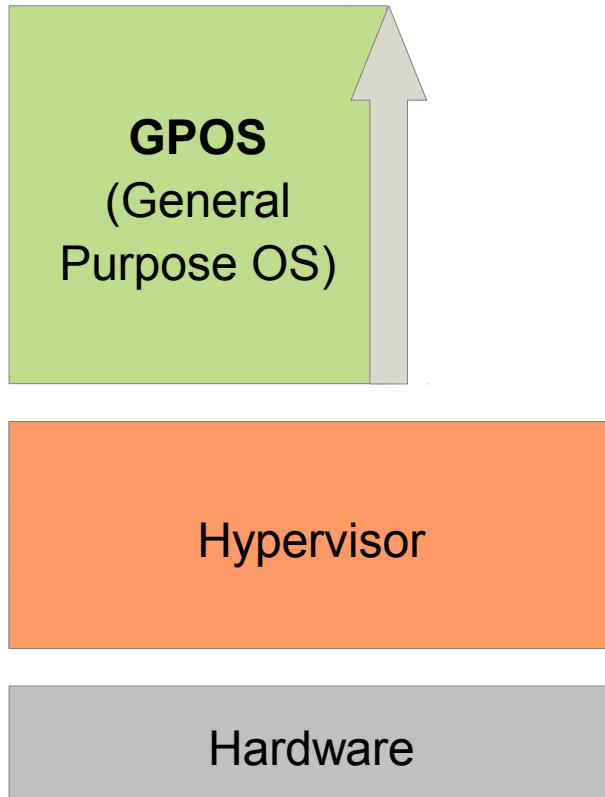
# No “one size fits it all”

Comparison of RT-virtualization approaches

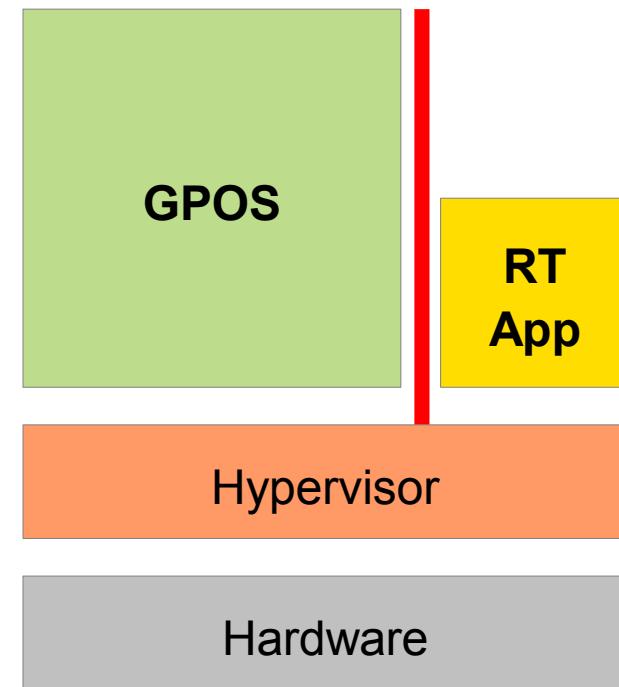
Partitioning hypervisor	RT-enhanced commodity hypervisor
<ul style="list-style-type: none"><li>• Reduced hypervisor complexity<ul style="list-style-type: none"><li>• Simplifies safety/security validation</li><li>• Can provide better real-time properties</li></ul></li><li>• Mostly static assignments</li><li>• No or limited overcommitment</li><li>• Duplication of system boot-up logic</li><li>• Often requires paravirtualization</li><li>• Non-standard management interfaces</li><li>• No industry-grade free solutions available</li></ul>	<ul style="list-style-type: none"><li>• Reuse of existing solutions<ul style="list-style-type: none"><li>• Work well with common guests</li><li>• Up-to-date virtualization features</li><li>• Large user base</li><li>• Free Linux-based solution available</li></ul></li><li>• Features vs. complexity</li><li>• Adding RT support can be non-trivial task</li></ul>

# A bare-metal hypervisor has to boot its guest

Classic type-1 hypervisor boot-up



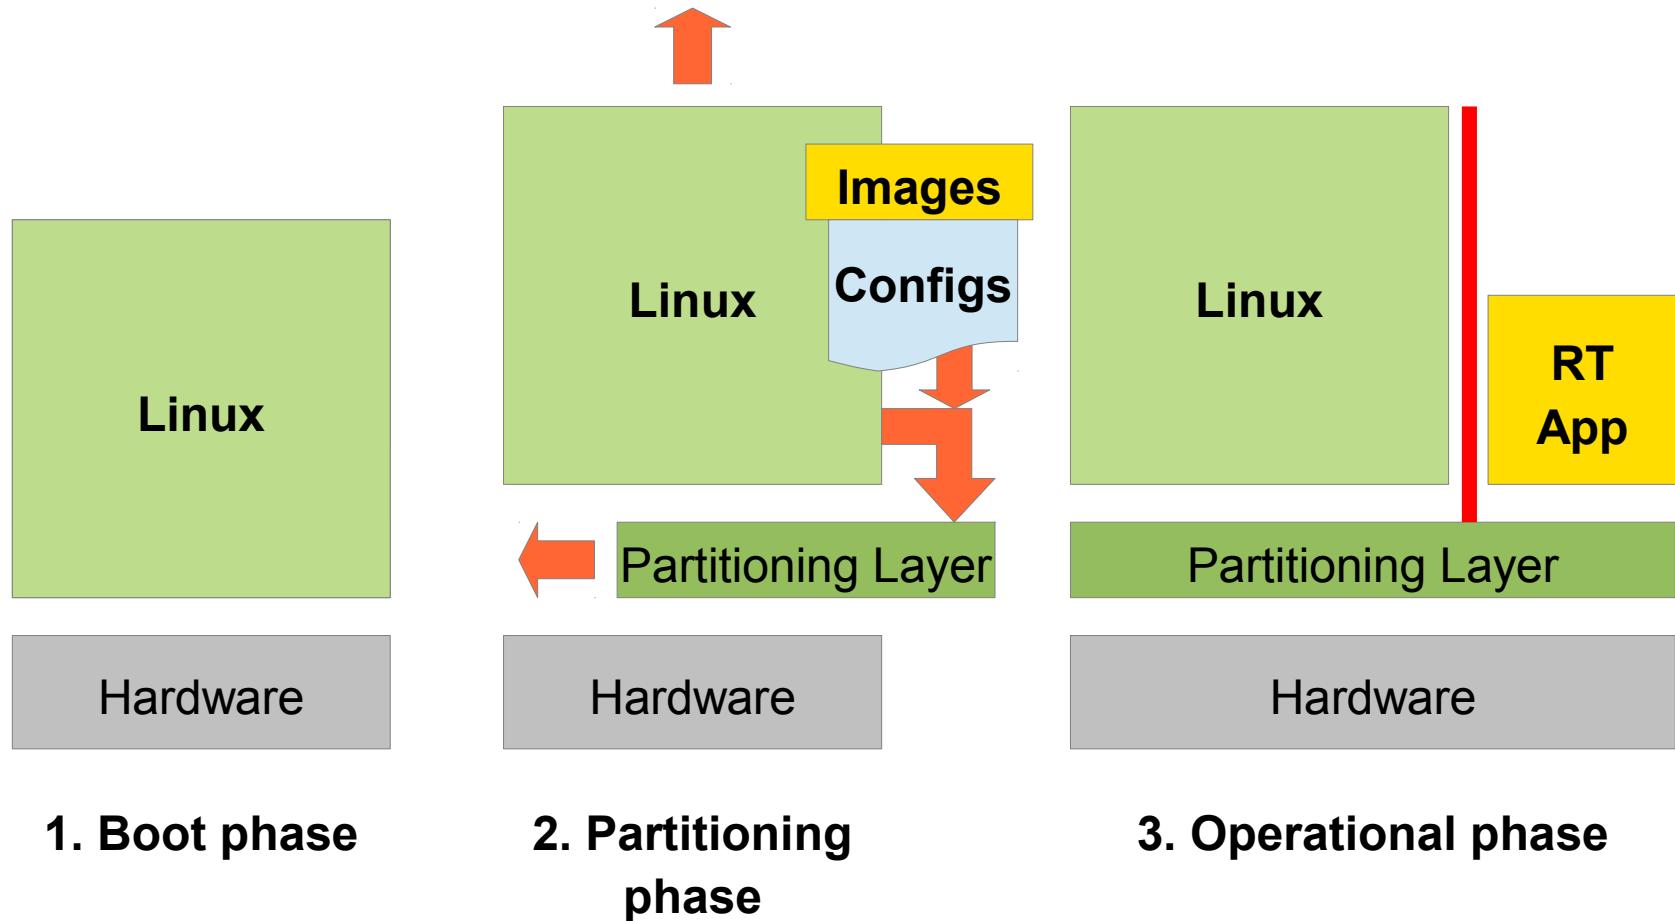
**1. Boot phase**



**2. Operational phase**

# What about postponing the hypervisor start?

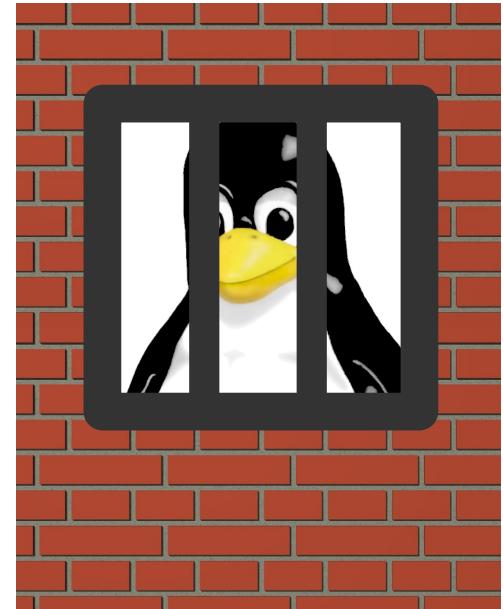
Basic concept of late partitioning



# Jailhouse: Keep it simple, keep it open

## The Philosophy of Jailhouse

- **Avoid emulation, focus on hardware assisted isolation**
    - No overcommitment, no scheduler, static partitioning
    - Directly assign physical devices, do not emulate them
    - You need more? Use KVM!
  - **Only expose resources that are required for operation**
    - No boot-up phase virtualization
    - Board initialization done by Linux
  - **Off-load uncritical tasks to Linux**
    - Initial setup / image loading
    - Reconfigurations while in non-operational mode
    - Monitoring, logging etc.
  - **Release as Open Source** (ongoing process – stay tuned)
- => Minimal-sized hypervisor with full CPU assignment and Linux look-and-feel**



# Virtualizing Real-Time Systems with Linux

## Agenda

Motivation & Background

Real-Time Virtualization Basics

Linux/KVM-based Setups

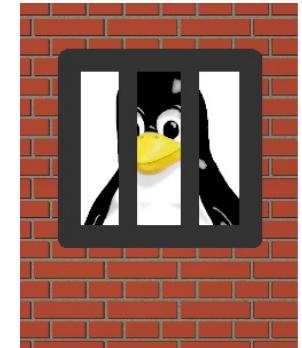
A New Approach for Lowest Latency

**Summary**

# Real-Time – The Next Level for Virtualization

## Summary

- **Virtualization provides benefits also for real-time systems**
    - Consolidation
    - Legacy migration (to multicore)
  - **Two approaches**
    - Enhance full-featured hypervisor
    - Provide dedicated partitioning hypervisor
  - **Open Source based solutions in reach**
    - PREEMPT-RT + KVM + device pass-through works
    - RT device emulation via QEMU is approaching
    - Static partitioning as Linux feature via Jailhouse is WiP
- => Real-time system engineering remains challenging – with or without virtualization**



**Any Questions?**

**Thank you!**

Jan Kiszka <[jan.kiszka@siemens.com](mailto:jan.kiszka@siemens.com)>