

# ELI: Bare-Metal Performance for I/O Virtualization

Abel Gordon<sup>×</sup> Nadav Amit<sup>¤</sup> Nadav Har'El<sup>×</sup>  
Muli Ben-Yehuda<sup>\*,¤</sup> Alex Landau<sup>×</sup> Assaf Schuster<sup>¤</sup> Dan Tsafrir<sup>¤</sup>

<sup>×</sup> IBM Research – Haifa

<sup>¤</sup> Technion – Israel Institute of Technology



Seventeenth International Conference on  
Architectural Support for Programming Languages and Operating  
Systems (ASPLOS 2012)  
London, UK, March 3–7 2012



## Why (not) Virtualization ?

### ■ Pros

- Consolidation
- Flexibility
- Cloud computing
- Simplified management
- Migration
- ....




### ■ Cons


- Performance degradation
- Performance degradation
- Performance degradation



## x86 Virtualization Performance



## The Cause of x86 I/O Virtualization Overhead




- Privileged instructions trap to the hypervisor
- Traps cause an *Exit* (switch to the hypervisor context)
- Hypervisor emulates their behavior
- Hypervisor resumes the guest (switch to the guest context)

I/O intensive workloads cause  
many exits = overhead!


## Direct Device Assignment

- Give the guest direct access to a physical device
- Devices nowadays know how to “self virtualize” (SR-IOV)
- Main I/O control and data paths do not require hypervisor intervention




## The Solution (?) to I/O Virtualization Overhead

- Device assignment is by far the best performing option, but it only achieves 60% of bare-metal performance




## Interrupt, Exits, Interrupt, Exits, Interrupt, Exits...




- No hypervisor intervention (?!)... **except for handling more than 48,000 interrupts per second!**
  - Exits due to external interrupts
  - Exits due to interrupt completions

## ELI: Exitless Interrupts for unmodified, untrusted guests




## Background: x86 Interrupt Handling

- I/O devices raise interrupts to notify the system software about incoming events
- CPU temporarily stops the currently executing code and jumps to a pre-specified interrupt handler




## x86 Interrupt Handling in Virtual Environments



- Host IDT: handles physical interrupts
- Guest IDT: handles virtual interrupts
- If a physical interrupt –**including interrupts from assigned devices**– arrives while the guest is running the **CPU forces an Exit!**

## ELI: Exitless Interrupts - Delivery

- ELI configures the CPU to deliver all interrupts to the guest
- ELI runs the guest using a shadow IDT
- Host interrupts are bounced back to the host in the form of exceptions
- ...without the guest being aware of it



## ELI: Exitless Interrupts - Completion

- Guests write to the LAPIC EOI register
- Old LAPIC interface:
  - Hypervisor traps memory accesses → **page granularity**
- New LAPIC interface (x2APIC), required for Exitless Completions
  - Hypervisor traps accesses to MSRs → **register granularity**


ELI gives direct access  
only to the EOI register



## Evaluation: Netperf


**102,000**  **800**

Exits/sec




**60%**  **98%**

Time in guest



**bare-metal  
performance!**



## Evaluation: Apache


91,000  1,100

Exits/sec

65%  97%

Time in guest

bare-metal  
performance!



## Evaluation Memcached


**123,000**  **1,000**

Exits/sec

**60%**  **100%**

Time in guest

bare-metal  
performance!




## Evaluation: Latency

Configuration	Avg. Latency	% of bare-metal
Baseline	36.14 $\mu$ s	129%
ELI delivery-only	30.10 $\mu$ s	108%
ELI	28.51 $\mu$ s	102%
Bare-metal	27.93 $\mu$ s	100%

Netperf UDP Request Response


ELI reduces the time it takes to deliver interrupts, critical for latency-sensitive workloads.

## Evaluation: Interrupt Coalescing (netperf)




Aggressive Interrupt coalescing → No Interrupt coalescing

## Evaluation: Computation vs. I/O Ratio (throughput)



## Evaluation: Computation vs. I/O Ratio (interrupt rate)



## Mitigating Interrupts Overhead - Related Work

- Polling
- Hybrid
- Interrupt Coalescing

- (1) ELI is **complementary** to these approaches. **Removes the virtualization overhead** caused by the costly exits and entries during interrupt handling.
  - (2) ELI allows the guest to choose the desired method and achieve bare-metal performance in any-case.

Coming Soon....

- Reduce frequency of exits caused by para-virtual I/O devices
- Reduce frequency of exits caused by non-assigned interrupts
- Reduce exits required to inject a virtual interrupt
- Improve performance of SMP workloads with high Inter-processor interrupt (IPI) rate
- Improve performance of Nested Virtualization

## Conclusions

- High virtualization performance requires the CPU to spend most of the time running the guest (useful work) and not the host (handling ~~exits~~ overhead)  

- x86 virtualization requires host involvement ~~(exits)~~ to handle interrupts (critical path for I/O workloads)  

- ELI lets the guest handle interrupts directly (no ~~exits~~) and securely, making it possible for untrusted and unmodified guests reach near bare-metal performance  




Questions ?